




Machine Learning

第12讲 强化学习

Reinforcement Learning

刘 峤

电子科技大学计算机科学与工程学院



Roadmap

- Introduction
 - ※ Terminologies.
 - ※ Rewards, Returns, and Value functions
 - ※ Play games using reinforcement learning
- Value-based learning
 - ※ Deep Qnetwork (DQN) for approximating $Q^*(s,a)$.
 - ※ Learn the network
 - ※ parameters using temporal different (TD).
- Policy-based learning.
 - ※ Policy network for approximating $\pi(a|s)$.
 - ※ Learn the network parameters using policy gradient.
- Actor-critic method. (Policy network + value network.) 自学

12.1 Introduction

12.1.1 Terminologies

Random Variable

- Random variable:
 - ※ a variable whose values depend on outcomes of a random event.
 - ※ Uppercase letter ***X*** for random variable.
 - ※ Lowercase letter ***x*** for an observed value.
 - ※ For example, we flipped a coin 4 times and observed:
 - ※ $x_1 = 1, \quad x_2 = 1, \quad x_3 = 0, \quad x_4 = 1$
- Probability Density Function (PDF)
 - ※ PDF provides a relative likelihood that the value of the random variable would equal to that sample.

Expectation

- Random variable **X** is in the domain \mathcal{X} .
- For continuous distribution, the expectation of $f(X)$ is:

$$\mathbb{E}[f(x)] = \int_{\mathcal{X}} p(x) \cdot f(x) dx$$

- For discrete distribution, the expectation of $f(X)$ is:

$$\mathbb{E}[f(x)] = \sum_{x \in \mathcal{X}} p(x) \cdot f(x)$$

Random Sampling

- Sample red ball w.p. 0.2, green ball w.p.0.5, and blue ball w.p.0.3.
 - ※ Randomly sample a ball.
 - ※ What will be the outcome?

```
from numpy.random import choice
```

```
samples = choice( ['R', 'G', 'B' ], size=100, p=[0.2, 0.5, 0.3])  
print (samples)
```

```
['G' 'G' 'G' 'R' 'R' 'G' 'B' 'G' 'G' 'B' 'B' 'G' 'G' 'R' 'R' 'G' 'R' 'B'  
'R' 'G' 'G' 'B' 'B' 'G' 'B' 'R' 'G' 'G' 'B' 'B' 'G' 'G' 'G' 'G' 'G'  
'B' 'G' 'B' 'B' 'G' 'R' 'G' 'G' 'G' 'B' 'B' 'R' 'G' 'R' 'G' 'B' 'G' 'B'  
'R' 'R' 'G' 'R' 'R' 'R' 'R' 'B' 'R' 'B' 'B' 'G' 'G' 'R' 'B' 'G' 'G' 'G'  
'G' 'G' 'R' 'B' 'B' 'G' 'G' 'G' 'B' 'B' 'B' 'B' 'B' 'B' 'R' 'B' 'B' 'G'  
'R' 'B' 'B' 'G' 'G' 'G' 'G' 'G' 'B' 'G']
```

Paradigm



**Supervised
Learning**



**Unsupervised
Learning**



**Reinforcement
Learning**

Objective

$$p_{\theta}(y|x)$$

$$p_{\theta}(x)$$

$$\pi_{\theta}(a|s)$$

Applications

→ **Classification**

→ **Regression**

→ **Inference**

→ **Generation**

→ **Prediction**


→ **Control**

Terminology: state and action

State s (this frame)

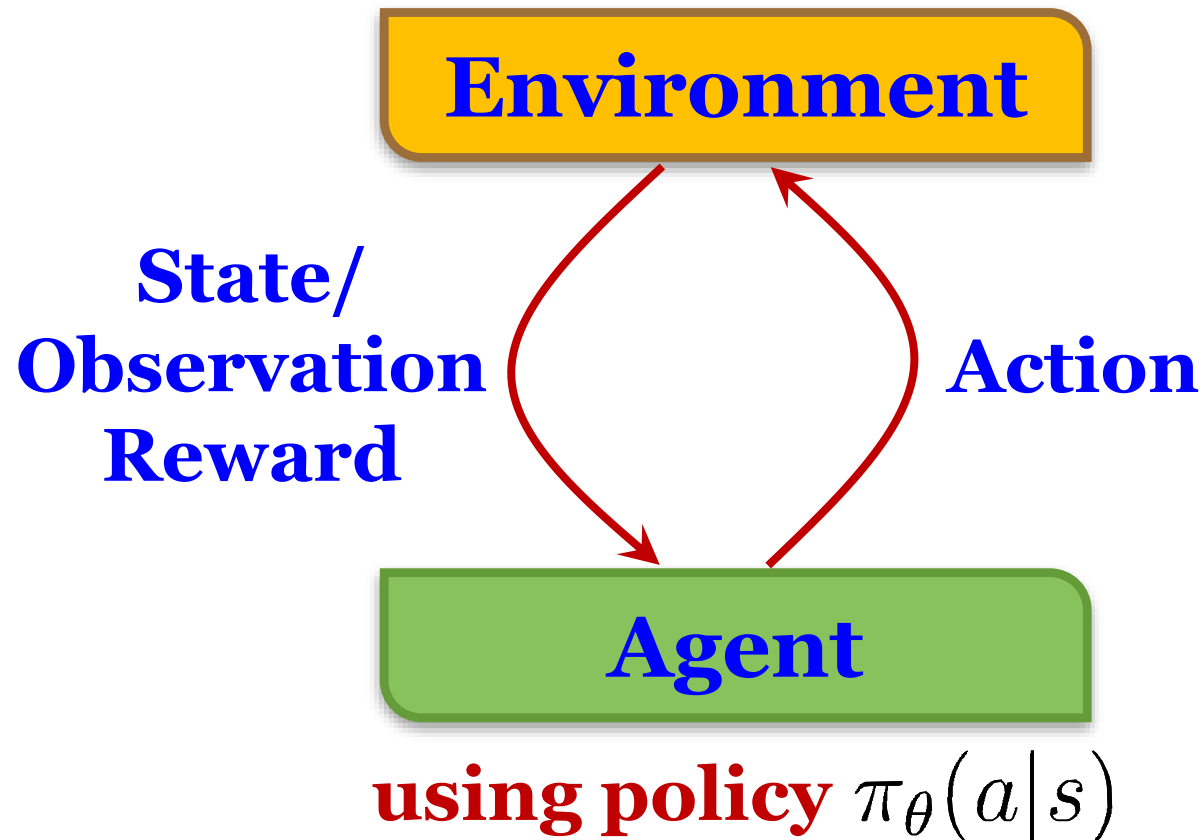


Action $a \in \{\text{left, right, up}\}$



Terminology: policy

Setting



policy π • Policy function $\pi : (s, a) \rightarrow [0, 1]$

$$\pi(a|s) = P(A = a|S = s)$$

- It is the probability of taking action $A = a$ given state s , e.g.,

$$\pi(\text{left}|s) = 0.2$$

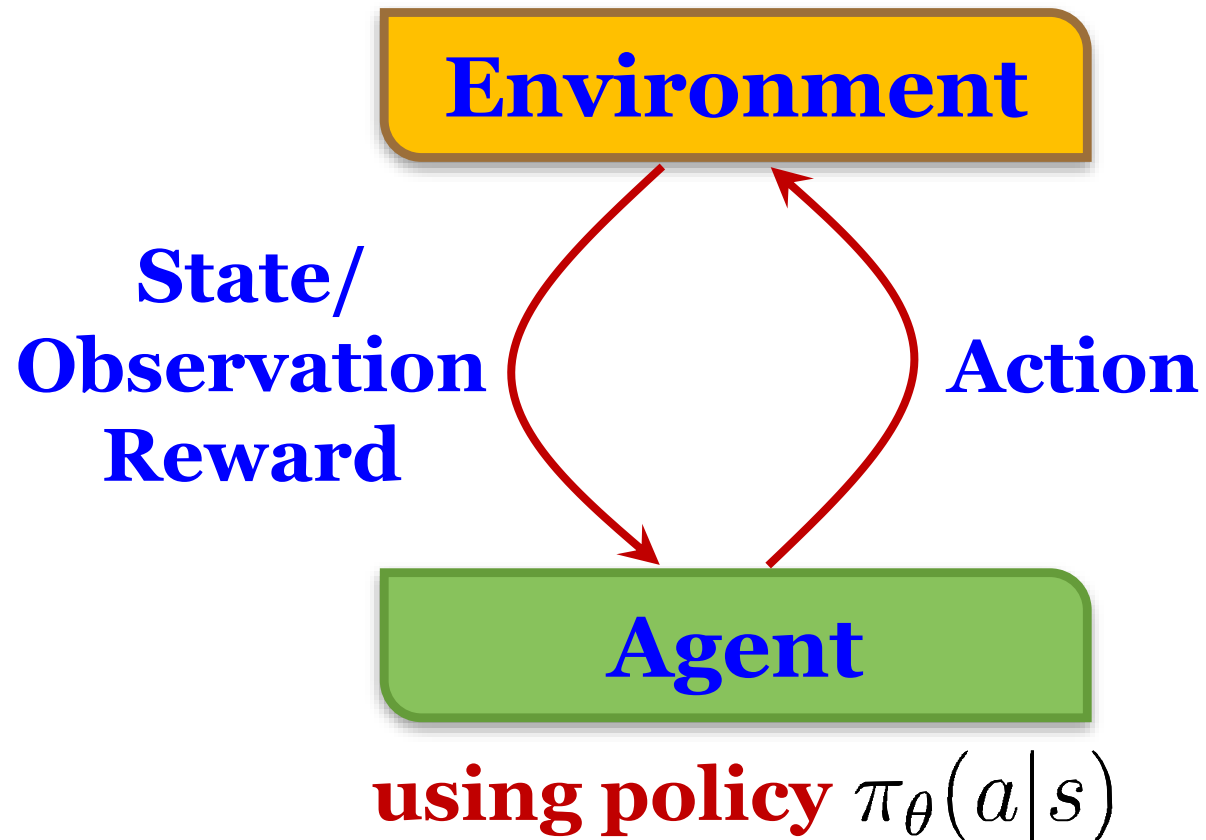
$$\pi(\text{right}|s) = 0.1$$

$$\pi(\text{up}|s) = 0.7$$

- Upon observing state $S = s$, the agent's action A can be random.

Terminology: reward

Setting



reward R

- Collect a coin: $R = +1$
- Win the game: $R = +10000$
- Touch a Goomba: $R = -10000$
※ game over
- Nothing happens: $R = 0$

Terminology: state transition

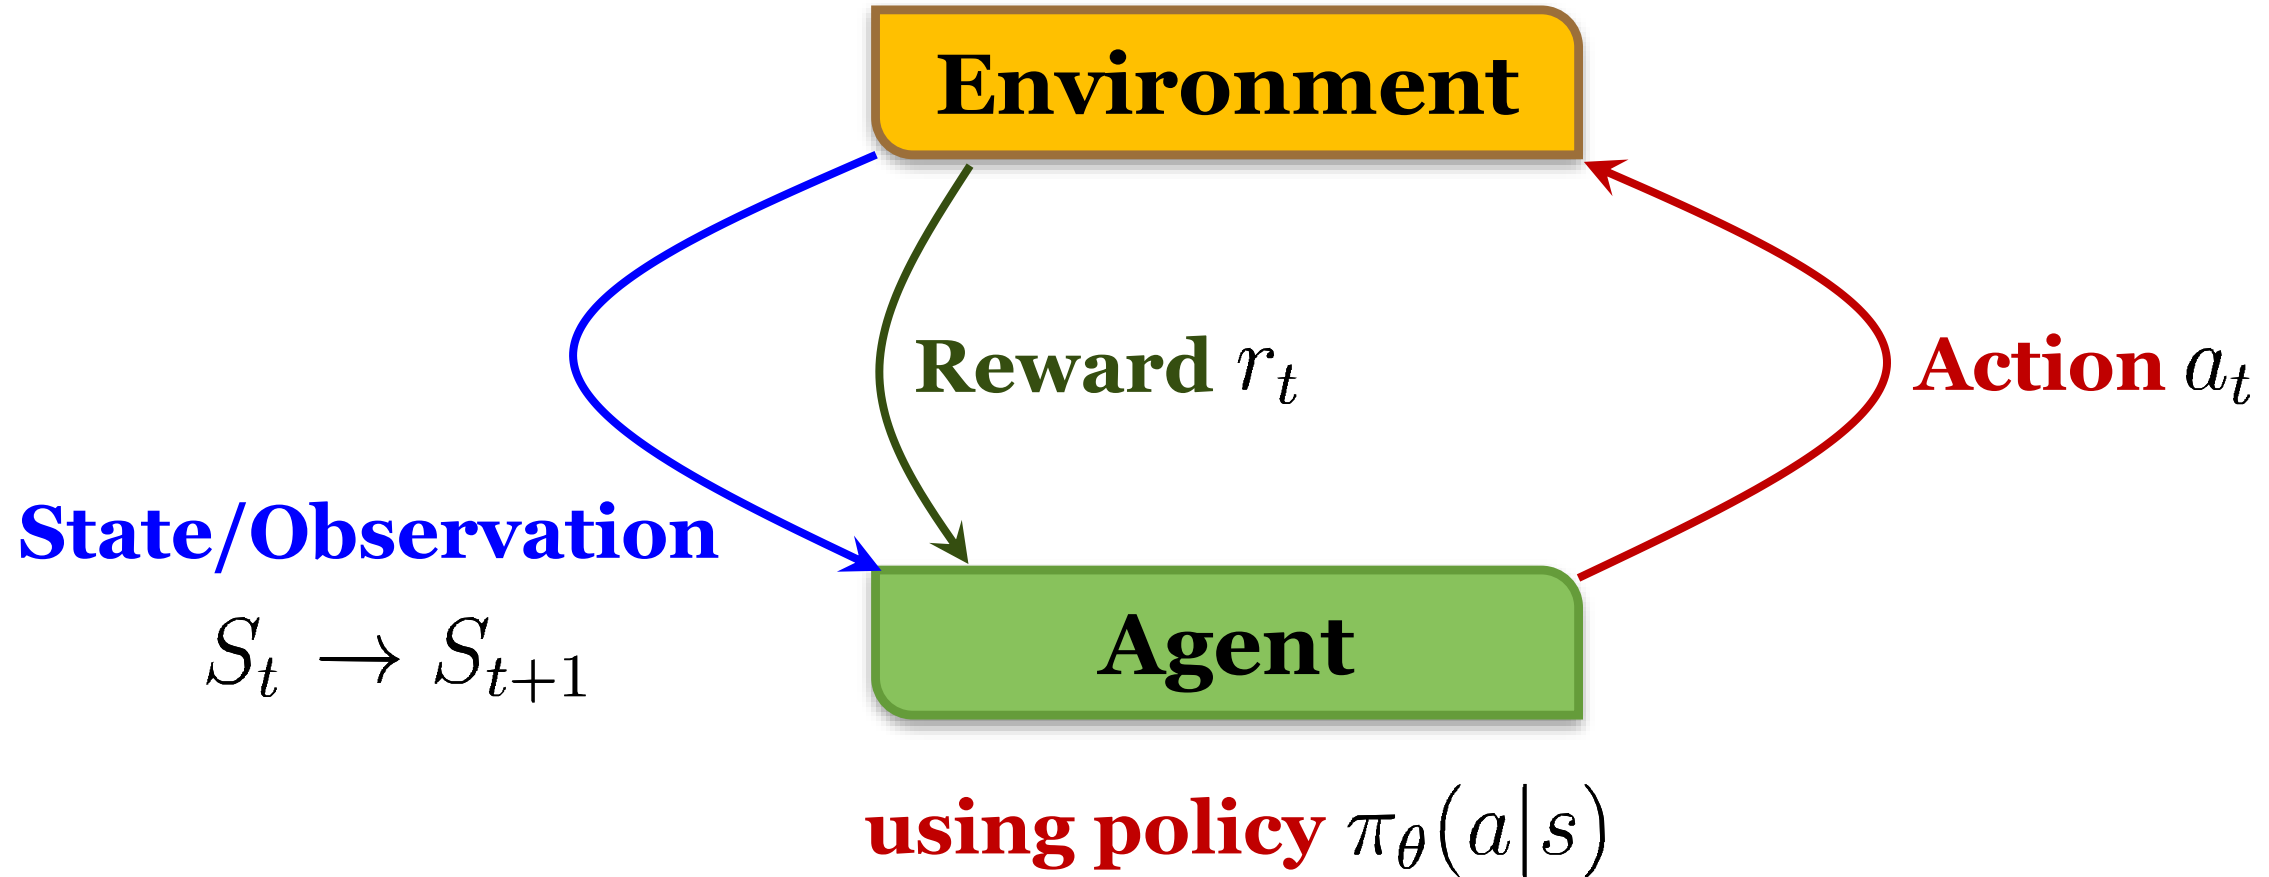
- state transition



- E.g., "up" action leads to a new state.
- State transition can be random.
 - ※ Randomness is from the environment.
 - ※ E.g., the Goombas' next move is random

$$p(s'|s, a) = P(S' = s' | S = s, A = a)$$

Terminology: agent environment interaction




Randomness in Reinforcement Learning

- **Actions** have randomness. $A \sim \pi(\cdot|s)$
 - ※ Given state s , the action can be random, e.g.,
$$\pi(\text{left}|s) = 0.2 \quad \pi(\text{right}|s) = 0.1 \quad \pi(\text{up}|s) = 0.7$$
- **State transitions** have randomness. $S' \sim (\cdot|s, a)$
 - ※ Given state $S = s$ and action $A = a$,
 - the environment randomly generates a new state S' .
- Play the game using AI.
 - ※ (**state**, **action**, **reward**) trajectory:
$$s_1, a_1, r_1, s_2, a_2, r_2, \dots, s_T, a_T, r_T.$$

Reinforcement Learning

- Learning to interact with an environment
 - ※ Robots, games, process control
 - ※ With limited human training
 - ※ Where the 'right thing' isn't obvious



- Reinforcement Learning:
 - ※ Goal: Maximize $\sum_{i=1}^{\infty} \text{Reward}(\text{State}_i, \text{Acton}_i)$
 - ※ Data: $\text{Reward}_i, \text{State}_{i+1} = \text{Interact}(\text{State}_i, \text{Acton}_i)$

12.1.2 Rewards and Returns

Return

- Definition: Return (aka **cumulative future reward**).
 - ※ $U_t = R_t + R_{t+1} + R_{t+2} + R_{t+3} + \dots$ until the game is over.
- Question: Are R and R_{t+1} equally important?
 - ※ Which of the followings do you prefer?
 - I give you \$100 right now.
 - I will give you \$100 one year later.
 - Or how about: I will give you \$200 one year later.
- Future reward is less valuable than present reward.
 - ※ R_{t+1} should be given less weight than R_t .

Return

- Definition: Return (aka **cumulative future reward**).
 - ※ $U_t = R_t + R_{t+1} + R_{t+2} + R_{t+3} + \dots$ until the game is over.
- Definition: Discounted return (aka **cumulative *discounted* future reward**).
 - ※ γ : *discount rate* -- tuning hyper-parameter.
 - ※ $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \dots$
- At time step t , the return U_t is random.
 - ※ Two sources of randomness:
 1. Action can be random: $P[A = a | S = s] = \pi(a | s)$
 2. New state can be random: $P[S' = s' | S = s, A = a] = p(s' | s, a)$

Randomness in Returns

- Definition: Discounted Return (at time step t).
 - ※ $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \dots$
- At time step t , the return U_t is random.
 1. Action can be random: $P[A = a | S = s] = \pi(a | s)$
 2. New state can be random: $P[S' = s' | S = s, A = a] = p(s' | s, a)$
- For any $i \geq t$, the reward R_i depends on S_i and A_i .
 - ※ Thus, given s_t , the return U_t depends on the random variables:
 - $A_t, A_{t+1}, A_{t+2}, \dots$ and S_{t+1}, S_{t+2}, \dots

12.1.3 Value functions

Action-Value Function $Q(s,a)$

- Definition: Discounted Return (cumulative discounted future reward).
 - ※ $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \dots$
- Definition: **Action-value function** for policy π
 - ※ U_t depends on actions $A_t, A_{t+1}, A_{t+2}, \dots$ and states S_{t+1}, S_{t+2}, \dots
 - ※ $Q_\pi(s_t, a_t) = \mathbb{E}[U_t | S_t = s_t, A_t = a_t]$
 - ※ Actions are random: $P[A = a | S = s] = \pi(a|s)$ (**Policy function**)
 - ※ States are random: $P[S' = s' | S = s, A = a] = p(s'|s, a)$ (**State transition**)

Action-Value Function $Q(s,a)$

- Definition: Discounted Return (cumulative discounted future reward).
$$\times U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \dots$$
- Definition: **Action-value function** for policy π
$$\times Q_\pi(s_t, a_t) = \mathbb{E}[U_t | S_t = s_t, A_t = a_t]$$
- Definition: **Optimal action-value function**
$$\times Q^*(s_t, a_t) = \max_{\pi} Q_\pi(s_t, a_t)$$

State-Value Function $V(s)$

- Definition: Discounted Return (cumulative discounted future reward).
 - ※ $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \dots$
- Definition: **Action-value function** for policy π
 - ※ $Q_\pi(s_t, a_t) = \mathbb{E}[U_t | S_t = s_t, A_t = a_t]$
- Definition: **State-value function**
 - ※ $V_\pi(s_t) = \mathbb{E}_A[Q_\pi(s_t, A)] \quad A \sim \pi(\cdot | s)$
 - ※ $V_\pi(s_t) = \mathbb{E}_A[Q_\pi(s_t, A)] = \sum_a \pi(a | s_t) \cdot Q_\pi(s_t, a) \quad (\text{if Actions are discrete})$
 - ※ $V_\pi(s_t) = \mathbb{E}_A[Q_\pi(s_t, A)] = \int \pi(a | s_t) \cdot Q_\pi(s_t, a) da \quad (\text{Actions are continuous})$

Understanding the Value Functions

- Action-value function: $Q_{\pi}(s_t, a_t) = \mathbb{E}[U_t | S_t = s_t, A_t = a_t]$
 - ※ For policy π , $Q_{\pi}(\mathbf{s}, \mathbf{a})$ evaluates how good it is for an agent to pick action \mathbf{a} while being in state \mathbf{s} .
- State-value function: $V_{\pi}(s_t) = \mathbb{E}_A[Q_{\pi}(s_t, A)]$
 - ※ For fixed policy π , $V_{\pi}(\mathbf{s})$ evaluates how good the situation is in state \mathbf{s} .
 - ※ $\mathbb{E}_{\mathbf{s}}[V_{\pi}(\mathbf{s})]$ evaluates how good the policy π is.


12.1.4 Play games using reinforcement learning

How does AI control the agent?


- Suppose we have a good policy $\pi(a|s)$.
 - ⌘ Upon observe the state s_t ,
 - ⌘ random sampling: $a_t \sim \pi(\cdot|s_t)$.
- Suppose we know the optimal action-value function $Q^*(s, a)$
 - ⌘ Upon observe the state s_t ,
 - ⌘ choose the action that maximizes the value: $a_t = \operatorname{argmax}_a Q^*(s_t, a)$.

OpenAI Gym

- Gym is a toolkit for developing reinforcement learning algorithms.
 - ※ <https://gym.openai.com/>
- Problem setting 1: Classical control problems




Cart Pole




Pendulum

OpenAI Gym


- Gym is a toolkit for developing reinforcement learning algorithms.
- Problem setting 2: Atari Games



Pong




Space Invader




Breakout

OpenAI Gym


- Gym is a toolkit for developing reinforcement learning algorithms.
- Problem setting 3: MuJoCo (Advanced Physics Simulation)



Ant



HalfCheetah




Humanoid

Play CartPole Game

```
import gym
env = gym.make('CartPole-v0') # 生成环境
state = env.reset()
for t in range(100):
    env.render() # 弹出环境渲染窗口
    print(state) # [0.01850658 0.01749877 -0.03132206 0.01806279]

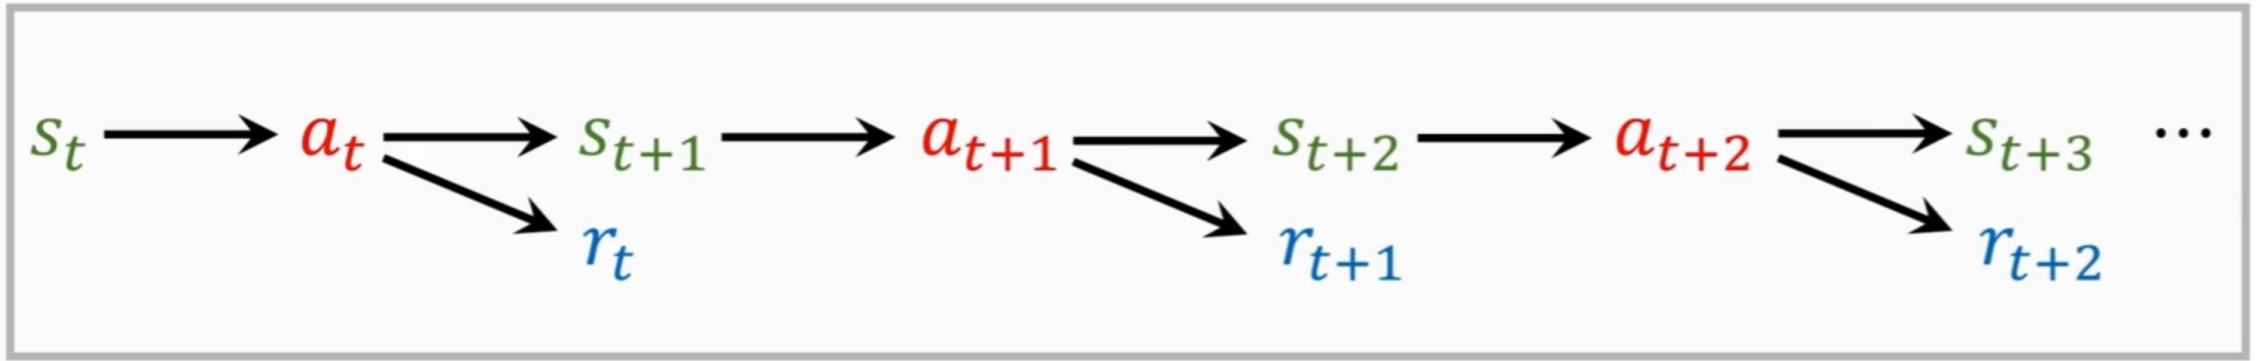
    action = env.action_space.sample() # take a random action
    state, reward, done, info = env.step(action)

    if done: # done == 1 means finished (win or loose)
        print('Mission terminated.')
        break
env.close()
```



Play games using reinforcement learning

- Observe state s_t , make action a_t , environment gives s_{t+1} and reward r_t .



- The agent can be controlled by either $\pi(a|s)$ or $Q^*(s, a)$.

12.2 Value-based Reinforcement learning

Recall: Discounted Return & Action-Value Function

- Definition: **Discounted Return** (cumulative discounted future reward).
 - ※ $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \dots$
- Definition: **Action-value function** for policy π
 - ※ U_t depends on actions $A_t, A_{t+1}, A_{t+2}, \dots$ and states S_{t+1}, S_{t+2}, \dots
 - ※ $Q_\pi(s_t, a_t) = \mathbb{E}[U_t | S_t = s_t, A_t = a_t]$ $Q^*(s_t, a_t) = \max_{\pi} Q_\pi(s_t, a_t)$
 - ※ Actions are random: $P[A = a | S = s] = \pi(a|s)$ (**Policy function**)
 - ※ States are random: $P[S' = s' | S = s, A = a] = p(s'|s, a)$ (**State transition**)

Recall: Action-Value Function $Q(s,a)$

- Definition: Discounted Return (cumulative discounted future reward).
 - ※ $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \dots$
- Definition: **Action-value function** for policy π
 - ※ $Q_\pi(s_t, a_t) = \mathbb{E}[U_t | S_t = s_t, A_t = a_t]$
- Definition: **Optimal action-value function**
 - ※ $Q^*(s_t, a_t) = \max_{\pi} Q_\pi(s_t, a_t)$
 - ※ Whatever policy function T is used, the result of taking a_t at state s_t cannot be better than $Q^*(s_t, a_t)$.

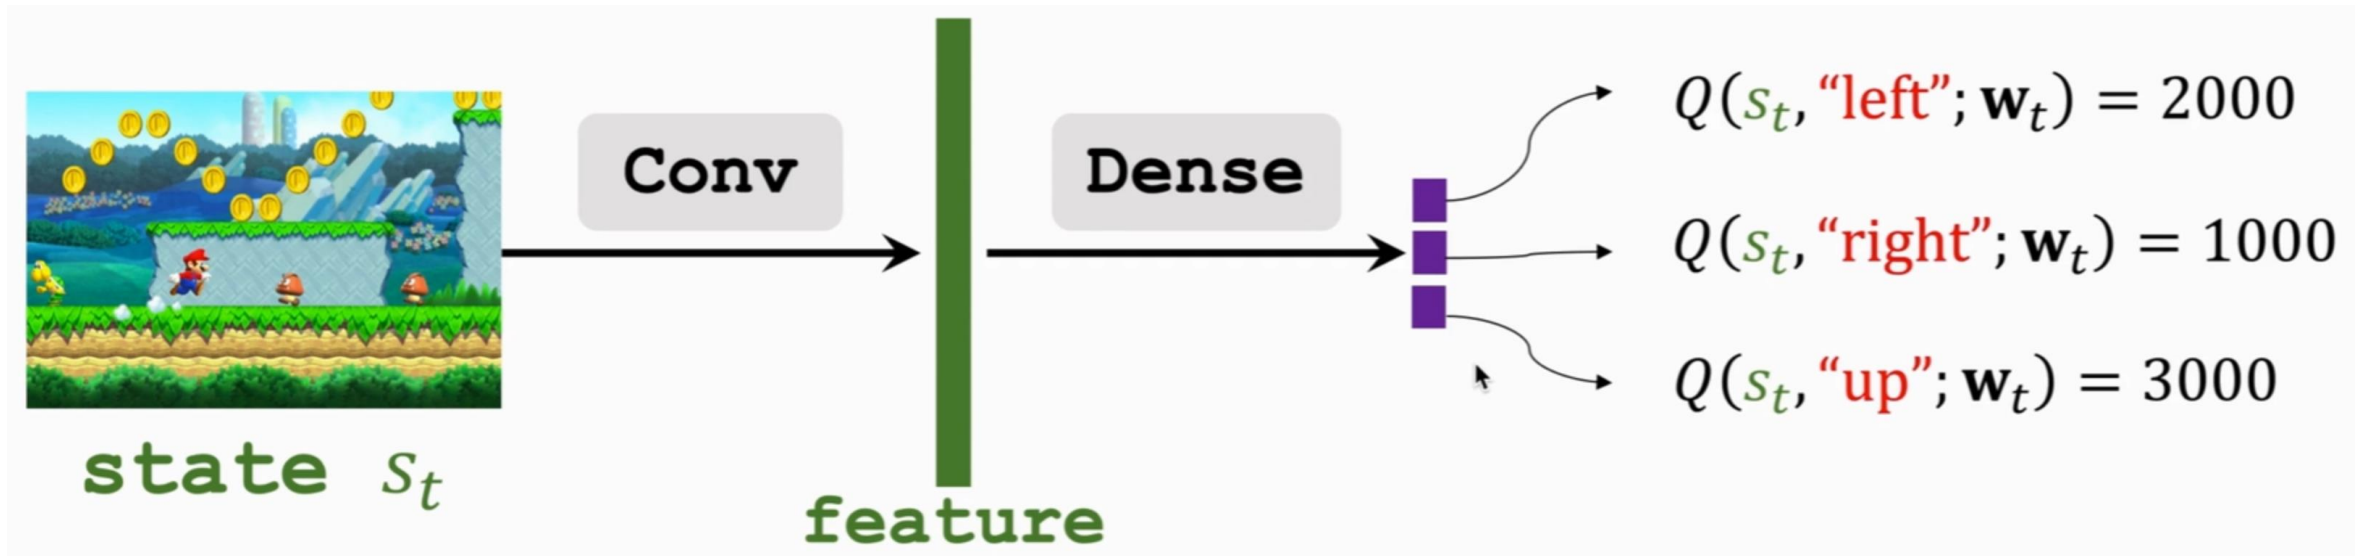
12.2.1 Deep Q Network (DQN)

Approximate the Q Function

- Goal: Win the game (\approx maximize the total reward.)
- Question: If we know $Q^*(s, a)$, what is the best action?
 - ⌘ Obviously, the best action is: $a^* = \underset{a}{\operatorname{argmax}} Q^*(s, a)$
- Q^* is an indication for how good it is
 - ⌘ for an agent to pick action a while being in state s .
- Challenge: We do not know $Q^*(s, a)$.
- Solution: Deep Q Network (DQN)
 - ⌘ Use neural network $Q(s, a; \mathbf{w})$ to approximate $Q^*(s, a)$.

Deep Q Network (DQN)

- Input shape: size of the screenshot.
- Output shape: dimension of action space.




- Question: Based on the predictions, what should be the action?

Apply DQN to Play Game

$$s_{t+2} \sim p(\cdot | s_{t+1}, a_{t+1})$$

$$s_{t+1} \sim p(\cdot | s_t, a_t)$$


$$s_{t+3} \sim p(\cdot | s_{t+2}, a_{t+2})$$



12.2.2 Temporal Different (TD)

Example

- Alice want to drive from NYC to Atlanta.
 - ※ Model $Q(\mathbf{w})$ estimates the time cost, e.g., 1000 minutes.
- Question: How do I update the model?
 - ※ Make a prediction: $q = Q(\mathbf{w})$, e.g., $q = 1000$.
 - ※ Finish the trip and get the target y , e.g., $y = 860$.
 - ※ Loss: $L = \frac{1}{2}(q - y)^2$
 - ※ Gradient: $\frac{\partial L}{\partial \mathbf{w}} = \frac{\partial q}{\partial \mathbf{w}} \cdot \frac{\partial L}{\partial q} = (q - y) \cdot \frac{\partial Q(\mathbf{w})}{\partial \mathbf{w}}$
 - ※ Gradient descent: $\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha \cdot \frac{\partial L}{\partial \mathbf{w}} \big|_{\mathbf{w}=\mathbf{w}_t}$



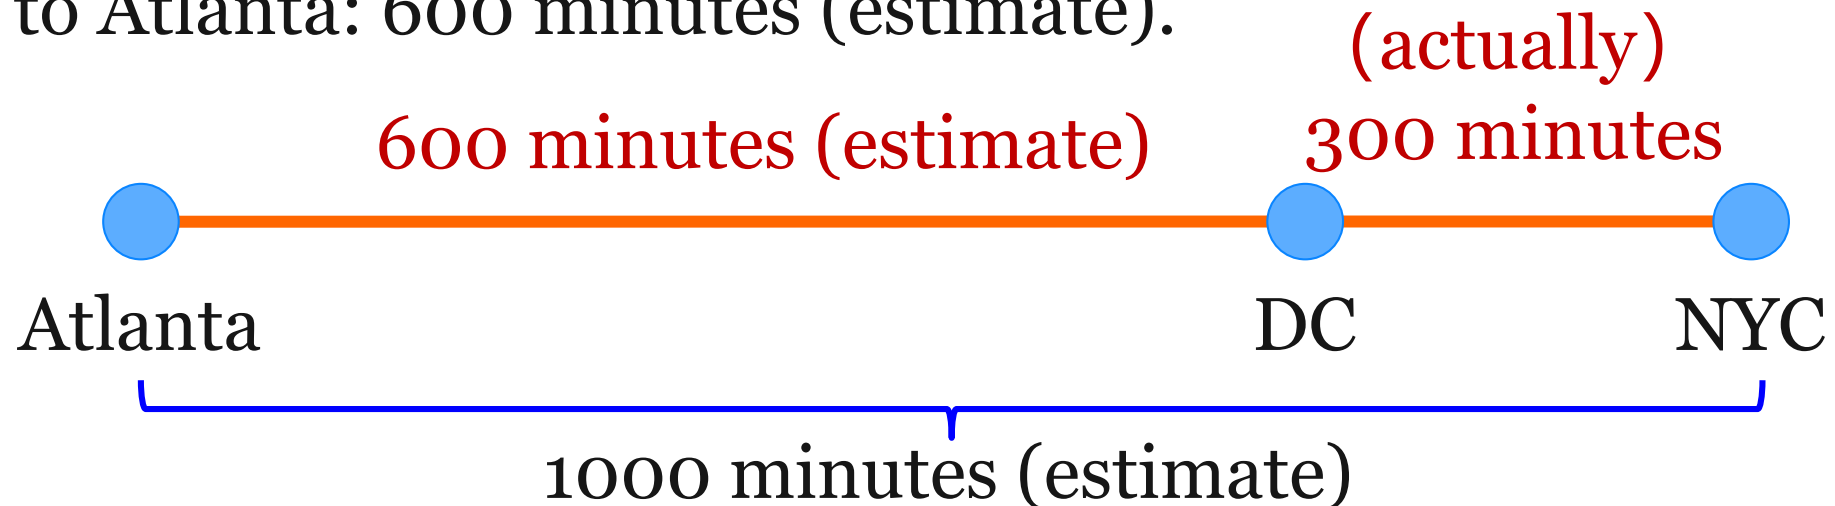
Example

- Alice want to drive from NYC to Atlanta.
 - ※ Model $Q(w)$ estimates the time cost, e.g., 1000 minutes.
- Question: How do I update the model?
 - ※ Can we update the model before finishing the trip?
 - ※ Can we get a better w as soon as we arrived DC?



Temporal Different (TD) Learning

- Model's estimate:
 - ※ NYC to Atlanta: 1000 minutes (estimate).
- Alice arrived at DC; actual time cost:
 - ※ NYC to DC: 300 minutes (actual).
- Model now updates its estimate:
 - ※ DC to Atlanta: 600 minutes (estimate).



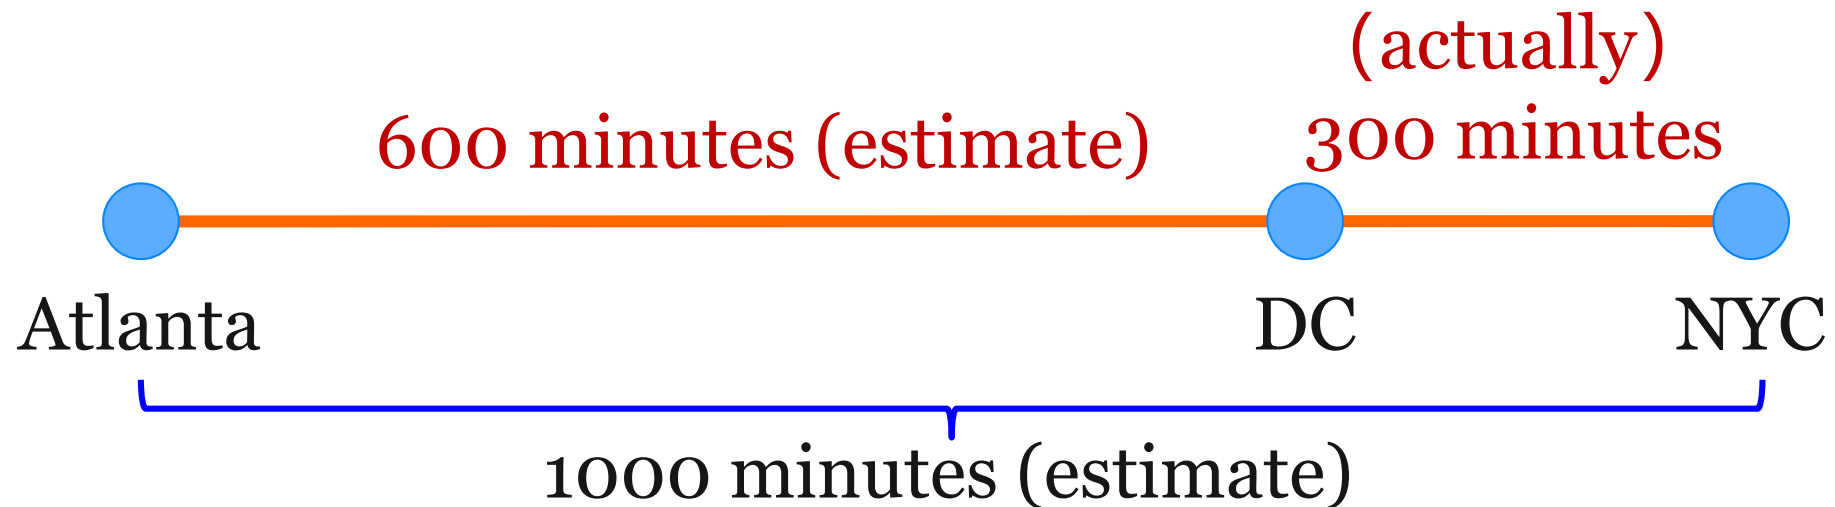
Temporal Different (TD) Learning

- Model's estimate: $Q(\mathbf{w}) = 1000$ minutes.. 900 minutes --> TD target
 - ※ Updated estimate: $300 + 600 = 900$ minutes.
- TD target $y = 900$ is a more reliable estimate than 1000.
 - ※ Loss: $L = \frac{1}{2}(Q(\mathbf{w}) - y)^2$
 - ※ Gradient: $\frac{\partial L}{\partial \mathbf{w}} = (1000 - 900) \cdot \frac{\partial Q(\mathbf{w})}{\partial \mathbf{w}}$
 - ※ Gradient descent: $\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha \cdot \frac{\partial L}{\partial \mathbf{w}} \big|_{\mathbf{w}=\mathbf{w}_t}$



Why does TD Learning Work

- Model's estimates:
 - ※ NYC to Atlanta: 1000 minutes.
 - ※ DC to Atlanta: 600 minutes.
 - ※ \rightarrow NYC to DC: 400 minutes.
- Ground truth:
 - ※ NYC to DC: 300 minutes.
 - ※ TD error: $\delta = 400 - 300 = 100$



How to apply TD learning to DQN?

- In the "driving time" example, we have the equation:

$$T_{\text{NYC} \rightarrow \text{ATL}} \approx T_{\text{NYC} \rightarrow \text{DC}} + T_{\text{DC} \rightarrow \text{ATL}}$$

$$\text{Model's estimate} = \text{Actual time} + \text{Model's estimate}$$

- In deep reinforcement learning:

$$Q(s_t, a_t; \mathbf{w}) \approx r_t + \gamma \cdot Q(s_{t+1}, a_{t+1}; \mathbf{w})$$

How to apply TD learning to DQN?

- Recall: Definition: **Discounted Return**

$$\begin{aligned}\times U_t &= R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \gamma^4 R_{t+4} + \dots \\ &= R_t + \gamma \cdot (R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 R_{t+4} + \dots) \\ &= R_t + \gamma \cdot U_{t+1}\end{aligned}$$

- TD learning for DQN:

$$\begin{aligned}\times \text{DQN's output, } Q(s_t, a_t; \mathbf{w}), \text{ is estimate of } E[U_t]. \\ \times \text{DQN's output, } Q(s_{t+1}, a_{t+1}; \mathbf{w}), \text{ is estimate of } E[U_{t+1}]. \\ \times \text{Thus, } \underbrace{Q(s_t, a_t; \mathbf{w})}_{\text{Prediction}} \approx \underbrace{r_t + \gamma \cdot Q(s_{t+1}, a_{t+1}; \mathbf{w})}_{\text{TD Target}}\end{aligned}$$

Train DQN using TD learning

- Prediction: $Q(s_t, a_t; \mathbf{w}_t)$
- TD target:

$$\begin{aligned} y_t &= r_t + \gamma \cdot Q(s_{t+1}, a_{t+1}; \mathbf{w}_t) \\ &= r_t + \gamma \cdot \max_a Q(s_{t+1}, a; \mathbf{w}_t) \end{aligned}$$

- Loss: $L_t = \frac{1}{2}[Q(s_t, a_t; \mathbf{w}) - y_t]^2$
- Gradient descent: $\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha \cdot \frac{\partial L_t}{\partial \mathbf{w}} \big|_{\mathbf{w}=\mathbf{w}_t}$

Summary: Value-Based Reinforcement Learning

- Definition: Optimal action-value function

$$Q^*(s_t, a_t) = \max_{\pi} \mathbb{E}[U_t | S_t = s_t, A_t = a_t]$$

- DQN: Approximate $Q^*(s, a)$ using a neural network (DQN).
 - ⊗ $Q(s, a; \mathbf{w})$ is a neural network parameterized by \mathbf{w} .
 - ⊗ Input: observed state \mathbf{s} .
 - ⊗ Output: scores for every action $a \in \mathcal{A}$

Temporal Difference (TD) Learning

- Algorithm: One iteration of TD learning.
 1. Observe state $S_t = s_t$ and action $A_t = a_t$.
 2. Predict the value: $q_t = Q(s_t, a_t; \mathbf{w}_t)$
 3. Differentiate the value network: $d_t = \frac{\partial Q(s_t, a_t; \mathbf{w})}{\partial \mathbf{w}} \Big|_{\mathbf{w}=\mathbf{w}_t}$
 4. Environment provides new state \mathbf{s}_{t+1} and reward \mathbf{r}_t .
 5. Compute TD target: $y_t = r_t + \gamma \cdot \max_a Q(s_{t+1}, a; \mathbf{w}_t)$
 6. Gradient descent: $\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha \cdot (q_t - y_t) \cdot d_t$

DeepMind's DQN playing Breakout

12.3 Policy-based Reinforcement learning

12.3.1 Policy network

Policy Function Approximation

Recall: Policy function $\pi(a|s)$

- Policy function $\pi(a|s)$ is a probability density function (PDF)
- It takes state s as input.
- It output the probabilities for all the actions, e.g.,

$$\pi(\text{left}|s) = 0.2$$


$$\pi(\text{right}|s) = 0.1$$

$$\pi(\text{up}|s) = 0.7$$

- The agent performs an action a random drawn from the distribution.

Policy Network $\pi(a|s, \theta)$

- Policy network: Use a neural net to approximate $\pi(a|s)$
 - ※ Use policy network $\pi(a|s, \theta)$ to approximate $\pi(a|s)$.
 - ※ θ : trainable parameters of the neural net.

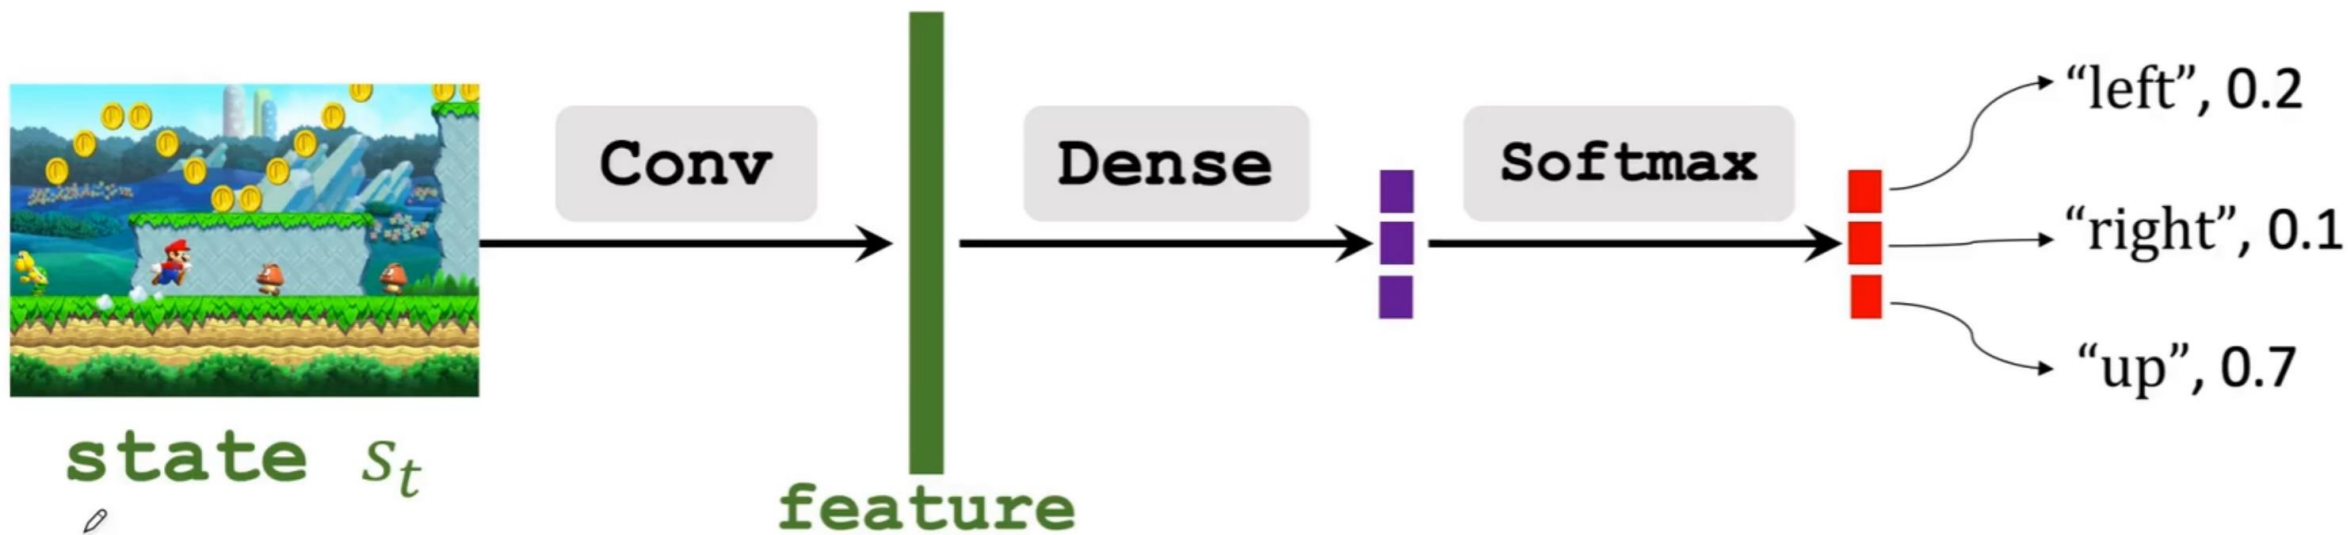


Policy Network $\pi(a|s, \theta)$

$$\sum_{a \in \mathcal{A}} \pi(a|s, \theta) = 1$$

- Where: $\mathcal{A} = \{\text{"left"}, \text{"right"}, \text{"up"}\}$ is the set of all actions.

※ That is why we use *softmax* activation



12.3.2 State-Value Function Approximation

Recall: State-Value Function $V(s)$

- Definition: Discounted Return (cumulative discounted future reward).
 - ※ $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \dots$
- Definition: **Action-value function** for policy π
 - ※ $Q_\pi(s_t, a_t) = \mathbb{E}[U_t | S_t = s_t, A_t = a_t]$
- Definition: **State-value function**
 - ※ $V_\pi(s_t) = \mathbb{E}_A[Q_\pi(s_t, A)] \quad A \sim \pi(\cdot | s)$
 - ※ $V_\pi(s_t) = \mathbb{E}_A[Q_\pi(s_t, A)] = \sum_a \pi(a | s_t) \cdot Q_\pi(s_t, a) \quad (\text{if Actions are discrete})$
 - ※ $V_\pi(s_t) = \mathbb{E}_A[Q_\pi(s_t, A)] = \int \pi(a | s_t) \cdot Q_\pi(s_t, a) da \quad (\text{Actions are continuous})$

Policy-Based Reinforcement Learning

- Definition: State-value function.

$$\text{※ } V_{\pi}(s_t) = \mathbb{E}_A[Q_{\pi}(s_t, A)] = \sum_a \pi(a|s_t) \cdot Q_{\pi}(s_t, a)$$

- Approximate state-value function.

※ Approximate policy function $\pi(a|s_t)$ by policy network $\pi(a|s_t, \boldsymbol{\theta})$

※ Approximate value function $V_{\pi}(s_t)$ by:

$$V(s_t; \boldsymbol{\theta}) = \sum_a \pi(a|s_t; \theta) \cdot Q_{\pi}(s_t, a)$$

Policy-Based Reinforcement Learning

- Definition: Approximate state-value function.

$$V(s; \theta) = \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a)$$

- Policy-based learning: Learn θ that maximizes $J(\theta) = \mathbb{E}_S[V(S; \theta)]$
- How to improve θ ? **Policy gradient ascent!**

※ observe state **s**.

※ Update policy by: $\theta \leftarrow \theta + \beta \cdot \frac{\partial V(s; \theta)}{\partial \theta}$

12.3.3 Policy gradient

Policy Gradient

- Definition: Approximate state-value function.

$$V(s; \theta) = \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a)$$

- Policy gradient: Derivative of $V(s; \theta)$ w.r.t. θ

$$\frac{\partial V(s; \theta)}{\partial \theta} = \frac{\partial \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a)}{\partial \theta}$$

$$= \sum_a \frac{\partial \pi(a|s; \theta) \cdot Q_\pi(s, a)}{\partial \theta} \quad \text{Push derivative inside the summation}$$

$$= \sum_a \frac{\partial \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a) \quad \text{Pretend } Q_\pi \text{ is independent of } \theta. \\ \text{(It may not be true.)}$$

Policy Gradient

- Definition: Approximate state-value function.

$$V(s; \theta) = \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a)$$

- Policy gradient: Derivative of $V(s; \theta)$ w.r.t. θ

$$\begin{aligned} \frac{\partial V(s; \theta)}{\partial \theta} &= \sum_a \frac{\partial \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a) \\ &= \sum_a \pi(a|s; \theta) \cdot \frac{\partial \log \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a) \end{aligned}$$

Chain rule: $\frac{\partial \log[\pi(\theta)]}{\partial \theta} = \frac{1}{\pi(\theta)} \cdot \frac{\partial \pi(\theta)}{\partial \theta}$

$$\pi(\theta) \cdot \frac{\partial \log[\pi(\theta)]}{\partial \theta} = \pi(\theta) \cdot \frac{1}{\pi(\theta)} \cdot \frac{\partial \pi(\theta)}{\partial \theta} = \frac{\partial \pi(\theta)}{\partial \theta}$$

Policy Gradient

- Definition: Approximate state-value function.

$$V(s; \theta) = \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a)$$

- Policy gradient: Derivative of $V(s; \theta)$ w.r.t. θ

$$\begin{aligned} \frac{\partial V(s; \theta)}{\partial \theta} &= \sum_a \frac{\partial \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a) \\ &= \sum_a \pi(a|s; \theta) \cdot \frac{\partial \log \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a) \\ &= \mathbb{E}_A \left[\frac{\partial \log[\pi(A|s; \theta)]}{\partial \theta} \cdot Q_\pi(s, A) \right] \end{aligned}$$

Note: This derivation is over-simplified and not rigorous.

Policy Gradient

- Two forms of policy gradient:

※ Form 1:
$$\frac{\partial V(s; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} = \sum_a \frac{\partial \pi(a|s; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \cdot Q_{\pi}(s, a)$$

※ Form 2:
$$\frac{\partial V(s; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} = \mathbb{E}_{A \sim \pi(\cdot|s; \boldsymbol{\theta})} \left[\frac{\partial \log \pi(A|s; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \cdot Q_{\pi}(s, A) \right]$$

Calculate Policy Gradient for Discrete Actions

- If the actions are discrete, e.g., action space $\mathcal{A} = \{\text{"left"}, \text{"right"}, \text{"up"}\}$
- Use Form 1:
$$\frac{\partial V(s; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} = \sum_a \frac{\partial \pi(a|s; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \cdot Q_{\pi}(s, a)$$
 1. Calculate $f(a, \theta) = \frac{\partial \pi(a|s; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \cdot Q_{\pi}(s, a)$, for every action $a \in \mathcal{A}$
 2. Policy gradient: $\frac{\partial V(s; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} = f(\text{"left"}, \boldsymbol{\theta}) + f(\text{"right"}, \boldsymbol{\theta}) + f(\text{"up"}, \boldsymbol{\theta})$
- This approach does not work for continuous actions.

Calculate Policy Gradient for Continuous Actions

- If the actions are continuous, e.g., action space $\mathcal{A} = [0, 1], \dots$
- Use Form 2:
$$\frac{\partial V(s; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} = \mathbb{E}_{A \sim \pi(\cdot | s; \boldsymbol{\theta})} \left[\frac{\partial \log \pi(A | s; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \cdot Q_{\pi}(s, A) \right]$$
 1. Randomly sample an action a according to the PDF $\pi(\cdot | s; \boldsymbol{\theta})$
 2. Calculate $g(\hat{a}, \boldsymbol{\theta}) = \frac{\partial \log \pi(\hat{a} | s; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \cdot Q_{\pi}(s, \hat{a})$
- Obviously, $\mathbb{E}_A[g(A, \boldsymbol{\theta})] = \frac{\partial V(s; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}}$
- $g(\hat{a}, \boldsymbol{\theta})$ is an unbiased estimate of $\frac{\partial V(s; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}}$

Calculate Policy Gradient for Continuous Actions

- If the actions are continuous, e.g., action space $\mathcal{A} = [0, 1], \dots$
- Use Form 2:
$$\frac{\partial V(s; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} = \mathbb{E}_{A \sim \pi(\cdot | s; \boldsymbol{\theta})} \left[\frac{\partial \log \pi(A | s; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \cdot Q_{\pi}(s, A) \right]$$
 1. Randomly sample an action a according to the PDF $\pi(\cdot | s; \boldsymbol{\theta})$
 2. Calculate $g(\hat{a}, \boldsymbol{\theta}) = \frac{\partial \log \pi(\hat{a} | s; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \cdot Q_{\pi}(s, \hat{a})$
 3. Use $g(\hat{a}, \boldsymbol{\theta})$ as an approximation to the policy gradient $\frac{\partial V(s; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}}$
- This approach work for discrete actions.

Update policy network using policy gradient

Algorithm

1. Observe the state s_t .
2. Randomly sample action a_t according to $\pi(\cdot|s_t; \boldsymbol{\theta}_t)$
3. Compute $q_t \approx Q_\pi(s_t, a_t)$ (some estimate).
4. Differentiate policy network: $d_{\boldsymbol{\theta}_t, t} = \frac{\partial \log \pi(a_t|s_t; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \Big|_{\boldsymbol{\theta}=\boldsymbol{\theta}_t}$
5. (Approximate) policy gradient: $g(a_t, \boldsymbol{\theta}_t) = q_t \cdot d_{\boldsymbol{\theta}_t, t}$
6. Update policy network: $\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \beta \cdot g(a_t, \boldsymbol{\theta}_t)$

Algorithm

3. Compute $q_t \approx Q_\pi(s_t, a_t)$ (some estimate). **How?**

- **Option 1: REINFORCE.**

- ※ Play the game to the end and generate the trajectory:

$$s_1, a_1, r_1, s_2, a_2, r_2, \dots, s_T, a_T, r_T$$

- ※ Compute the discounted return $u_t = \sum_{k=t}^T \gamma^{k-t} r_k$, for all t.

- ※ Since $Q_\pi(s_t, a_t) = \mathbb{E}[U_t]$, we can use u_t to approximate $Q_\pi(s_t, a_t)$

- ※ Therefore: Use $q_t = u_t$

Algorithm

3. Compute $q_t \approx Q_\pi(s_t, a_t)$ (some estimate). How?

- Option 2: Approximate Q_π using a neural network.

- ✂ This leads to the actor-critic method.

Policy-Based Method

- If a good policy function π is known, the agent can be controlled by the policy: randomly sample $a_t \sim \pi(\cdot | s_t)$.
- Approximate policy function $\pi(a|s)$ by policy network $\pi(a|s; \theta)$
- Learn the policy network by policy gradient.
- Policy gradient algorithm learn θ that maximizes $\mathbb{E}_S[V(S; \theta)]$

12.4 Actor-Critic Methods

12.4.1 Value Network and Policy Network

State-Value Function Approximation

- Definition: State-value function.

$$V_{\pi}(s) = \sum_a \pi(a|s) \cdot Q_{\pi}(s, a) \approx \sum_a \pi(a|s; \theta) \cdot q(s, a, \mathbf{w})$$

- Policy network (actor):

- ✧ Use neural net $\pi(a|s; \theta)$ to approximate $\pi(a|s)$

- ✧ θ : trainable parameters of the neural net.

- Value network (critic):

- ✧ Use neural net $q(s, a; \mathbf{w})$ to approximate $Q_{\pi}(s, a)$


- ✧ \mathbf{w} : trainable parameters of the neural net.

Actor: Policy Network $\pi(a|s, \theta)$

- Input: state s , e.g., a screenshot of Super Mario.
- Output: probability distribution over the actions.
- Let $\mathcal{A} = \{\text{"left"}, \text{"right"}, \text{"up"}\}$ be the set of all actions

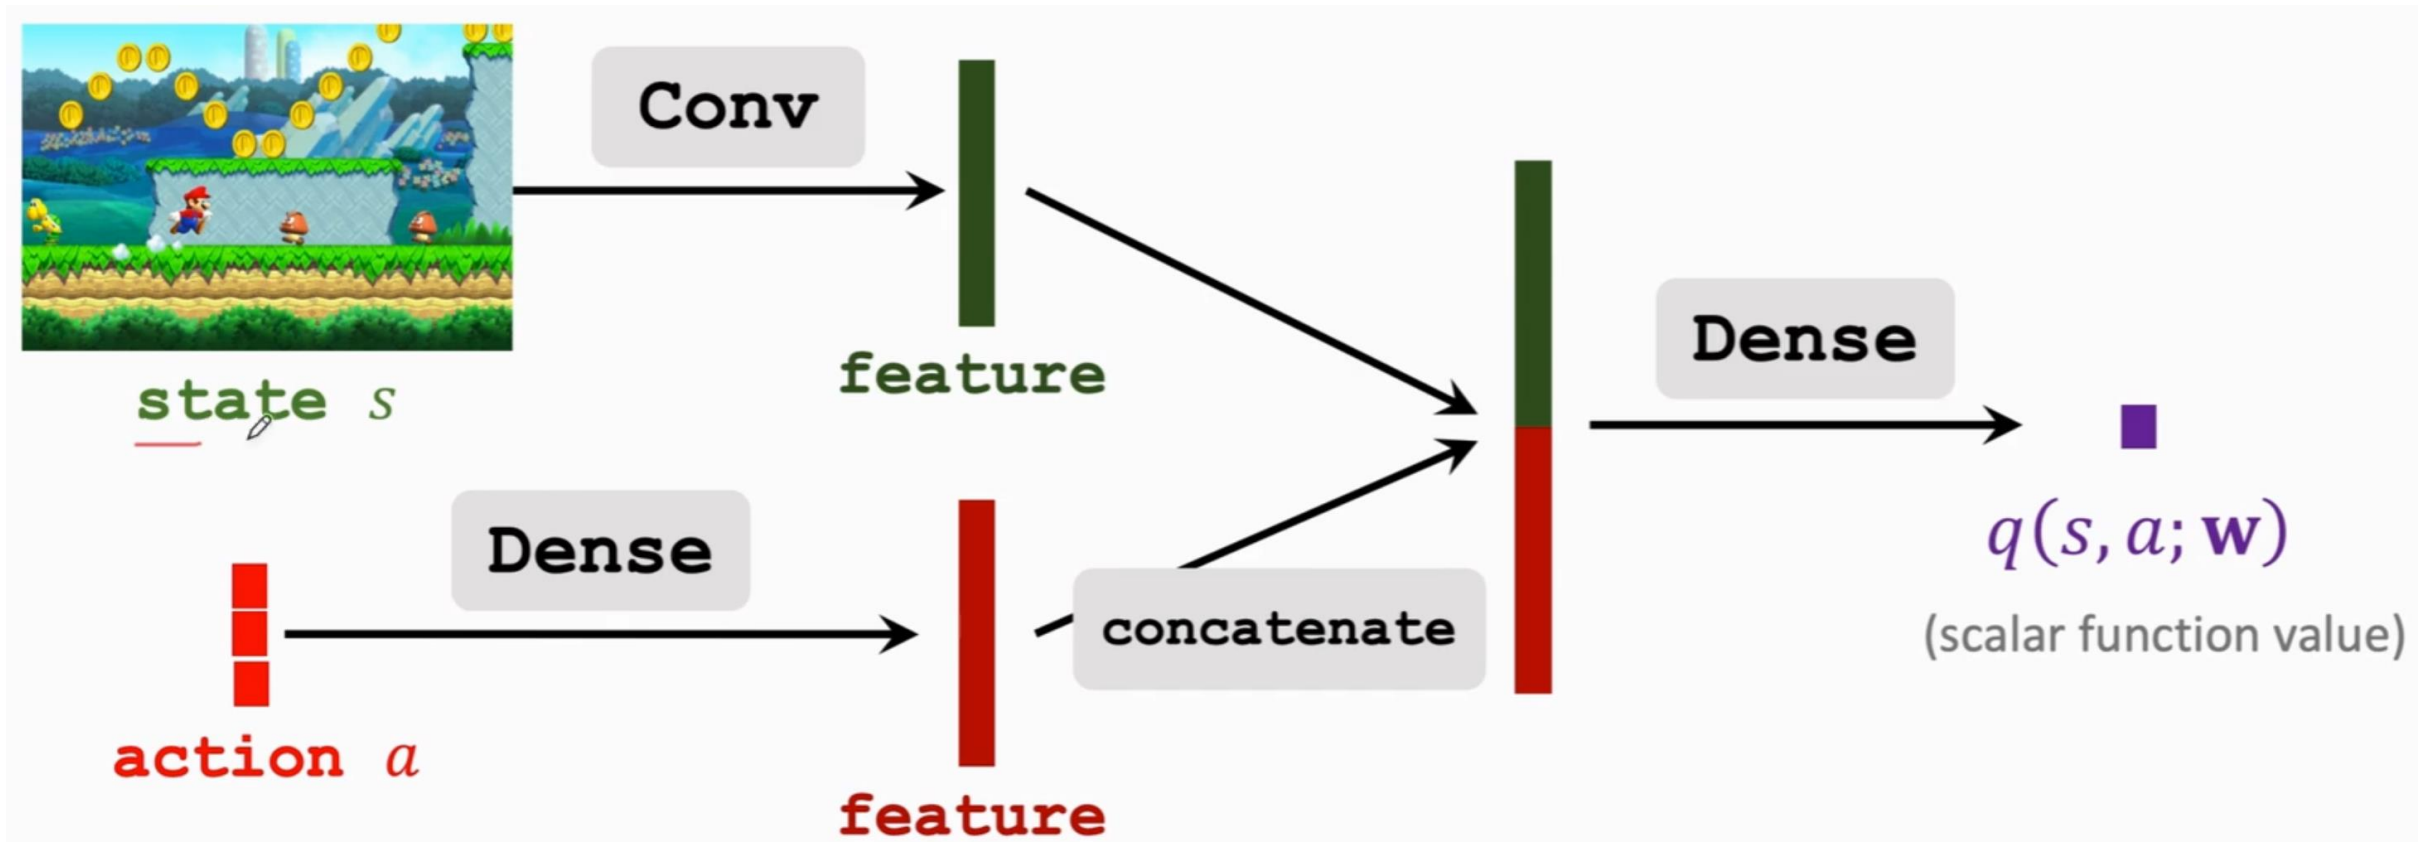
$$\sum_{a \in \mathcal{A}} \pi(a|s, \theta) = 1$$

That is why we use softmax activation



Critic: Value Network $q(s, a; \mathbf{w})$

- Inputs: state s and action a .
- Output: approximate action-value (scalar).



Train the networks

- Definition: State-value function approximated using neural networks.

$$V(s; \boldsymbol{\theta}, \mathbf{w}) = \sum_a \pi(a|s; \boldsymbol{\theta}) \cdot q(s, a; \mathbf{w})$$

- Training: Update the parameters $\boldsymbol{\theta}$ and \mathbf{w} .
 - ※ Update policy network $\pi(a|s; \boldsymbol{\theta})$ to increase the state-value $V(s; \boldsymbol{\theta}, \mathbf{w})$.
 - Actor gradually performs better.
 - Supervision is purely from the value network (critic).
 - ※ Update value network $q(s, a; \mathbf{w})$ to better estimate the return.
 - Critic's judgement becomes more accurate.
 - Supervision is purely from the rewards.

Train the networks

- Definition: State-value function approximated using neural networks.

$$V(s; \boldsymbol{\theta}, \mathbf{w}) = \sum_a \pi(a|s; \boldsymbol{\theta}) \cdot q(s, a; \mathbf{w})$$

- Training: Update the parameters $\boldsymbol{\theta}$ and \mathbf{w} .
 1. Observe the state s_t .
 2. Randomly sample action a_t according to $\pi(\cdot | s_t; \boldsymbol{\theta}_t)$.
 3. Perform a_t and observe new state s_{t+1} and reward r_t .
 4. Update \mathbf{w} (in value network) using temporal difference (TD).
 5. Update $\boldsymbol{\theta}$ (in policy network) using policy gradient.

Update value network q using TD

- Compute $q(s_t, a_t; \mathbf{w}_t)$ and $q(s_{t+1}, a_{t+1}; \mathbf{w}_t)$.
- TD target: $y_t = r_t + \gamma \cdot q(s_{t+1}, a_{t+1}; \mathbf{w}_t)$
- Loss: $L(\mathbf{w}) = \frac{1}{2} [q(s_t, a_t; \mathbf{w}) - y_t]^2$
- Gradient descent: $\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha \cdot \frac{\partial L(\mathbf{w})}{\partial \mathbf{w}} \Big|_{\mathbf{w}=\mathbf{w}_t}$

Update policy network π using policy gradient

- Definition: State-value function approximated using neural networks.

$$V(s; \boldsymbol{\theta}, \mathbf{w}) = \sum_a \pi(a|s; \boldsymbol{\theta}) \cdot q(s, a; \mathbf{w})$$

- Policy gradient: Derivative of $V(s_t; \boldsymbol{\theta}, \mathbf{w})$ w.r.t. $\boldsymbol{\theta}$.

- ※ Let $g(a, \boldsymbol{\theta}) = \frac{\partial \log \pi(a|s, \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \cdot q(s_t, a; \mathbf{w})$

- ※ $\frac{\partial V(s; \boldsymbol{\theta}, \mathbf{w})}{\partial \boldsymbol{\theta}} = \mathbb{E}_A[g(A, \boldsymbol{\theta})]$

- Algorithm: Update policy network using stochastic policy gradient.

- ※ Random sampling: $a \sim \pi(\cdot|s_t; \boldsymbol{\theta}_t)$ (Thus $g(a, \boldsymbol{\theta})$ is unbiased.)

- ※ Stochastic gradient ascent: $\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \beta \cdot g(a, \boldsymbol{\theta}_t)$

Summary of Algorithm

1. Observe state s_t and randomly sample $a_t \sim \pi(\cdot | s_t; \theta_t)$.
2. Perform a_t ; then environment gives new state s_{t+1} and reward r_t .
3. Randomly sample $\tilde{a}_{t+1} \sim \pi(\cdot | s_{t+1}; \theta_t)$. (Do not perform a_{t+1} !)
4. Evaluate value network: $q_t = q(s_t, a_t; \mathbf{w}_t)$ and $q_{t+1} = q(s_{t+1}, a_{t+1}; \mathbf{w}_t)$
5. Compute TD error: $\delta_t = q_t - (r_t + \gamma \cdot q_{t+1})$
6. Differentiate value network: $d_{\mathbf{w},t} = \frac{\partial q(s_t, a_t; \mathbf{w})}{\partial \mathbf{w}} \Big|_{\mathbf{w}=\mathbf{w}_t}$
7. Update value network: $\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha \cdot \delta_t \cdot d_{\mathbf{w},t}$
8. Differentiate policy network: $d_{\theta,t} = \frac{\partial \log \pi(a_t | s_t, \theta)}{\partial \theta} \Big|_{\theta=\theta_t}$
9. Update policy network: $\theta_{t+1} = \theta_t + \beta \cdot q_t \cdot d_{\theta,t}$

Summary of Algorithm

.....

5. Compute TD error: $\delta_t = q_t - (r_t + \gamma \cdot q_{t+1})$

.....

9. Update policy network: $\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \beta \cdot q_t \cdot d_{\boldsymbol{\theta},t}$

9. Update policy network: $\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \beta \cdot \delta_t \cdot d_{\boldsymbol{\theta},t}$

Summary : Policy Network and Value Network

- Definition: State-value function.

$$V_{\pi}(s) = \sum_a \pi(a|s) \cdot Q_{\pi}(s, a)$$

- Definition: function approximation using neural networks.
 - ※ Approximate policy function $\pi(a|s)$ by $\pi(a|s; \boldsymbol{\theta})$ (**actor**).
 - ※ Approximate value function $Q_{\pi}(s, a)$ by $q(s, a; \mathbf{w})$ (**critic**).

Roles of Actor and Critic

- During training
 - ⌘ Agent is controlled by policy network (actor): $a_t \sim \pi(\cdot \mid s_t; \boldsymbol{\theta})$.
 - ⌘ Value network q (critic) provides the actor with supervision.
- After training
 - ⌘ Agent is controlled by policy network (actor): $a_t \sim \pi(\cdot \mid s_t; \boldsymbol{\theta})$.
 - ⌘ Value network q (critic) will not be used.

Training

- Learning: Update the policy network (actor) by policy gradient.
 - ※ Seek to increase state-value: $V(s; \boldsymbol{\theta}, \mathbf{w}) = \sum_a \pi(a|s; \boldsymbol{\theta}) \cdot q(s, a; \mathbf{w})$
 - ※ Compute policy gradient: $\frac{\partial V(s; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} = \mathbb{E} \left[\frac{\partial \log \pi(A|s, \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \cdot q(s, A; \mathbf{w}) \right]$
 - ※ Perform gradient ascent.
- Learning: Update the value network (critic) by TD learning.
 - ※ Predicted action-value: $q_t = q(s_t, a_t; \mathbf{w})$
 - ※ TD target: $y_t = r_t + \gamma \cdot q(s_{t+1}, a_{t+1}; \mathbf{w}_t)$
 - ※ Gradient: $\frac{\partial (q_t - y_t)^2 / 2}{\partial \mathbf{w}} = (q_t - y_t) \cdot \frac{\partial q(s_t, a_t; \mathbf{w})}{\partial \mathbf{w}}$
 - ※ Perform gradient descent.

References

- Reinforcement Learning: An Introduction
 - ※ By Richard S. Sutton and Andrew G. Barto (代码)
 - ※ <http://incompleteideas.net/sutton/book/the-book-2nd.html>
- Statistical Reinforcement Learning: Modern Machine Learning Approaches
 - ※ By Masashi Sugiyama (链接)
- Deep Reinforcement Learning (open course)
 - ※ By Shusen Wang
 - ※ <https://www.youtube.com/watch?v=vmkRMvhCW5c>



End of the course, wish you have a pleasant journey!