

# Tony Hoare on Design Principles of ALGOL 60



In his Turing Award lecture

“The first principle was *security*: ... A consequence of this principle is that *every subscript was checked at run time against both the upper and the lower declared bounds of the array*. Many years later we asked our customers whether they wished us to provide an option to switch off these checks in the interests of efficiency. Unanimously, they urged us not to - they knew how frequently subscript errors occur on production runs where failure to detect them could be disastrous.

*I note with fear and horror that even in 1980, language designers and users have not learned this lesson. In any respectable branch of engineering, failure to observe such elementary precautions would have long been against the law.”*

[ C.A.R.Hoare, The Emperor’s Old Clothes, Communications of the ACM, 1980]

# Overview

1. How buffer overflows – and other memory related problems – work?
2. **How to spot such problems in C(++) code?**
3. **What can the compiler do about it?**
4. **What can we do about it?**

# Essence of the problem

Suppose in a C program we have an array of length 4

```
char buffer[4];
```

What happens if we execute the statement below ?

```
buffer[4] = 'a';
```

This is UNDEFINED! *ANYTHING* can happen !

If the data written (ie. the “a”) is user input that can be controlled by an attacker, this vulnerability can be exploited: *anything that the attacker* wants can happen.

Not only writing outside array bounds is dangerous, but so is reading (remember Heartbleed)

## Solution to this problem

- Check array bounds at runtime
  - Algol 60 proposed this back in 1960!
- Unfortunately, C and C++ have not adopted this solution, for efficiency reasons.  
(Perl, Python, Java, C#, and even Visual Basic have)
- As a result, buffer overflows have been the no 1 security problem in software ever since

## Other root causes

Apart from going outside array bounds there are other root causes

- malloc & free – ie. memory management by program(mer) rather than the language platform (using eg garbage collector)
- integer overflows
- format string attacks
- more generally: weak type systems, esp. for strings
- data races
- ....

# Problems caused by buffer overflows

- The first Internet worm, and all subsequent ones (CodeRed, Blaster, ...), exploited buffer overflows
- Buffer overflows cause in the order of 50% of all security alerts
  - Eg check out CERT, [cve.mitre.org](http://cve.mitre.org), or bugtraq
- Trends
  - Attacks are getting **cleverer**
    - defeating ever more clever countermeasures
  - Attacks are getting **easier** to do, by script kiddies

## Any C(++) code acting on **untrusted input** is at risk

- code taking input over **untrusted network**
  - eg. sendmail, web browser, wireless network driver,...
- code taking input from **untrusted user** on multi-user system,
  - esp. services running with high privileges (as ROOT on Unix/Linux, as SYSTEM on Windows)
- code acting on **untrusted files**
  - that have been downloaded or emailed
- also **embedded software** -
  - eg. in devices with (wireless) network connections such as mobile phones, RFID card, airplane navigation systems, ...

How does buffer overflow work?



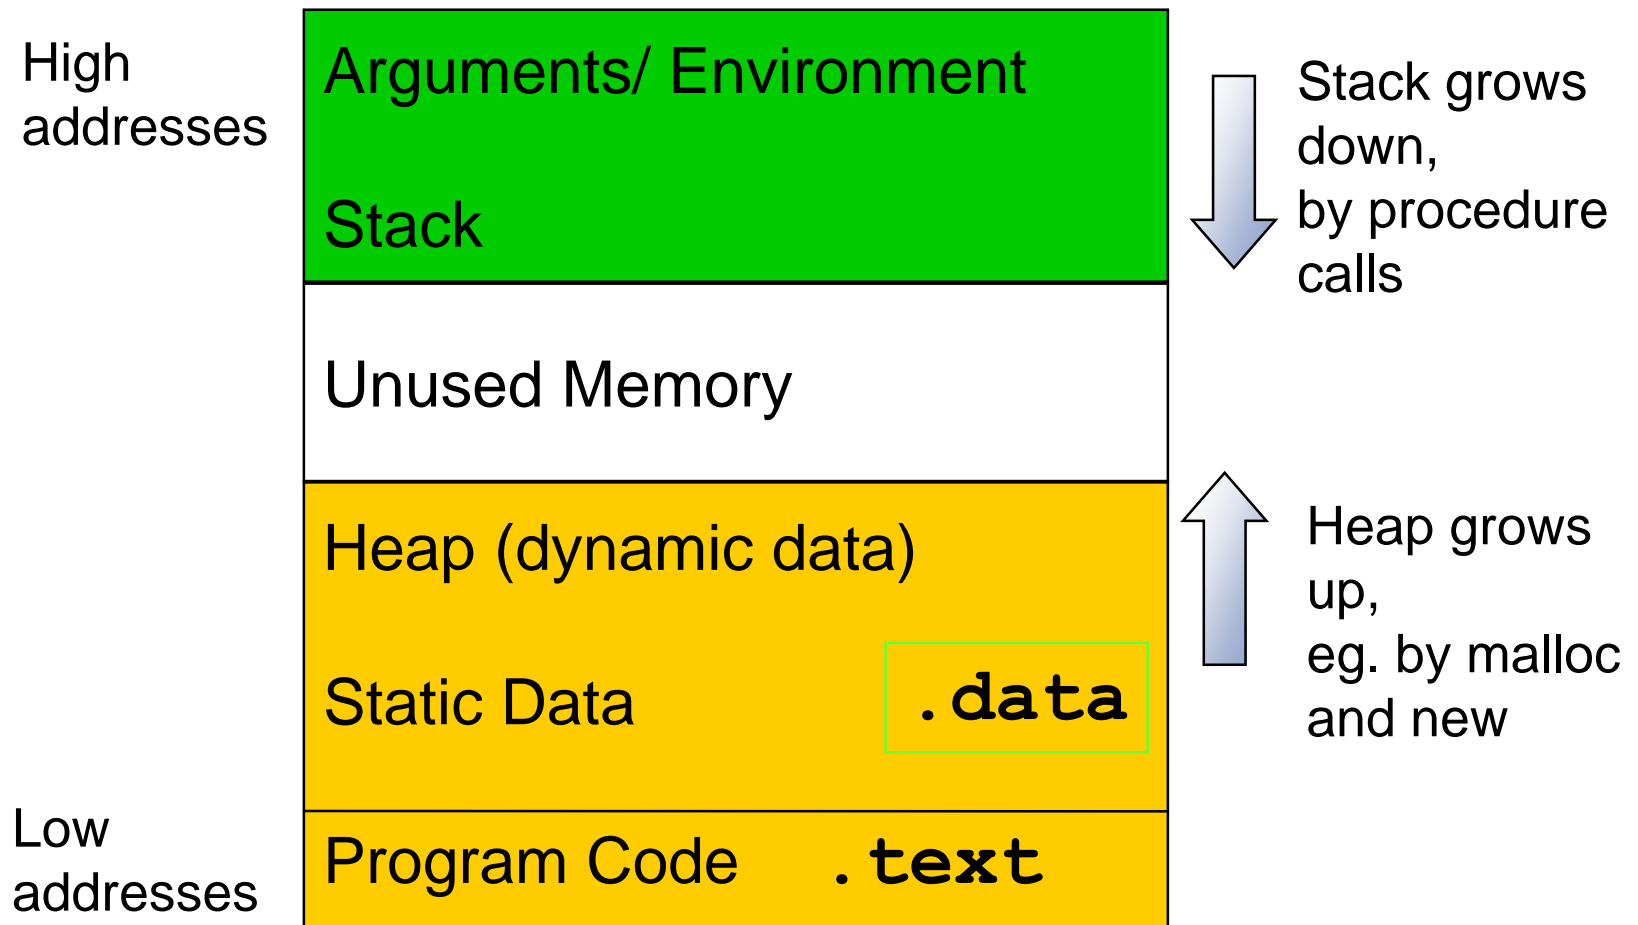
# Memory management in C/C++

- Program responsible for its memory management
- Memory management is very error-prone
  - *Who here has had a C(++) program crash with a segmentation fault?*

Technical term: C and C++ do not offer **memory-safety**  
(see lecture notes on language-based security, §3.1-3.2)

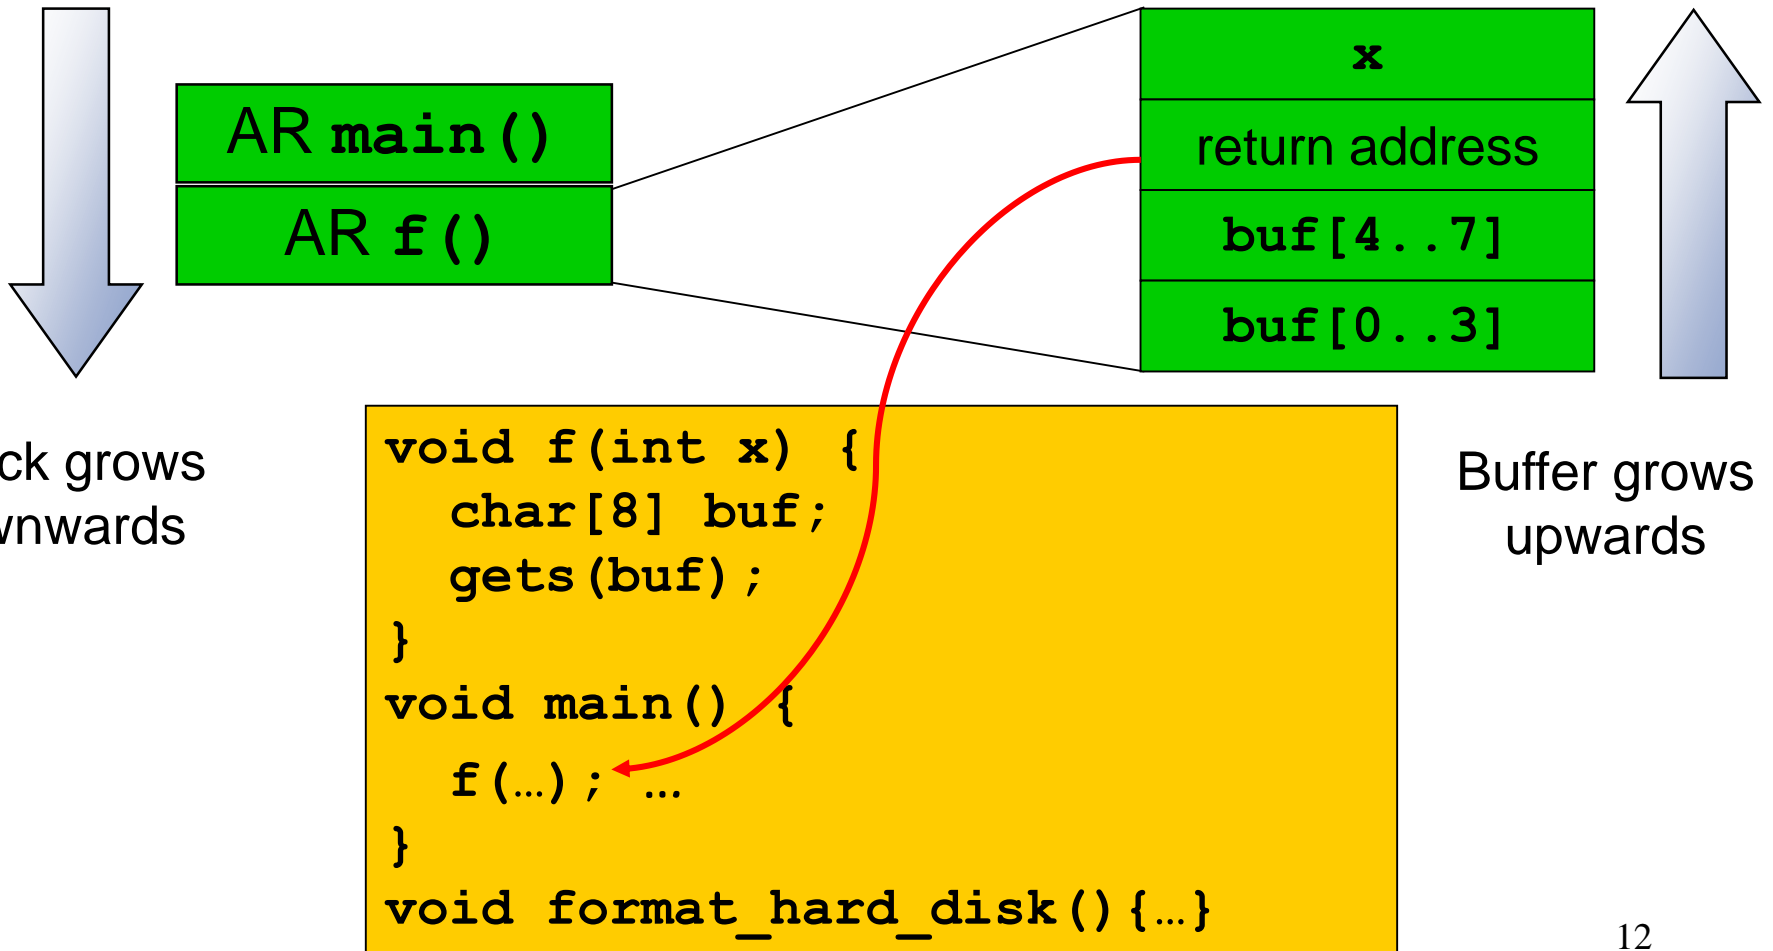
- Typical bugs:
  - Writing past the bound of an array
  - Pointer trouble
    - missing initialisation, bad pointer arithmetic, use after de-allocation (use after free), double de-allocation, failed allocation, forgotten de-allocation (memory leaks)...
- For efficiency, these bugs are not detected at run time:
  - behaviour of a buggy program is *undefined*

# Process memory layout



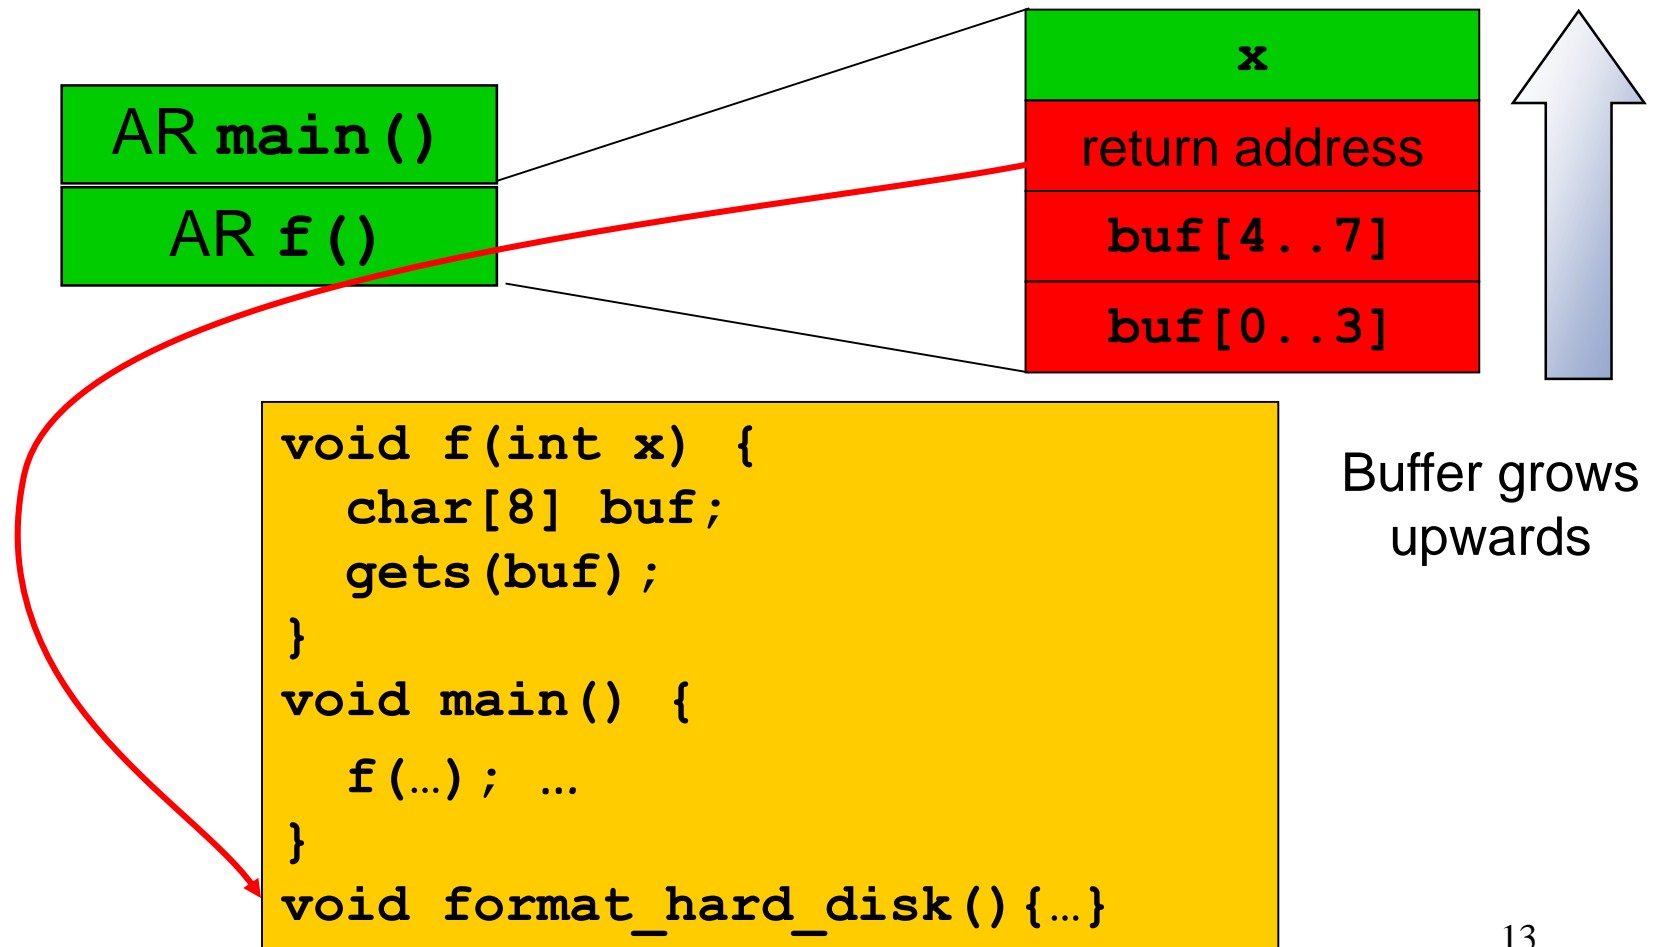
# Stack overflow

The stack consists of **Activation Records**:



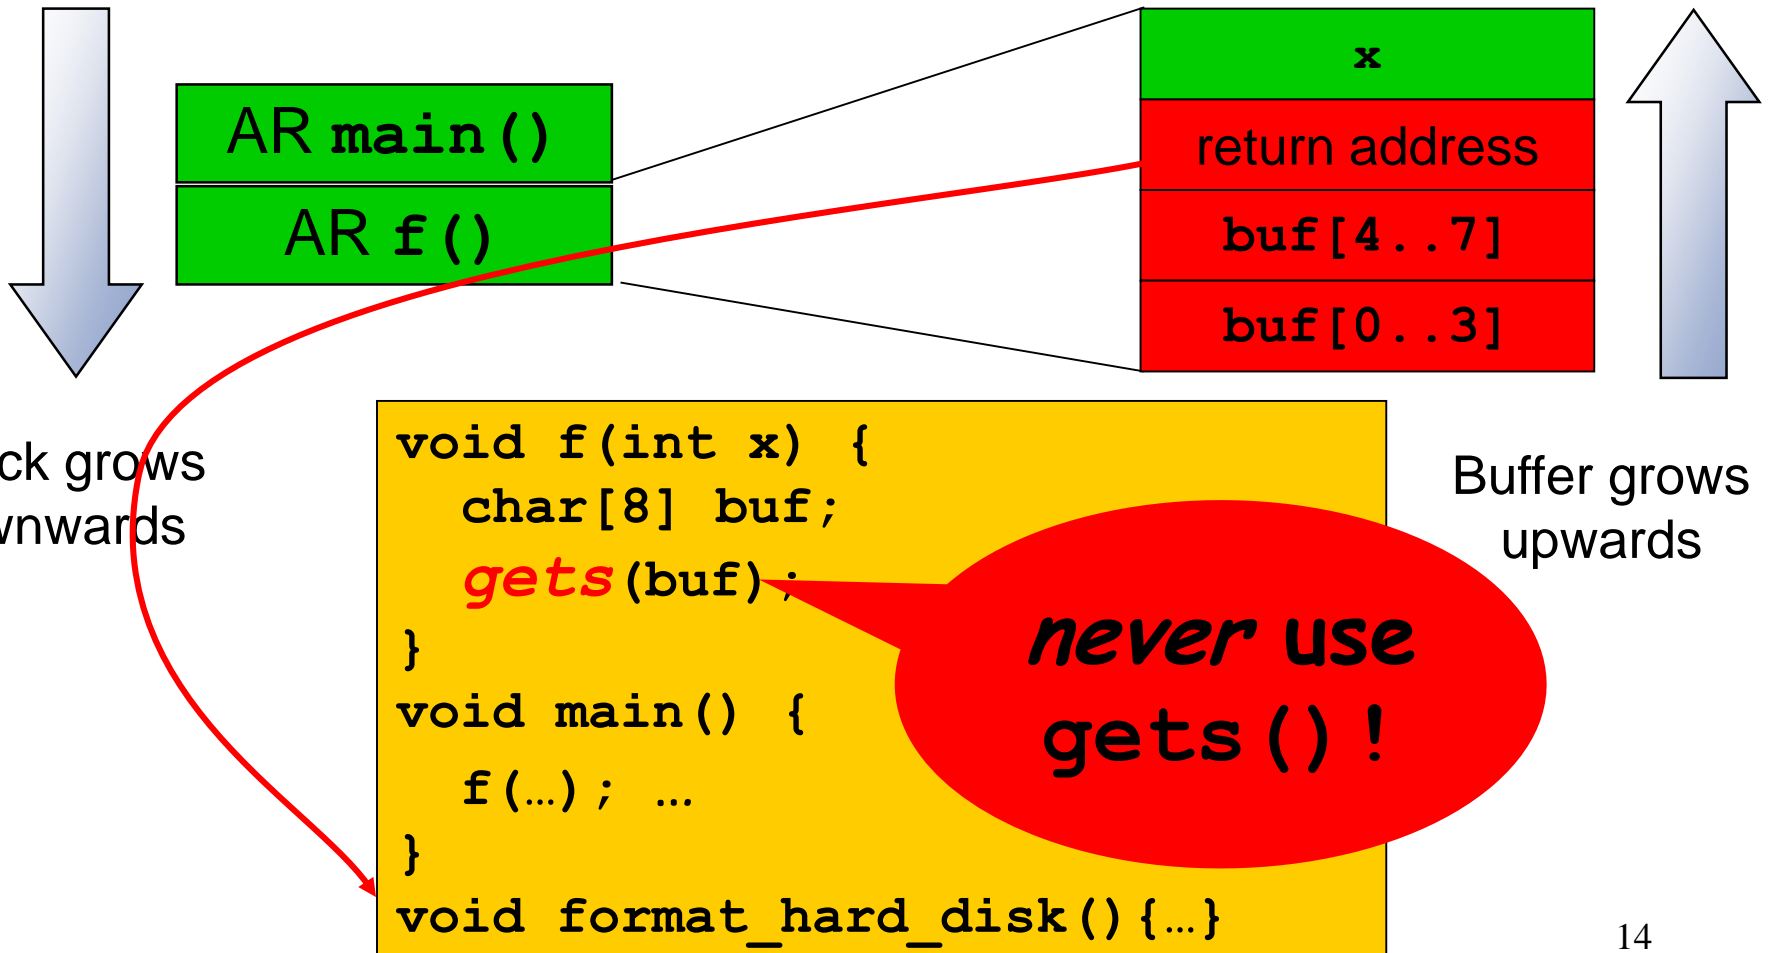
# Stack overflow

What if gets() reads **more than 8 bytes** ?



# Stack overflow

What if gets() reads **more than 8 bytes** ?



# Recurring problem: mixing control & data

In 1950s, Joe Engressia showed the telephone network could be hacked by **phone phreaking**:


ie. whistling at right frequencies

The root cause: **in-band signaling**



<http://www.youtube.com/watch?v=vVZm7l1CTBs>

In 1970s, before founding Apple together with Steve Jobs,  
Steve Wozniak sold Blue Boxes for phone phreaking at university



# Stack overflow, ie buffer overflow on the stack

- *How* the attacks works: overflowing buffers to corrupt data
- Lots of details to get right:
  - No nulls in (character-)strings
  - Filling in the correct return address:
    - Fake return address must be precisely positioned
    - Attacker might not know the address of his own string
  - Other overwritten data must not be used before return from function
  - ...
- Variant: **Heap overflow** of a buffer allocated on the heap instead of the stack

## What to attack? Fun with the stack

```
void f(char* buf1, char* buf2, bool b1) {  
    int i;  
    bool b2;  
    void (*fp) (int) ;  
    char[] buf3;  
    ....  
}
```

Overflowing stack-allocated buffer **buf3** to

- corrupt **return address**
  - ideally, let this return address point to another buffers where the attack code is placed
- corrupt **function pointers**, such as **fp**
- corrupt **any other data on the stack**, eg. **b2**, **i**, **a**, ...



Spotting the problem

## Example: gets

```
char buf[20];  
gets(buf); // read user input until  
           // first EoL or EOF character
```

- *Never* use `gets`
- Use `fgets(buf, size, stdin)` instead

## Example: strcpy

```
char dest[20];
```

```
strcpy(dest, src); // copies string src to dest
```

- `strcpy` assumes `dest` is long enough ,  
and assumes `src` is null-terminated
- Use `strncpy(dest, src, size)` instead

## Spot the defect! (1)

```
char buf[20];  
char prefix[] = "http://";  
...  
strcpy(buf, prefix);  
    // copies the string prefix to buf  
strncat(buf, path, sizeof(buf));  
    // concatenates path to the string buf
```

## Spot the defect! (1)

```
char buf[20];  
char prefix[] = "http://";  
...  
strcpy(buf, prefix);  
    // copies the string prefix to buf  
strncat(buf, path, sizeof(buf));  
    // concatenates path to the string buf
```

↑ strncat's 3rd parameter is number of  
chars to copy, not the buffer size

Another common mistake is giving `sizeof(path)` as 3rd argument...

## Spot the defect! (2)


```
char src[9];  
char dest[9];  
  
char* base_url = "www.ru.nl";  
strncpy(src, base_url, 9);  
    // copies base_url to src  
strcpy(dest, src);  
    // copies src to dest
```

## Spot the defect! (2)

```
char src[9];  
char dest[9];
```

`base_url` is 10 chars long, incl. its  
null terminator, so `src` will not be  
null-terminated

```
char* base_url = "www.ru.nl";  
strncpy(src, base_url, 9);  
    // copies base_url to src  
strcpy(dest, src);  
    // copies src to dest
```



## Spot the defect! (2)

```
char src[9];  
char dest[9];
```

`base_url` is 10 chars long, incl. its  
null terminator, so `src` will not be  
null-terminated

```
char* base_url = "www.ru.nl";  
strncpy(src, base_url, 9);  
    // copies base_url to src  
strcpy(dest, src);  
    // copies src to dest
```

so `strcpy` will overrun the buffer `dest`



## Example: `strcpy` and `strncpy`

- Don't replace

`strcpy(dest, src)`

with

`strncpy(dest, src, sizeof(dest))`

but with

`strncpy(dest, src, sizeof(dest)-1)`

`dst[sizeof(dest)-1] = `\\0`;`

if `dest` should be null-terminated!

- NB: a **strongly typed programming language** would guarantee that strings are always null-terminated...

## Spot the defect! (3)

```
char *buf;  
int i, len;  
  
read(fd, &len, sizeof(int));  
    // read sizeof(int) bytes, ie. an int,  
    // and store these at &len, ie. the  
    // memory address of the variable len  
buf = malloc(len);  
read(fd, buf, len); // read len bytes into buf
```

## Spot the defect! (3a)

```
char *buf;  
int i, len;
```

We forget to check for bytes  
representing a negative int,  
so `len` might be given a negative value

```
read(fd, &len, sizeof(int));  
    // read sizeof(int) bytes, ie. an int,  
    // and store these at &len, ie. the  
    // memory address of the variable len  
buf = malloc(len);  
read(fd, buf, len); // read len bytes into buf
```

`len` cast to unsigned and negative length overflows  
`read` then goes beyond the end of `buf`

## Spot the defect! (3b)

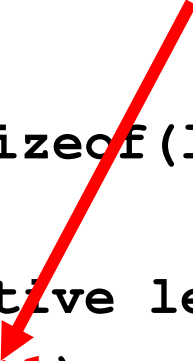
```
char *buf;  
int i, len;  
  
read(fd, &len, sizeof(len));  
if (len < 0)  
    {error ("negative length"); return; }  
buf = malloc(len);  
read(fd, buf, len);
```

Remaining problem may be that `buf` is not null-terminated

## Spot the defect! (3c)

```
char *buf;  
int i, len;  
  
read(fd, &len, sizeof(len));  
if (len < 0)  
    {error ("negative length"); return; }  
buf = malloc(len+1);  
read(fd, buf, len);  
buf[len] = '\0'; // null terminate buf
```

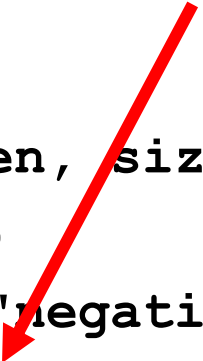
May result in **integer overflow**;  
we should check that  
**len+1** is positive



## Spot the defect! (3d)

```
char *buf;  
int i, len;  
  
read(fd, &len, sizeof(len));  
if (len < 0)  
    {error ("negative length"); return; }  
buf = malloc(len+1);  
read(fd, buf, len);  
buf[len] = '\\0'; // null terminate buf
```

What if the malloc() fails?  
(because we are out of memory)



## Spot the defect! (3e)

```
char *buf;
int i, len;

read(fd, &len, sizeof(len));
if (len < 0)
    {error ("negative length"); return; }
buf = malloc(len+1);
if (buf==NULL) { exit(EXIT_FAILURE);}
                // or something a bit more graceful
read(fd,buf,len);
buf[len] = '\0'; // null terminate buf
```

## Spot the defect! (3f) - Better still?

```
char *buf;
int i, len;

read(fd, &len, sizeof(len));
if (len < 0)
    {error ("negative length"); return; }
buf = calloc(len+1);
    //to initialise allocate memory to 0
if (buf==NULL) { exit(EXIT_FAILURE);}
    // or something a bit more graceful
read(fd,buf,len);
buf[len] = '\0'; // null terminate buf
```



## NB (Absence of) Language-Level Security

Note that in other (nicer?) programming languages, we might not have to worry about

- **writing past array bounds** (the language could throw an `IndexOutOfBoundsException` instead)
- **implicit conversions from signed to unsigned integers**
- **malloc possibly returning null** (the language could throw an `OutOfMemoryException` instead)
- **malloc not initialising memory**
- **integer overflow** (the language could throw an `IntegeroverflowException` instead)
- ...

## Spot the defect! (4)

```
#ifdef UNICODE
#define _sntprintf _snwprintf
#define TCHAR wchar_t
#else
#define _sntprintf _snprintf
#define TCHAR char
#endif

TCHAR buff[MAX_SIZE];
_sntprintf(buff, sizeof(buff), "%s\n", input);
```

[slide from presentation by Jon Pincus]

## Spot the defect! (4)

```
#ifdef UNICODE
#define _sntprintf _snwprintf
#define TCHAR wchar_t
#else
#define _sntprintf _snprintf
#define TCHAR char
#endif
```

`_sntprintf`'s 2<sup>nd</sup> param is # of chars in buffer, not # of bytes

```
TCHAR buff[MAX_SIZE];
_sntprintf(buff, sizeof(buff), "%s\n", input);
```



The CodeRed worm exploited such a mismatch, where code written under the assumption that 1 char was 1 byte allowed buffer overflows after the move from ASCII to Unicode

[slide from presentation by Jon Pincus]

## Spot the defect! (5)

```
#define MAX_BUF 256

void BadCode (char* input)
{   short len;
    char buf[MAX_BUF];

    len = strlen(input);

    if (len < MAX_BUF) strcpy(buf,input);
}
```

## Spot the defect! (5)

```
#define MAX_BUF 256
```

```
void BadCode (char* input)
```

```
{  short len;
```

```
  char buf[MAX_BUF];
```

```
  len = strlen(input);
```

```
  if (len < MAX_BUF) strcpy(buf, input);
```


```
}
```

What if **input** is longer than 32K ?

len will be a negative number,  
due to **integer overflow**



hence: potential  
**buffer overflow**




The **integer overflow** is the root problem, but the (heap) **buffer overflow** that this enables make it exploitable

## Spot the defect! (6)

```
bool CopyStructs(InputFile* f, long count)
{
    structs = new Structs[count];
    for (long i = 0; i < count; i++)
        { if !(ReadFromFile(f,&structs[i])) )
            break;
        }
}
```

## Spot the defect! (6)

```
bool CopyStructs(InputFile* f, long count)
{
    structs = new Structs[count];
    for (long i = 0; i < count; i++)
        { if !(ReadFromFile(f, &structs[i])) )
            break;
        }
}
```



effectively does a  
`malloc(count*sizeof(type))`  
which may cause integer overflow

And this integer overflow can lead to a (heap) buffer overflow.

Since 2005 the Visual Studio C++ compiler adds check to prevent this

## Spot the defect! (7)

```
char buff1[MAX_SIZE], buff2[MAX_SIZE];
// make sure url a valid URL and fits in buff1 and
  buff2:
if (! isValid(url)) return;
if (strlen(url) > MAX_SIZE - 1) return;
// copy url up to first separator, ie. first '/', to
  buff1
out = buff1;
do {
  // skip spaces
  if (*url != ' ') *out++ = *url;
} while (*url++ != '/');
strcpy(buff2, buff1);
...
```

[slide from presentation by Jon Pincus]



## Spot the defect! (7)

### Loop termination (exploited by Blaster)

```
char buff1[MAX_SIZE], buff2[MAX_SIZE];  
// make sure url a valid URL and fits in buff1 and  
buff2:  
if (! isValid(url)) return;  
if (strlen(url) > MAX_SIZE - 1) return;  
// copy url up to first separator, ie. first '/', to  
buff1  
out = buff1;  
do {  
    // skip spaces  
    if (*url != ' ') *out++ = *url;  
} while (*url++ != '/');  
strcpy(buff2, buff1),  
...
```

length up to the first null

what if there is no '/' in the URL?

[slide from presentation by Jon Pincus]

## Spot the defect! (8)

```
#include <stdio.h>

int main(int argc, char* argv[])
{   if (argc > 1)
        printf(argv[1]);
    return 0;
}
```

This program is vulnerable to **format string attacks**, where calling the program with strings containing special characters can result in a buffer overflow attack.

# Format string attacks

New type of attack, invented discovered in 2000.

Another way of corrupting the stack.

- Strings can contain special characters, eg `%s` in  

```
printf("Cannot find file %s", filename);
```

  
Such strings are called **format strings**
- What happens if we execute the code below?  

```
printf("Cannot find file %s");
```
- What *may* happen if we execute  

```
printf(string)
```

  
where `string` is user-supplied ?  
Esp. if it contains special characters, eg `%s, %x, %n, %hn`?

# Format string attacks

- `%x` reads and prints some bytes from stack

# This can leak sensitive data

Eg

```
printf(s);
```

will dump the stack if attacker supplies as input **s** the string

%x %x %x %x %x %x %x %x %x %x %x %x %x %x %x %x %x %x  
 %x %x %x %x %x %x %x %x %x %x %x %x %x %x %x %x %x %x  
 %x %x %x %x %x %x %x %x %x %x %x %x %x %x %x %x %x %x  
 %x

- `%n` writes the number of characters printed so far onto the stack

This allows the attacker to corrupt the stack...

# Advanced format string attacks

An attacker can try to control which address is used for reading from memory using `%s` or for writing to memory using `%n` with specially crafted format strings

- `\xEF\xCD\xCD\xAB %x %x ... %x %s`

With the right number of `%x` characters, this will print the string located at address **ABCDCEFF**

- `\xEF\xCD\xCD\xAB %x %x ... %x %n`

With the right number of `%x` characters, this will write the number of characters printed so far to location **ABCDCEFF**

The tricky things are inserting the right number of `%x`, and choosing an interesting address

Format string attacks should be history...

Format string attacks are easy to spot & fix:

replace `printf(str)`

with `printf("%s", str)`

# Recap: buffer overflows

- buffer overflow is **#1 weakness** in C and C++ programs
  - because these language are not **memory-safe**
- **tricky to spot**
- typical cause: poor programming with **arrays** and **strings**
  - esp. **library functions for null-terminated strings**
- related attacks
  - **format string attack**: another way of corrupting stack
  - **integer overflows**: a stepping stone on the way to get buffers to overflows
- Other low level flaws can also cause memory problems, eg **data races**


## Runtime aka dynamic countermeasures

These are things that the compiler can do,  
without the programmer needing to know!



# stack canaries

- introduced in [StackGuard](#) in gcc
- a dummy value - [stack canary or cookie](#) - is written on the stack in front of the return address and checked when function returns
- a careless stack overflow will overwrite the canary, which can then be detected.



# stack canaries

Stack without canary

<b>x</b>
return address
buf[4..7]
buf[0..3]

Stack with canary

<b>x</b>
return address
<i>canary value</i>
buf[4..7]
buf[0..3]

# stack canaries

- introduced in **StackGuard** in gcc
- a dummy value - **stack canary or cookie** - is written on the stack in front of the return address and checked when function returns
- a careless stack overflow will overwrite the canary, which can then be detected
- a careful attacker can overwrite the canary with the correct value
- additional countermeasures:
  - use a random value for the canary
  - XOR this random value with the return address
  - include string termination characters in the canary value

# Further improvements of stack canaries

- **PointGuard**
  - also protects other data values, eg function pointers, with canaries
- **ProPolice's Stack Smashing Protection (SSP)** by IBM
  - also **re-orders stack elements** to reduce potential for trouble: swapping parameters x and y on the stack changes whether overrunning x can corrupt y  
this is especially dangerous if y is a function pointer
- **Stackshield** has a special stack for return addresses, and can disallow function pointers to the data segment

# Non-eXecutable memory (NX aka $W\oplus X$ )

Distinguish

- X: executable memory (for storing code)
  - W: writeable, non-executable memory (for storing data)
- and let processor refuses to execute non-executable code

This can be done for the stack, or for arbitrary memory pages

How does this help?

Attacker can no longer jump to his own attack code,  
as any input he provides as attack code will be non-executable

Aka DEP (Data Execution Prevention).

Intel calls it eXecute-Disable (XD) , AMD calls it Enhanced Virus Protection

# Return-to-libc attacks

Way to get around non-executable memory:

overflow the stack to jump to code that is already there,  
esp. library code in **libc**

instead of jumping to your own attack code.

**libc** is a rich library that offers many possibilities for attacker, eg.  
**system, exec, fork**

which provide the attacker with any functionality he wants..

# Return-oriented programming

Next stage in evolution of attacks, esp. on 64 bit x86, as people removed or protected dangerous library calls like `system()`

Instead of using entire library call, the attacker

- looks for **small snippets of code which end with a return** (so-called **gadgets**)  
`... ; ... ; ... ; ret`
- strings these gadgets together as subroutines to form a program that does what he wants

Most libraries contain enough gadgets to provide a **Turing complete** programming language.

Return-to-libc is essentially simple case of return oriented programming.

# Control Flow Integrity (CFI)

A return-to-libc attack can be spotted as an unusual call pattern

Eg a function `foo()` never calls library routine `bar()` - and `bar()` does not even occur in the code of `foo()` - but when supplied with malicious input `foo()` suddenly does call `bar()`

With some extra administrative overhead for function calls, such strange call patterns can be detected and stopped at runtime.



# Address Space Layout Randomisation (ASLR)

- Attacker needs detailed information on memory layout (eg to jump to specific piece of code)
- By randomising the layout every time we start a program
  - ie. moving the offset of the heap, stack, etc, by some random valuethe attacker's life becomes much harder
- When to randomise? Only when we start a new program, or should we re-randomise when we fork() a process?

# Dynamic countermeasures (recap)

- canaries
- non-executable memory ( $W\oplus X$ )
- address space layout randomisation (ASLR)
- Control Flow Integrity

None of these countermeasures are perfect!

A determined attacker can and will find a way around them.

eg by figuring out canary values, figuring out the offset used in address randomisation, returning to libc, etc.

Moreover, they do not protect against **heap overflows**

Other countermeasures

# Countermeasures

- We can take countermeasures at different points in time
  - before we even begin programming
  - during development
  - at compilation time
  - when testing
  - when executing code

to prevent, mitigate, or detect buffer overflows problems

# Generic defence mechanisms

- Reducing attack surface

Not running or even installing certain software, or enabling all features by default, mitigates the threat

- Mitigating impact by reducing permissions

Reducing OS permissions of software (or user) will restrict the damage that an attack can have


- principle of least privilege

# Prevention

- Don't use C or C++
  - you can write insecure code in any programming language, but some make it easier...
- Better programmer awareness & training

Read – and make other people read – books like

- Building Secure Software, J. Viega & G. McGraw, 2002
- Writing Secure Code, M. Howard & D. LeBlanc, 2002
- 24 deadly sins of software security, M. Howard, D LeBlanc & J. Viega, 2005
- Secure programming for Linux and UNIX HOWTO, D. Wheeler,
- Secure C coding, T. Sirainen
- CERT secure coding guidelines for C and C++  
[www.securecoding.cert.org](http://www.securecoding.cert.org)



# Dangerous C system calls

source: Building secure software, J. Viega & G. McGraw, 2002

## Extreme risk

- `gets`

## High risk

- `strcpy`
- `strcat`
- `sprintf`
- `scanf`
- `sscanf`
- `fscanf`
- `vfscanf`
- `vsscanf`
- `streadd`
- `strecpy`
- `strtrns`
- `realpath`
- `syslog`
- `getenv`
- `getopt`
- `getopt_long`
- `getpass`

## Moderate risk

- `getchar`
- `fgetc`
- `getc`
- `read`
- `bcopy`

## Low risk

- `fgets`
- `memcpy`
- `snprintf`
- `strccpy`
- `strcadd`
- `strncpy`
- `strncat`
- `vsnprintf`

## Prevention – use better string libraries

- there is a choice between using **statically** vs **dynamically** allocated buffers
  - **static** approach easy to get wrong, and chopping user input may still have unwanted effects
  - **dynamic** approach susceptible to out-of-memory errors, and need for failing safely



## Better string libraries (1)

- `libsafe.h` provides safer, modified versions of eg. `strcpy`
  - prevents buffer overruns beyond current stack frame in the dangerous functions it redefines
- `libverify` enhancement of `libsafe`
  - keeps copies of the stack return address on the heap, and checks if these match
- `strncpy(dst,src,size)` and `strncat(dst,src,size)`  
with `size` the size of `dst`, not the maximum length copied.  
Consistently used in OpenBSD

## Better string libraries (2)

- [glib.h](#) provides Gstring type for dynamically growing null-terminated strings in C
  - but failure to allocate will result in crash that cannot be intercepted, which may not be acceptable
- [Strsafe.h](#) by Microsoft guarantees null-termination and always takes destination size as argument
- [C++ string class](#)
  - but `data()` and `c-str()` return low level C strings, ie `char*`, with result of `data()` is not always null-terminated on all platforms...

# Runtime detection on instrumented binaries

There are many memory error detection tools that **instrument binaries** to allow runtime detection of memory errors, esp.

- **out-of-bounds access**
- **use-after-free** bugs on heap

with some overhead (time, memory space) but no false positives

For example **Valgrind (Memcheck), Dr. Memory, Purify, Insure++, BoundsChecker, Cachegrind, Intel Parallel Inspector, Discoverer, AddressSanitizer,...**

# Safer variants of C

Some approaches go further and propose safer dialects of C which include

- bound checks,
- type checks
- automated memory management, to ensure memory safety
  - by garbage collection or region-based memory management

Examples are Cyclone, CCured, Vault, Control-C, Fail-Safe C, ...

# Fuzzing aka fuzz testing

Testing for security can be difficult!

- How to hit the right cases?

A classic technique to find buffer weaknesses is **fuzz testing**

- send *random, very long* inputs, to an application
- if there are buffer overflow weaknesses, this is likely to crash the application with a segmentation fault

The nice thing is that this is easy to automate!

*More on fuzz testing in the security testing lecture.*

# Code review & static analysis

- **Code reviews**  
ie. someone reviewing the code manually  
Expensive & labour intensive
- **Code scanning tools** aka **static analysis**  
Automated tools that look for suspicious patterns in code;  
ranges for **CTRL-F** or **grep**, to advanced analyses

Incl. free tools

- **RATS** – also for PHP, Python, Perl
- **Flawfinder**, **ITS4**, **Deputy**, **Splint**
- **PREfix**, **PREfast** by Microsoft

plus other commercial tools

**Coverity**, **PolySpace**, **Klockwork**, ...

# (formal) verification

The most extreme form of static analysis:

- Program verification

proving by mathematical means (eg Hoare logic) that memory management of a program is safe

- extremely labour-intensive ☹
- eg hypervisor verification project by Microsoft & Verisoft:
  - <http://www.microsoft.com/emic/verisoft.msp>

*Beware: in industry “verification” means testing,  
in academia it means formal program verification*

# Conclusions



“C is a terse and unforgiving abstraction of silicon”



# Summary

- Buffer overflows are a top security vulnerability
- Any C(++) code acting on untrusted input is at risk  
so effectively  
Any C(++) code is at risk
- Getting rid of buffer overflow weaknesses in C(++) code is hard and may prove to be impossible
  - Ongoing arms race between countermeasures and ever more clever attacks.
  - Attacks are not only getting cleverer, using them is getting easier

# Common themes

Buffer overflows involve three more general problems:

1. lack of input validation
2. mixing data & code  
namely data and return address on the stack
3. believing in & relying on an **abstraction** that is not 100%  
guaranteed & enforced  
namely **types** and **procedure interfaces** in C  
`int f(float f, boolean b, char* buf);`