


Software-Based Fault Isolation

Jinseong Jeon

software-based fault isolation



- Software-based fault isolation is the act of separating something faulty.

Efficient Software-Based Fault Isolation


R. Vahbe, S. Lucco, T. E. Anderson, and S. L. Graham

SOSP '93

fault isolation?

- need to incorporate independently developed software modules
 - micro-kernel design
 - BSD network packet filter
 - application-specific virtual memory management
 - Active Messages
 - extensible software
 - MS object linking and embedding system
 - Quark Xpress desktop publishing system
 - high I/O processes
 - POSTGRES
- need to prevent faults in extension code from corrupting other codes or permanent data while cooperating
- Hence, fault isolation is an act of separating distrusted extensions.


hardware-based fault isolation



- place each software module in its own address space
- communicate through Remote Procedure Call (RPC)
 - trap into the OS kernel,
 - copying each argument from the caller to the callee,
 - saving and restoring registers,
 - switching hardware address space,
 - trap back to user level.

expensive

software-based fault isolation



- load extension codes and their data into their own *fault domain*
 - fault domain = code segment + data segment
- enforce security policies that
 - a distrusted module is prohibited from writing or jumping outside its fault domain.
 - i.e. those distrusted modules cannot modify/execute each other's data/code.
 - the only way to do is to use explicit cross fault-domain communication.

possible questions

- how to enforce such security policies?
 - by binary rewriting
- what to rewrite, and how?
 - unsafe instruction
 - that cannot be statically verified to be within the correct segment
 - use dedicated registers
 - segment matching
 - address sandboxing
- how to share process resources and data?
 - trusted arbitration code
 - virtual address aliasing (or, shared segment matching)
- how to communicate with other fault domains?
 - explicit cross-fault-domain RPC interface
 - stub and jump table

segment matching

- fault domain

- = code segment + data segment
- shares a unique pattern of upper bits, “segment identifier”

- insert checking code before every unsafe instruction

- indirect jumps or stores, i.e. via registers of which value is determined at runtime

- pseudo code

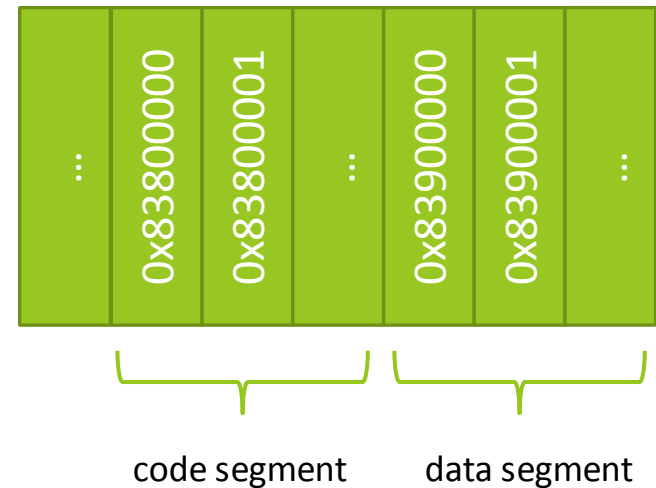
dedicated-reg \leftarrow target address

scratch-reg \leftarrow (dedicated-reg \gg shift-reg)

compare scratch-reg and segment-reg


trap if not equal

store/jump using dedicated-reg



address sandboxing

- instead of checking,
just setting the upper bits
to the correct segment identifier
- in the section “Ensure, don’t check”
at the next paper,
 - check = segment matching
 - ensure = address sandboxing




- pseudo code

$\text{dedicated-reg}^{x2} \leftarrow \text{target-reg} \ \& \ \text{and-mask-reg}$

$\text{dedicated-reg} \leftarrow \text{dedicated-reg} \mid \text{segment-reg}^{x2}$

store/jump using dedicated-reg

cross fault domain communication



- trusted stubs to handle RPC
 - for each pair of fault domains
 - stub: copy arguments, re/store registers, switch the exe. stack, validate dedicated regs but! no traps or address space switching (thus, cheaper than HW RPC)
- jump tables to transfer control
 - consists of jump instructions of which target address is legal, outside the domain