

Directed Greybox Fuzzing



Marcel Böhme



Thuan Pham



M.-D. Nguyen **Abhik Roychoudhury**



NUS
National University
of Singapore

Motivation

- Automated vulnerability detection techniques
 1. **Blackbox Fuzzing** (no program analysis, no feedback)
 2. **Whitebox Fuzzing** (mostly program analysis)
 3. **Greybox Fuzzing** (no program analysis, but coverage feedback)

Motivation

- Automated vulnerability detection techniques
 1. **Blackbox Fuzzing** (no program analysis, no feedback)

Seed Input



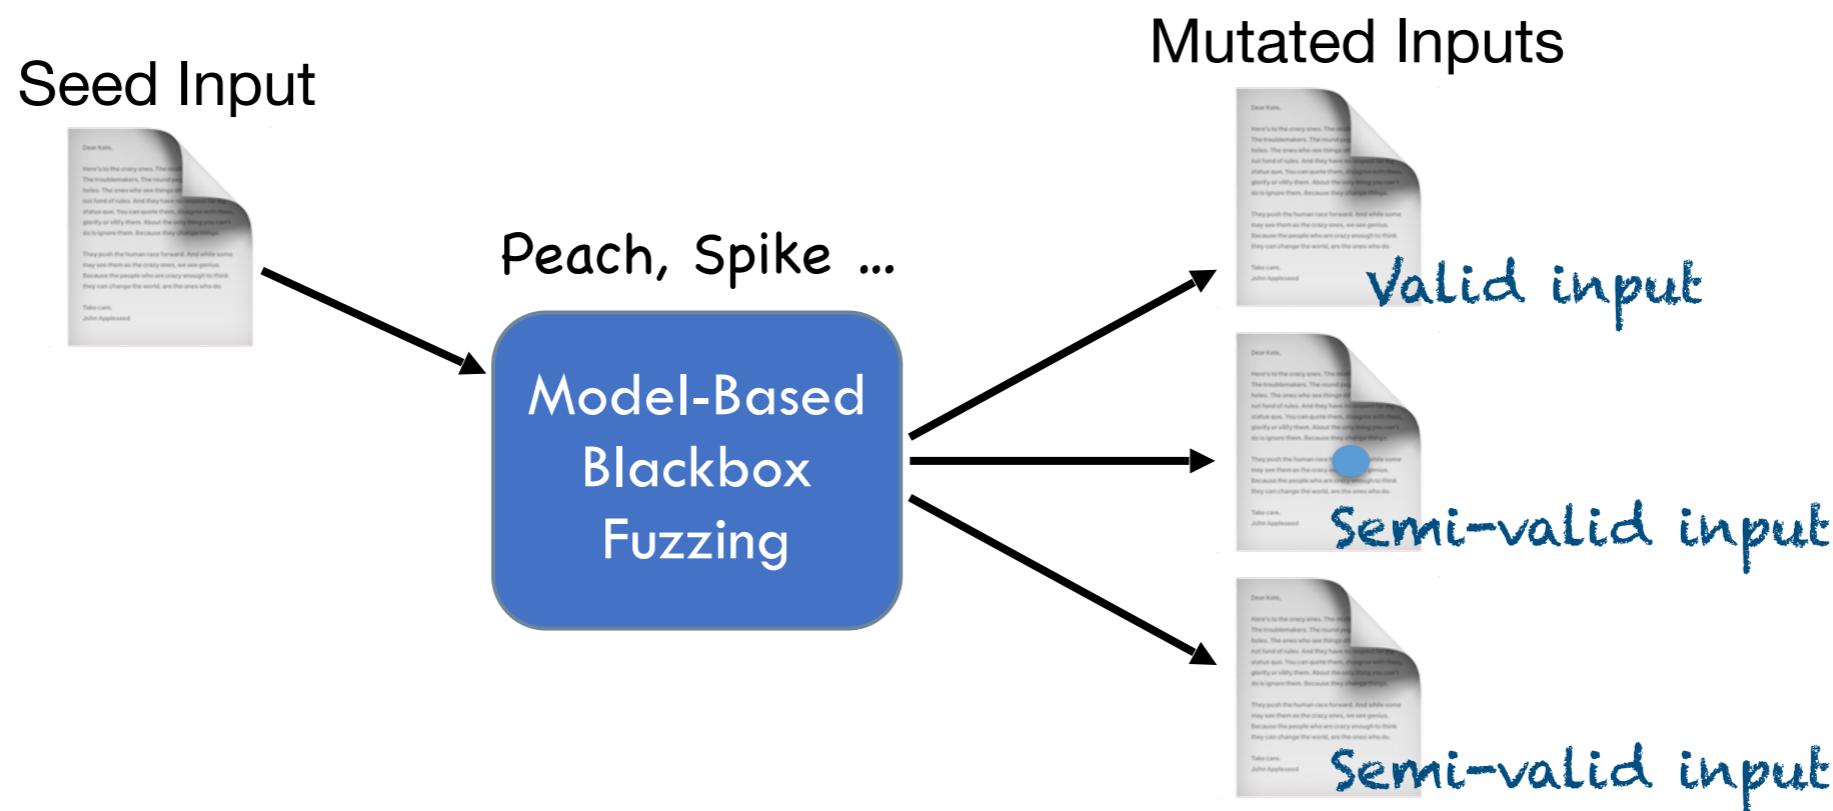
Peach, Spike ...

Model-Based
Blackbox
Fuzzing



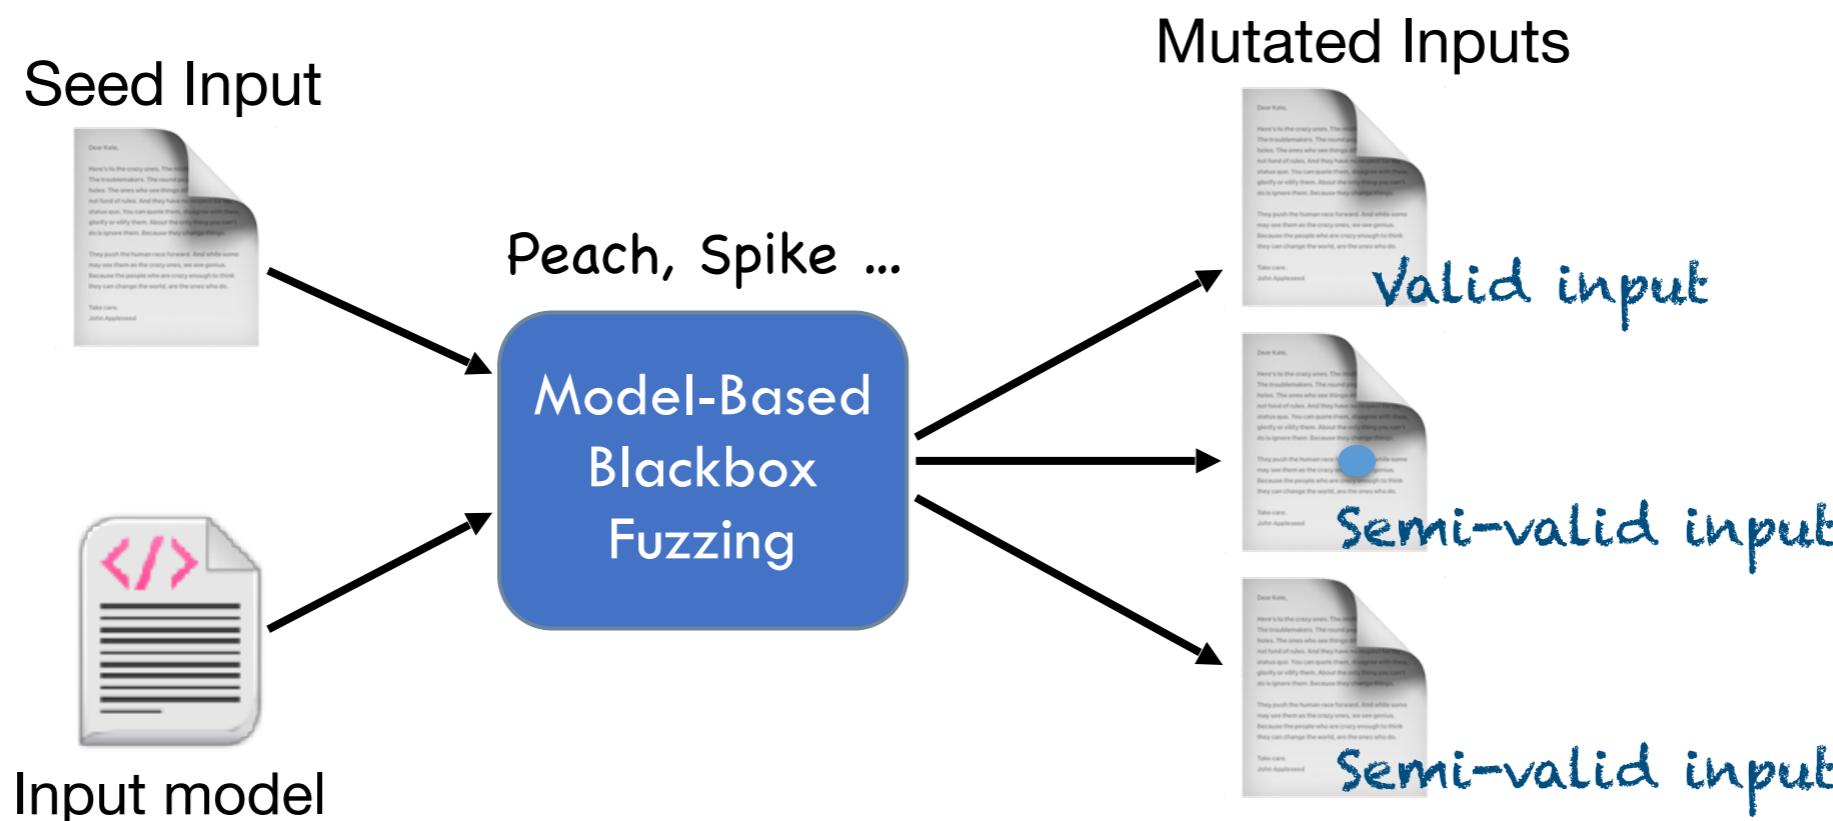
Motivation

- Automated vulnerability detection techniques
 - 1. **Blackbox Fuzzing** (no program analysis, no feedback)



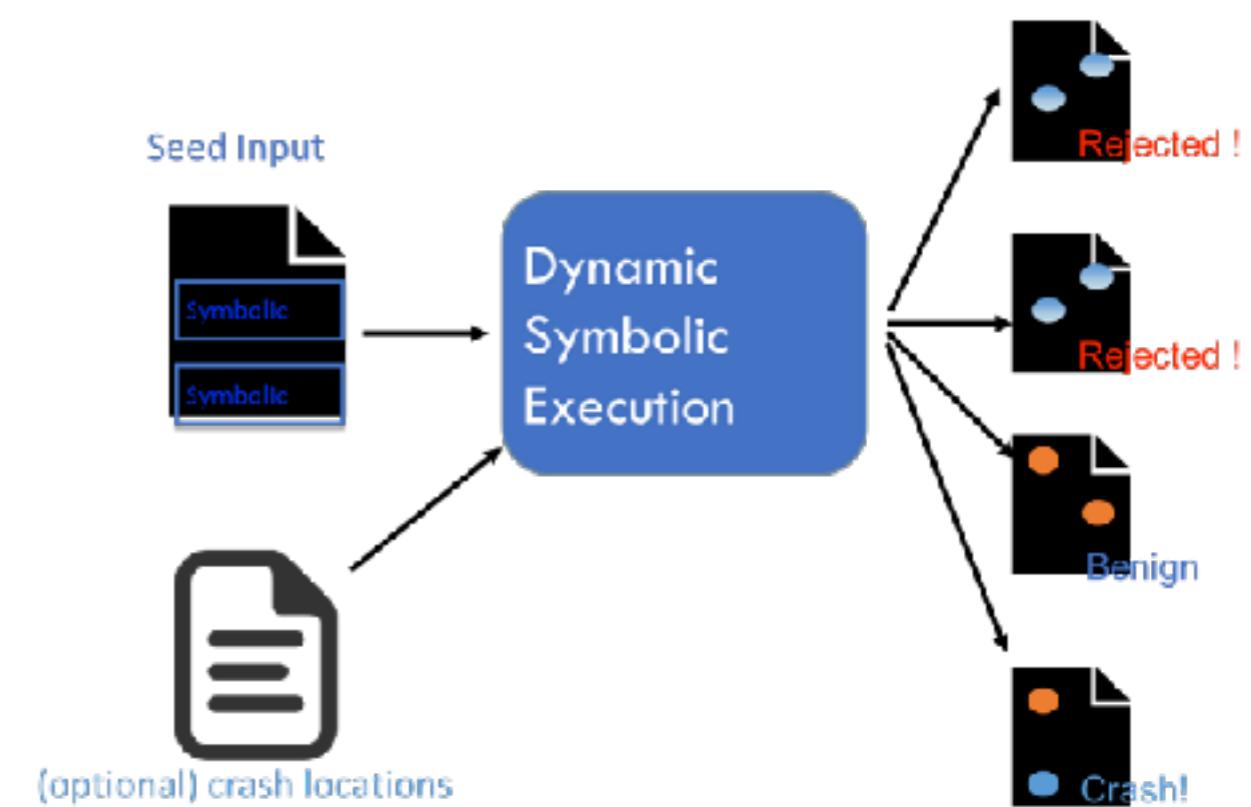
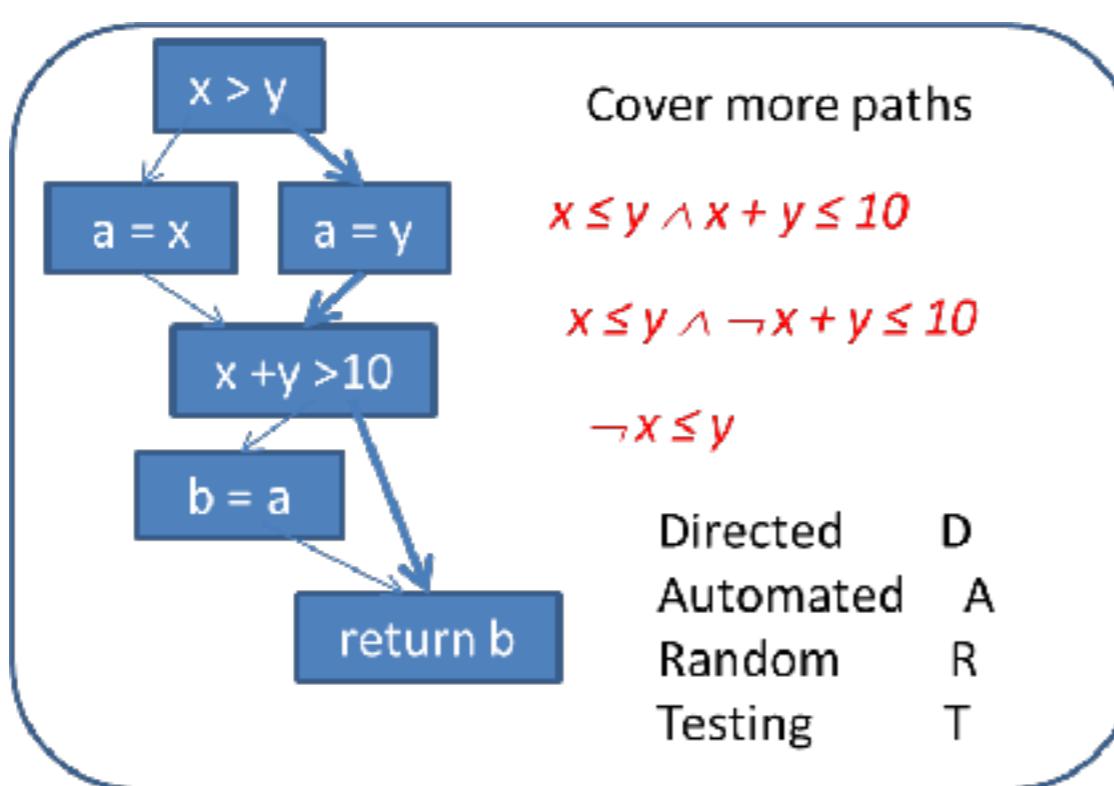
Motivation

- Automated vulnerability detection techniques
 - 1. **Blackbox Fuzzing** (no program analysis, no feedback)



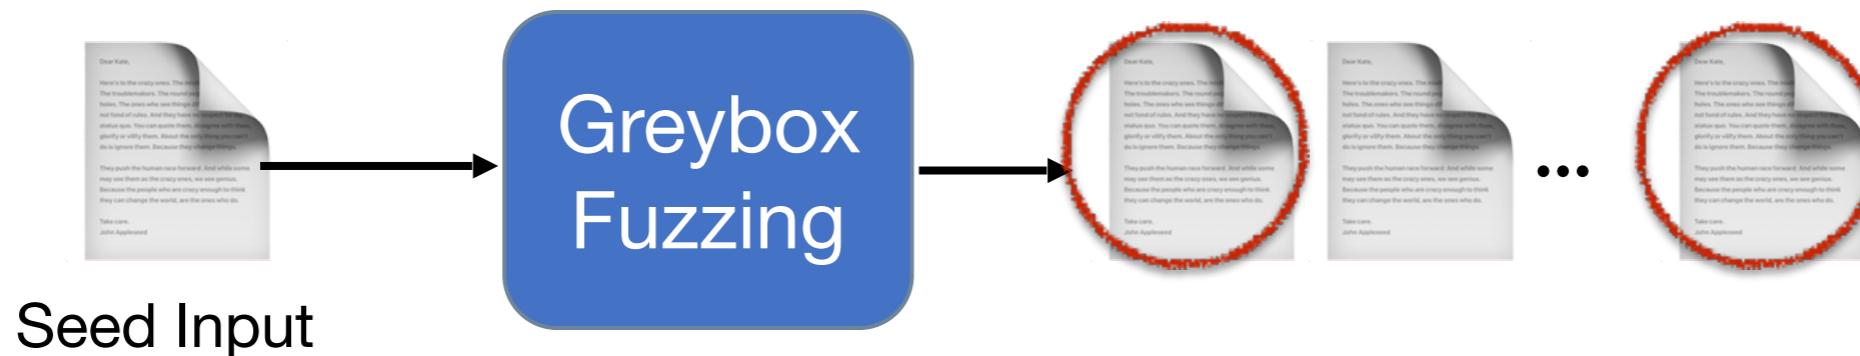
Motivation

- Automated vulnerability detection techniques
 - Blackbox Fuzzing** (no program analysis, no feedback)
 - Whitebox Fuzzing** (mostly program analysis)



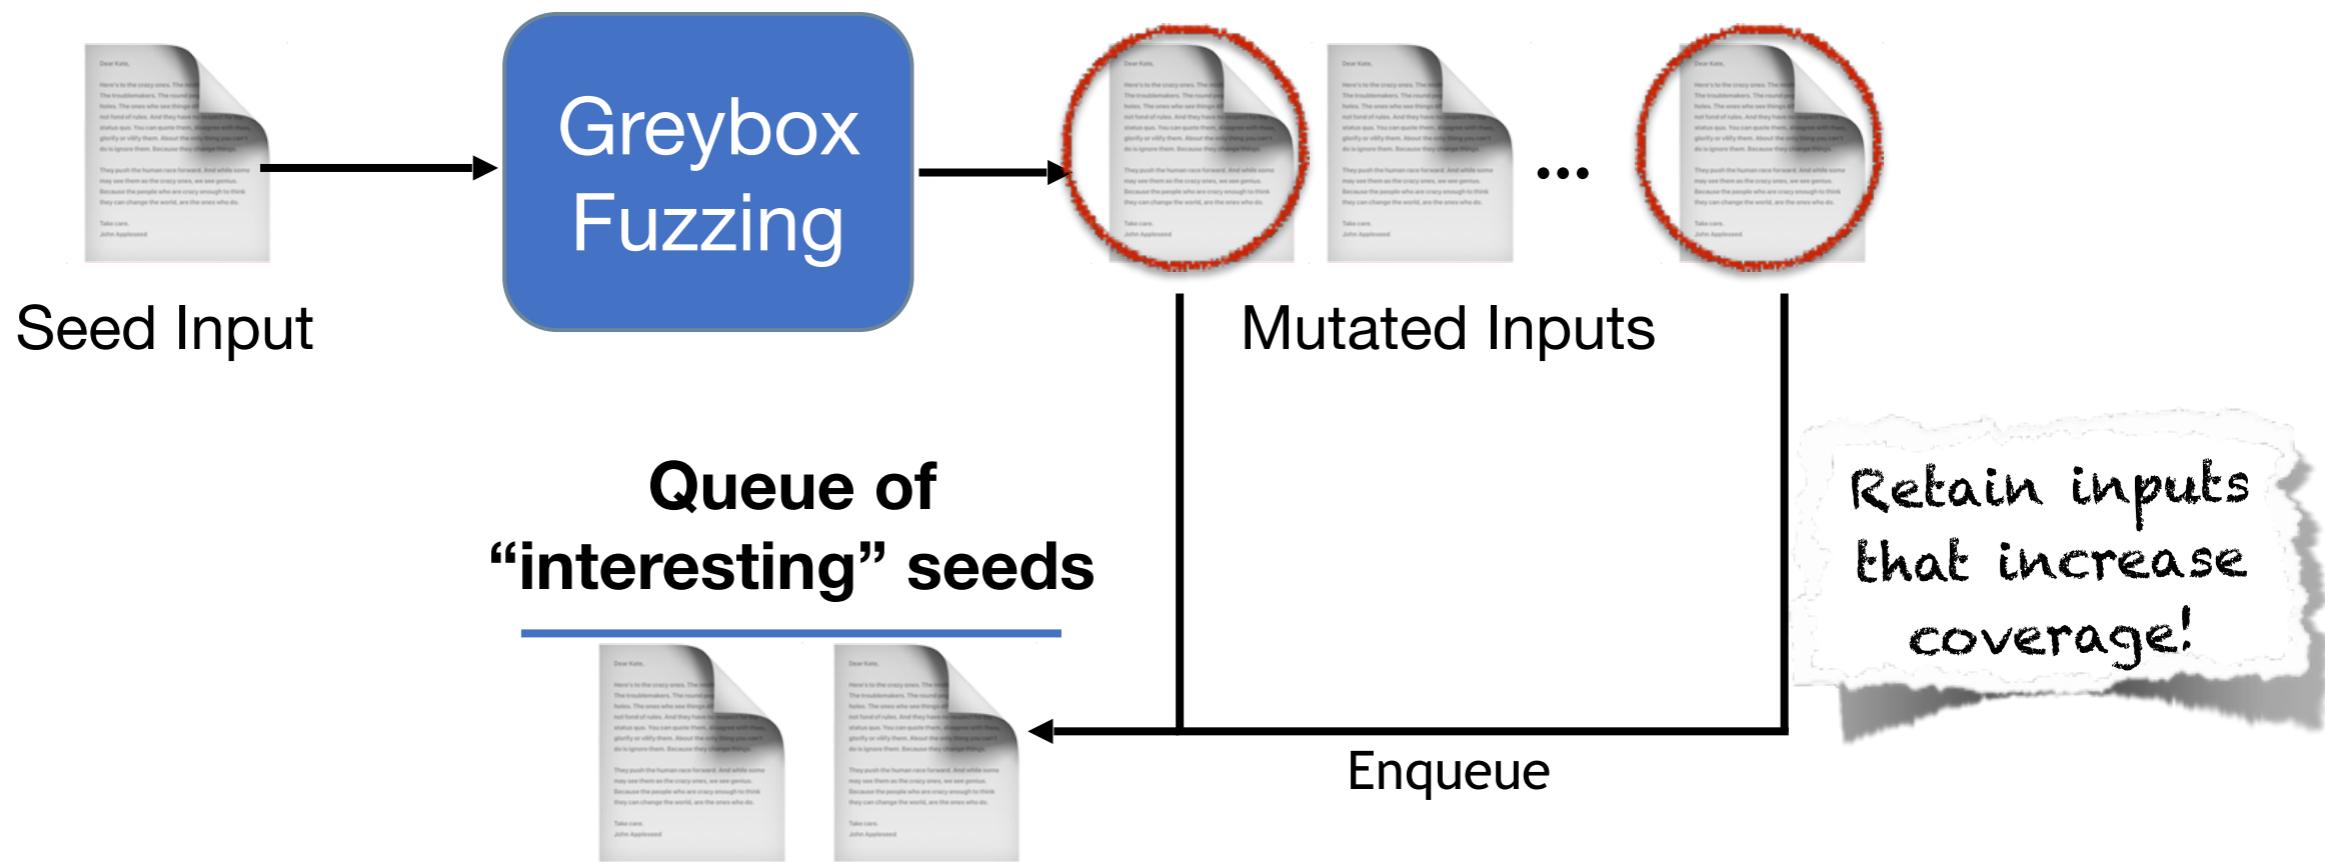
Motivation

- Automated vulnerability detection techniques
 1. **Blackbox Fuzzing** (no program analysis, no feedback)
 2. **Whitebox Fuzzing** (mostly program analysis)
 3. **Greybox Fuzzing** (no program analysis, but coverage feedback)



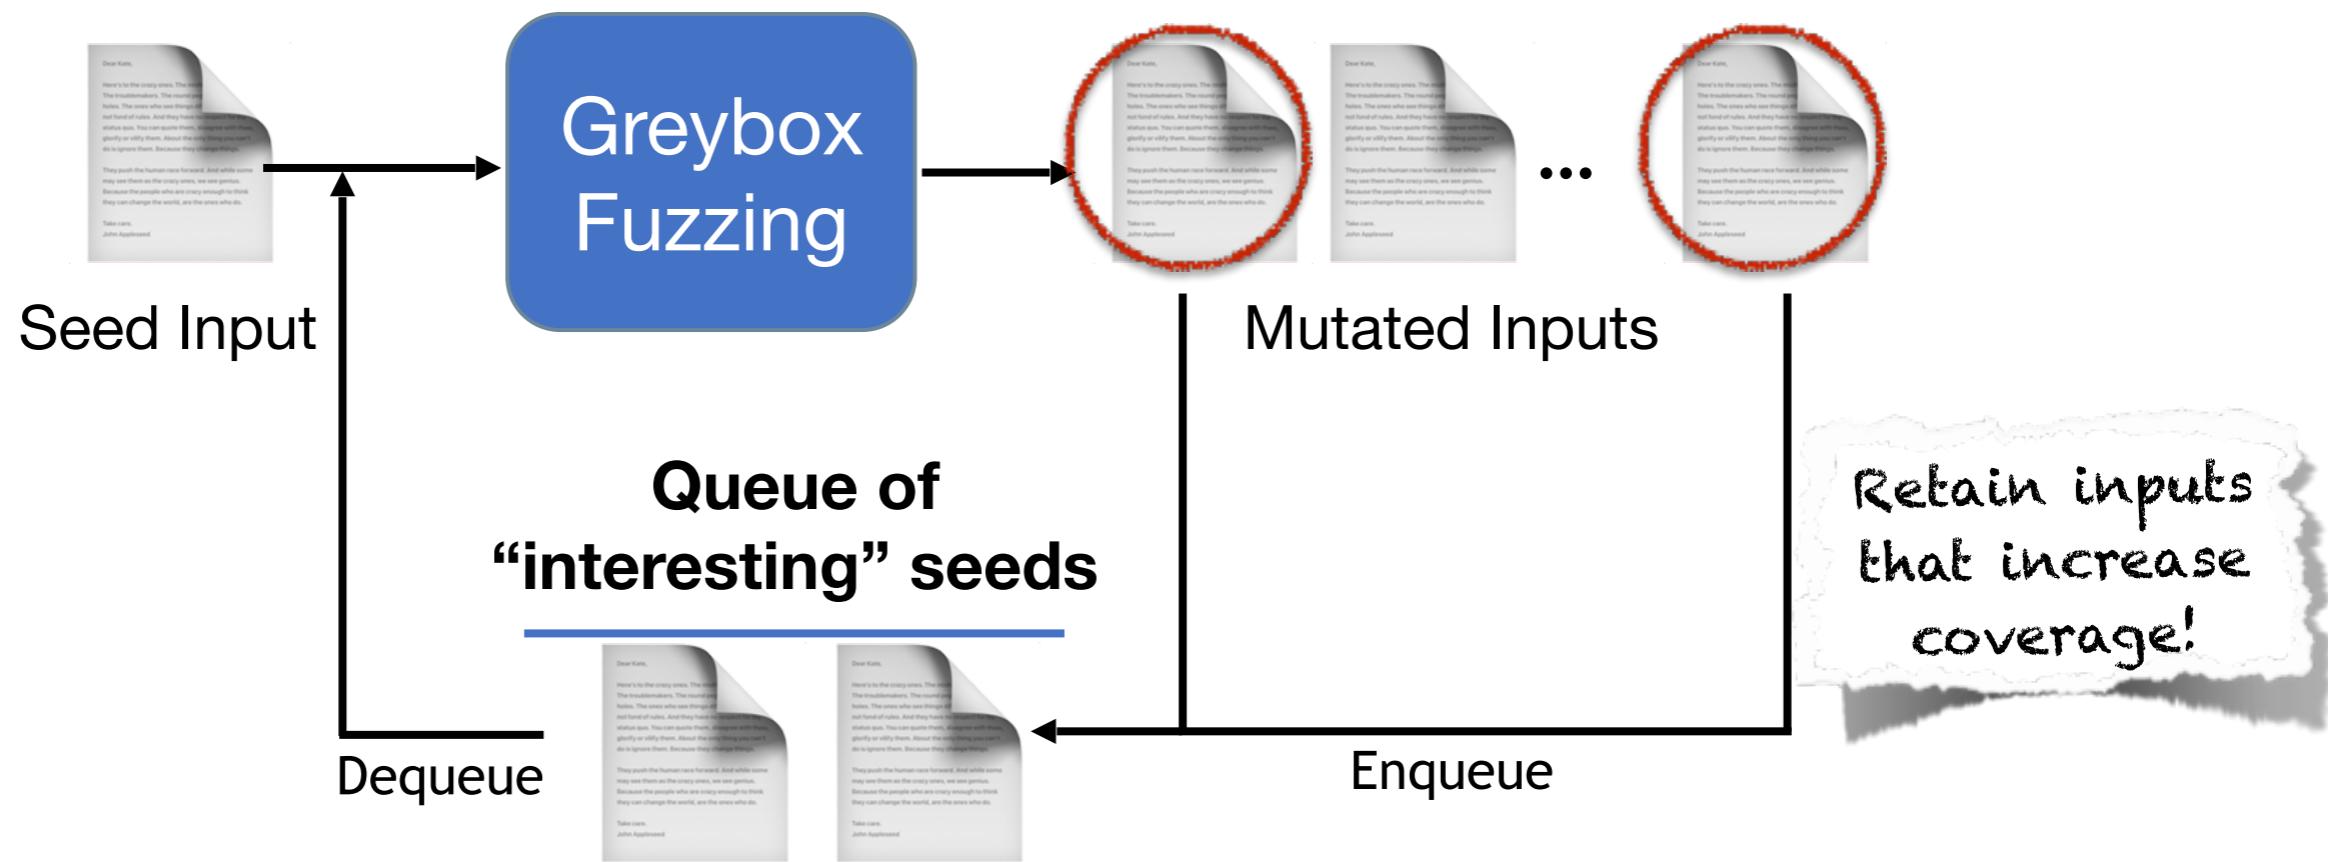
Motivation

- Automated vulnerability detection techniques
 1. **Blackbox Fuzzing** (no program analysis, no feedback)
 2. **Whitebox Fuzzing** (mostly program analysis)
 3. **Greybox Fuzzing** (no program analysis, but coverage feedback)



Motivation

- Automated vulnerability detection techniques
 1. **Blackbox Fuzzing** (no program analysis, no feedback)
 2. **Whitebox Fuzzing** (mostly program analysis)
 3. **Greybox Fuzzing** (no program analysis, but coverage feedback)



Motivation

- **Greybox Fuzzing** is frequently used
 - **State-of-the-art** in automated vulnerability detection
 - **Extremely efficient** coverage-based input generation
 - All program analysis before/at **instrumentation time**.
 - Start with a seed corpus, choose a seed file, **fuzz it**.
 - Add to corpus **only if new input increases coverage**.

Motivation

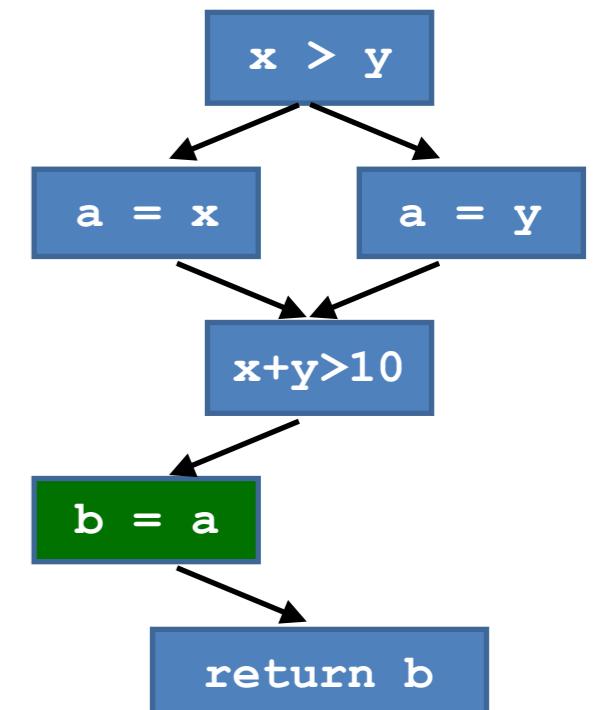
- **Greybox Fuzzing** is frequently used
 - **State-of-the-art** in automated vulnerability detection
 - **Extremely efficient** coverage-based input generation
 - All program analysis before/at **instrumentation time**.
 - Start with a seed corpus, choose a seed file, **fuzz it**.
 - Add to corpus **only if new input increases coverage**.
 - **Cannot be directed!**

Motivation

- **Directed Fuzzing** has many applications
 - **Patch Testing**: reach changed statements
 - **Crash Reproduction**: exercise stack trace
 - **SA Report Verification**: reach “dangerous” location
 - **Information Flow Detection**: exercise source-sink pairs

Motivation

- **Directed Fuzzing:** classical **constraint satisfaction prob.**
 - **Program analysis** to identify **program paths** that reach given program locations.
 - **Symbolic Execution** to derive **path conditions** for any of the identified paths.
 - **Constraint Solving** to find an **input** that
 - **satisfies the path condition** and thus
 - **reaches a program location** that was given.



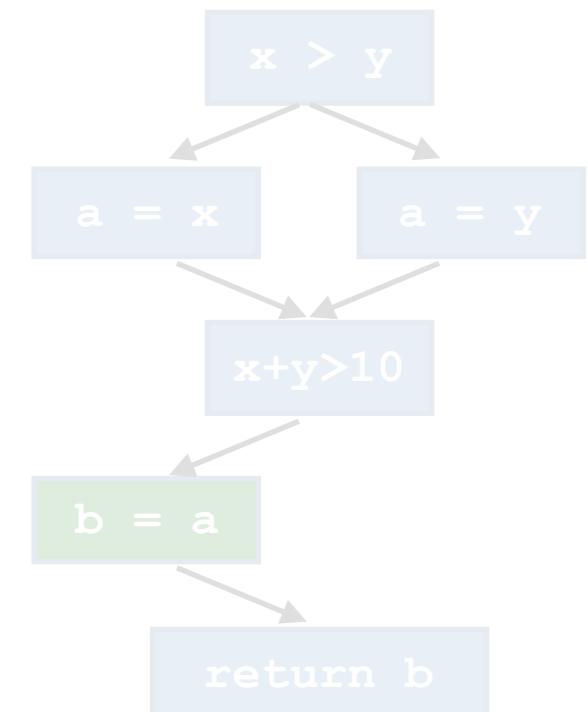
$$\varphi_1 = (x > y) \wedge (x + y > 10)$$

$$\varphi_2 = \neg(x > y) \wedge (x + y > 10)$$

Motivation

- **Directed Fuzzing:** classical **constraint satisfaction prob.**
 - Program analysis to identify program paths that reach given program locations.
 - Symbolic Execution to derive path conditions for any of the identified paths.
 - Constraint Solving to find an input that

Requires heavy-weight machinery!
 - satisfies the path condition and thus
 - reaches a program location that was given.



$$\varphi_1 = (x > y) \wedge (x + y > 10)$$

$$\varphi_2 = \neg(x > y) \wedge (x + y > 10)$$

Overview

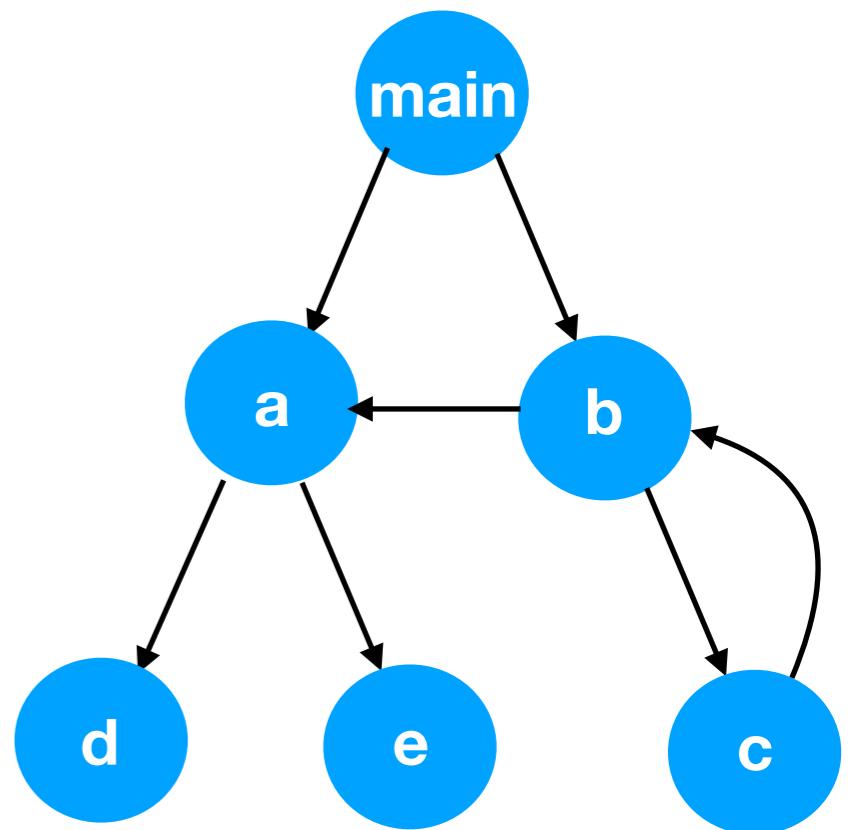
- **Directed Fuzzing as optimisation problem!**
 1. Instrumentation Time:
 1. Extract **call graph** (CG) and **control-flow graphs** (CFGs).
 2. For each **BB**, compute **distance** to target locations.
 3. Instrument program to **aggregate distance values**.

Overview

- **Directed Fuzzing as optimisation problem!**
 1. Instrumentation Time:
 1. Extract **call graph** (CG) and **control-flow graphs** (CFGs).
 2. For each **BB**, compute **distance** to target locations.
 3. Instrument program to **aggregate distance values**.
 2. Runtime, **for each input**
 1. collect coverage and distance **information**, and
 2. decide **how long to be fuzzed** based on distance.
 - If input is **closer** to the targets, it is fuzzed for **longer**.
 - If input is **further away** from the targets, it is fuzzed for **shorter**.

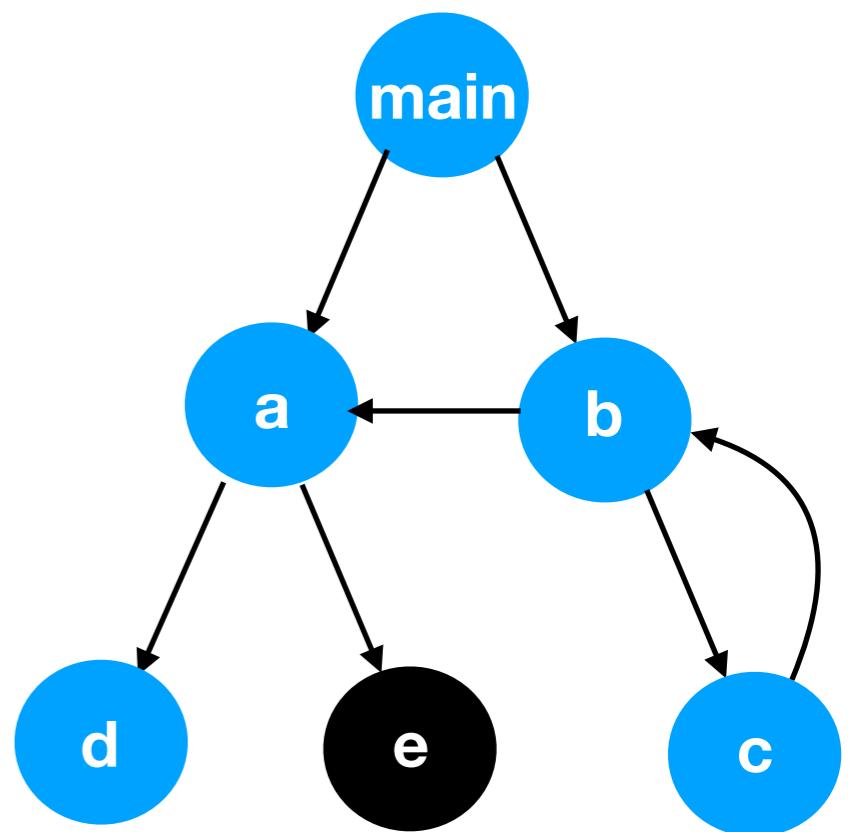
Instrumentation

- **Function-level target distance** using call graph (CG)



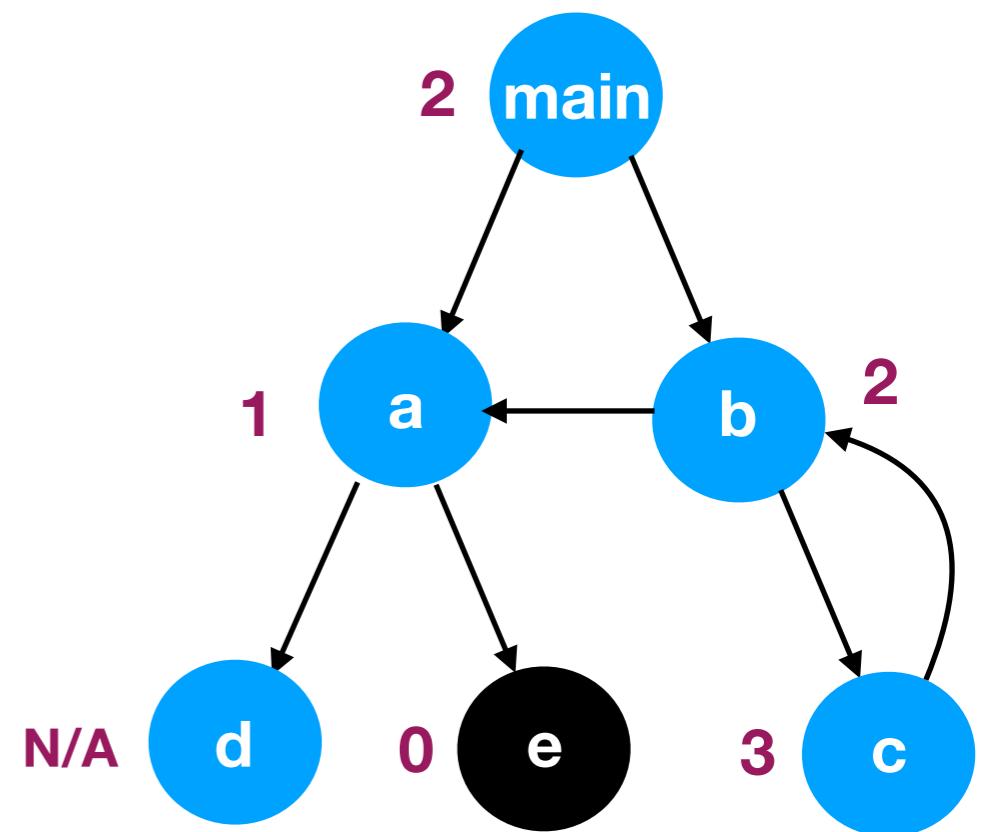
Instrumentation

- **Function-level target distance** using call graph (CG)
 1. Identify **target functions** in CG



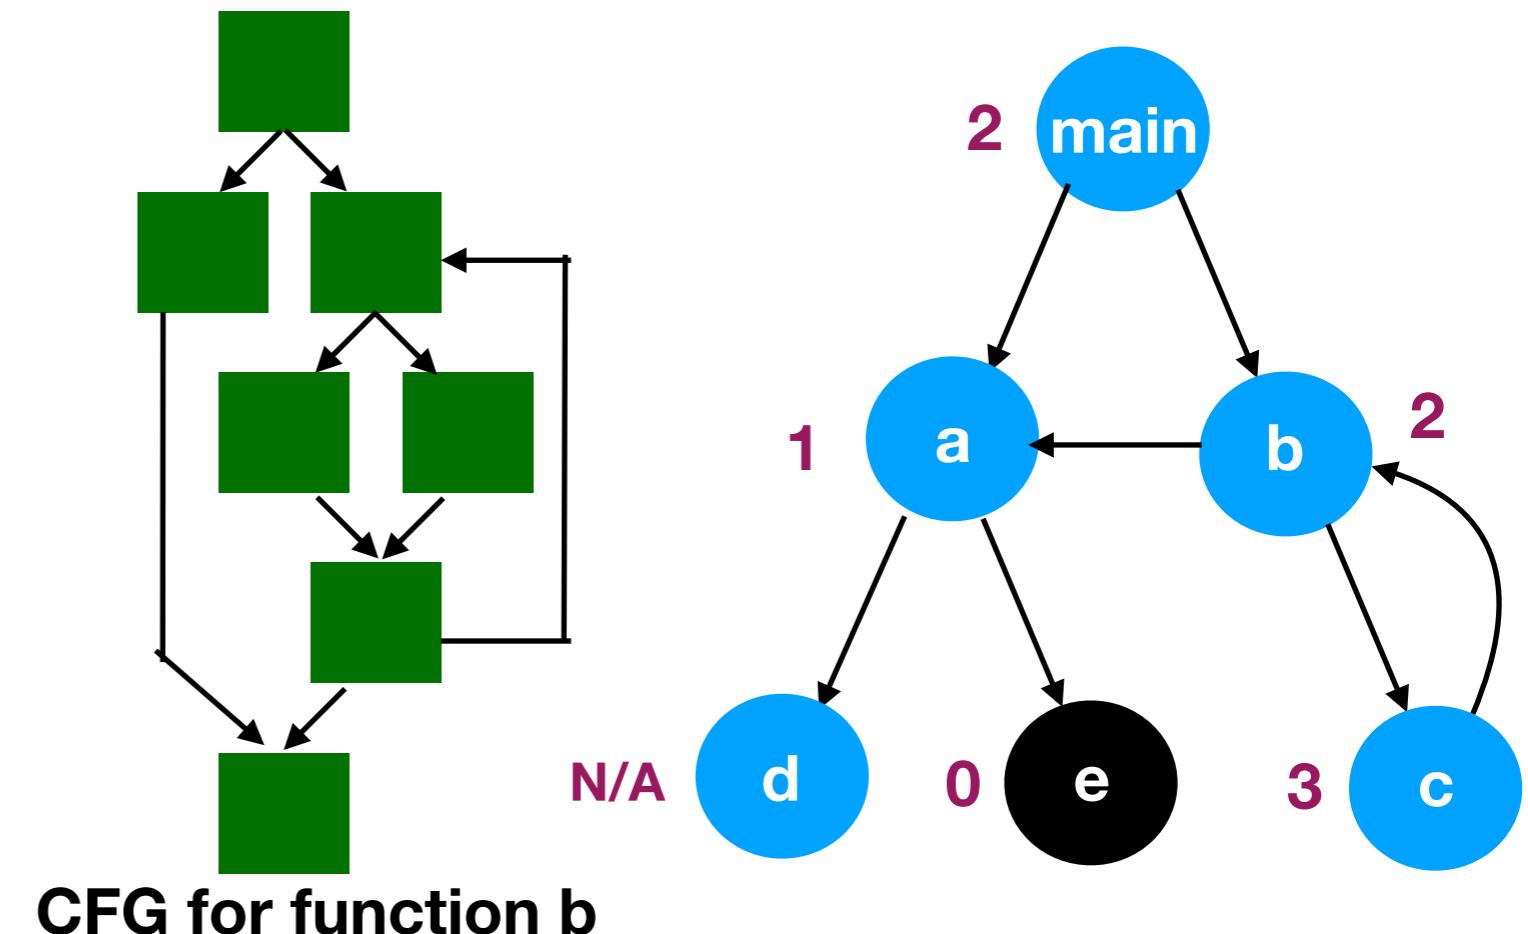
Instrumentation

- **Function-level target distance** using call graph (CG)
 1. Identify **target functions** in CG
 2. For each **function**, compute the harmonic mean of the **length** of the shortest path to targets



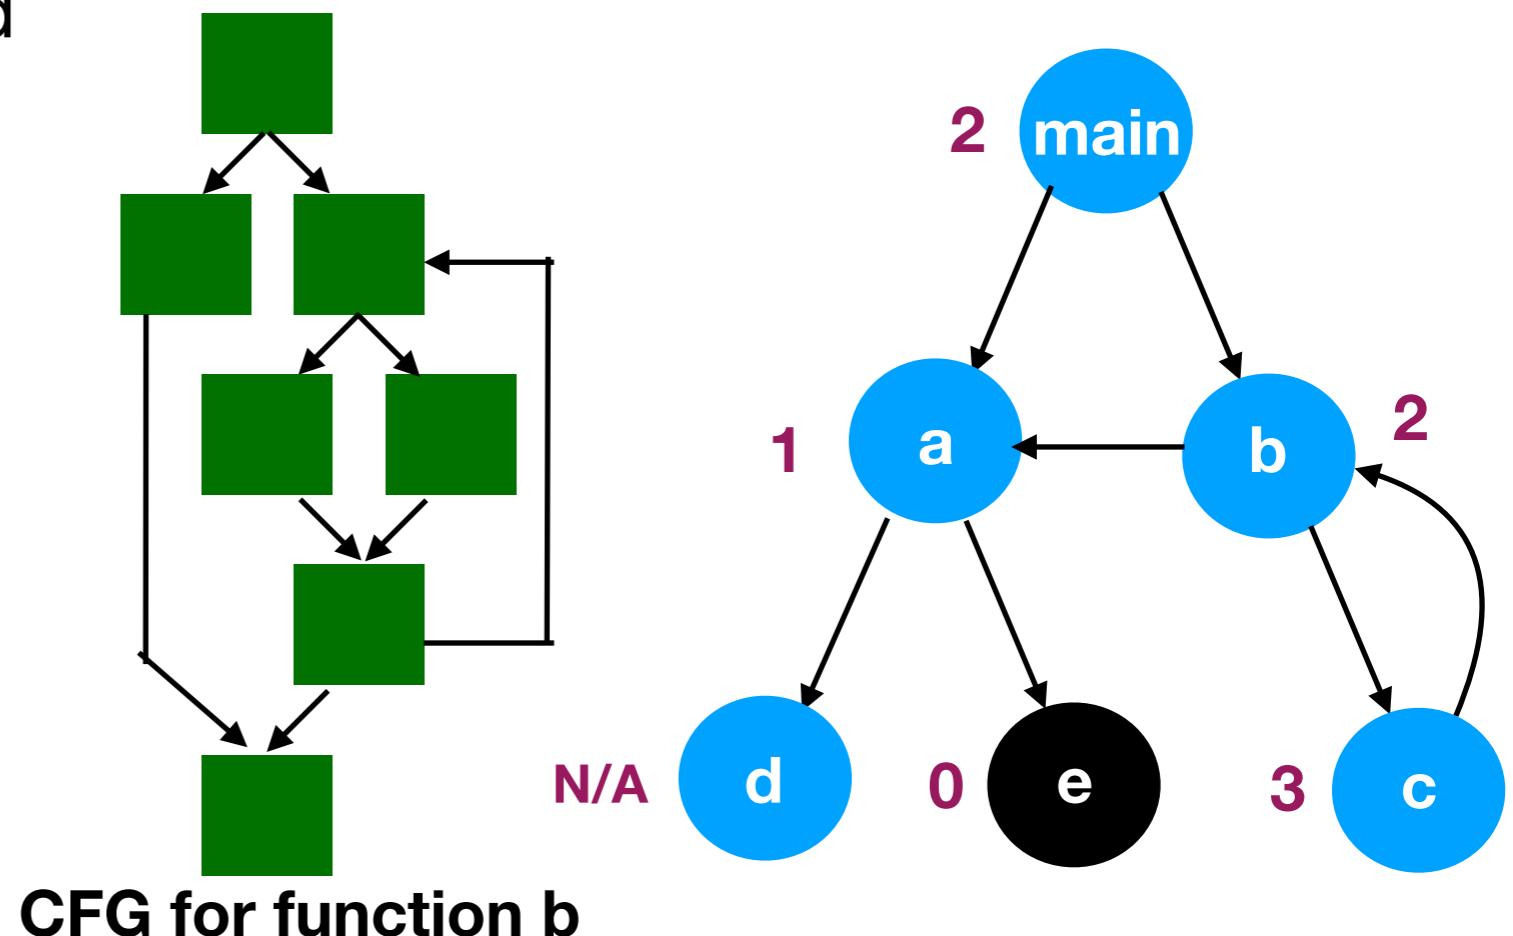
Instrumentation

- **Function-level target distance** using call graph (CG)
- **BB-level target distance** using CFG



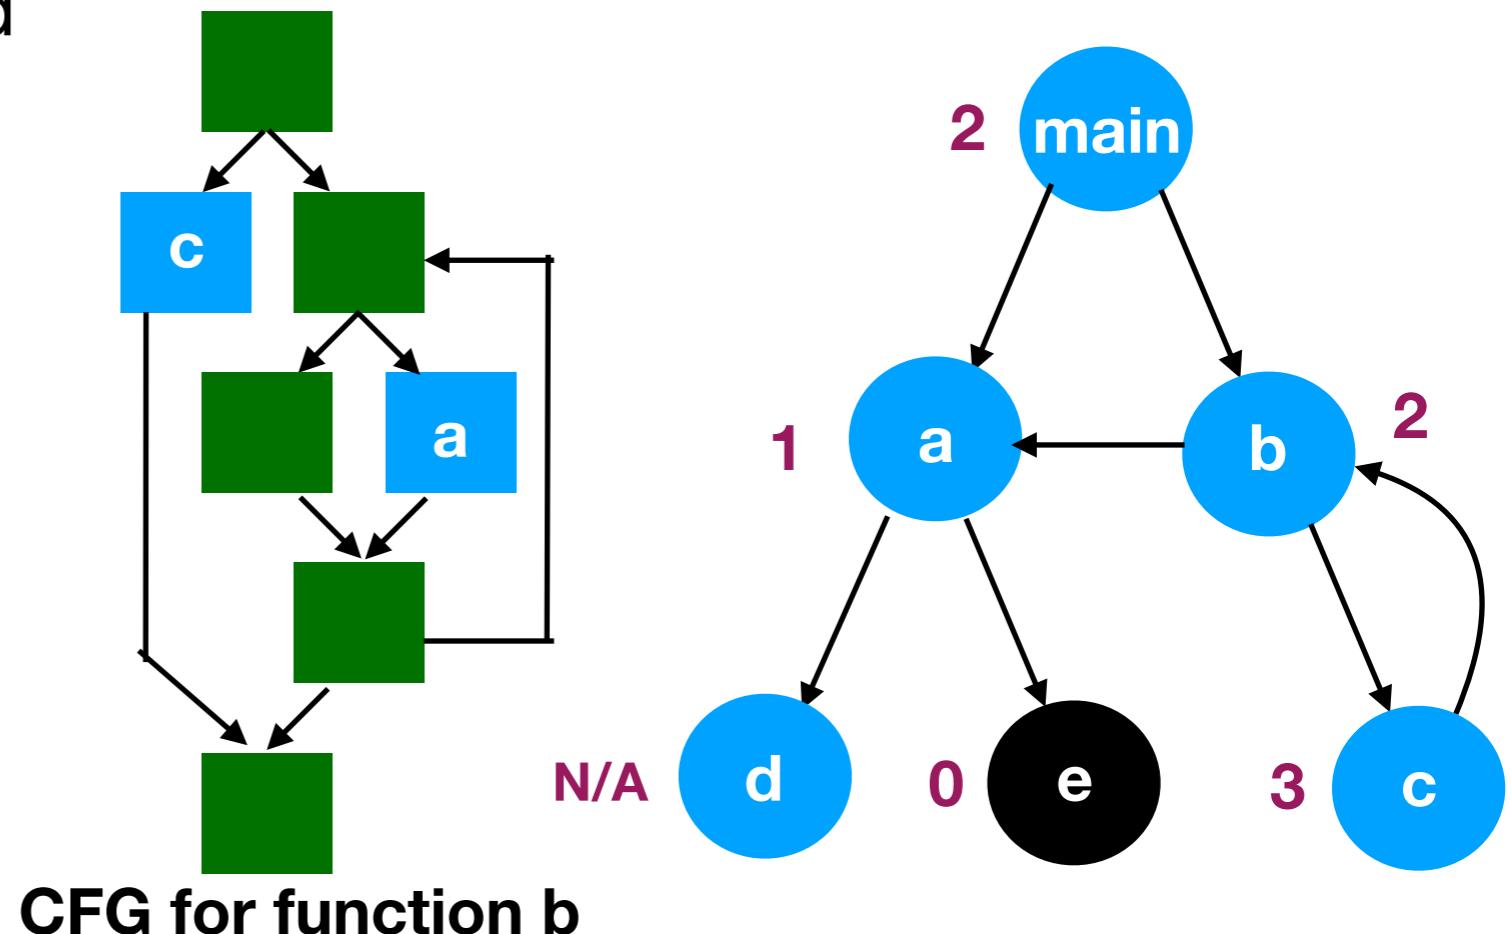
Instrumentation

- **Function-level target distance** using call graph (CG)
- **BB-level target distance** using control-flow graph (CFG)
 1. Identify **target BBs** and assign **distance 0**
(none in function b)



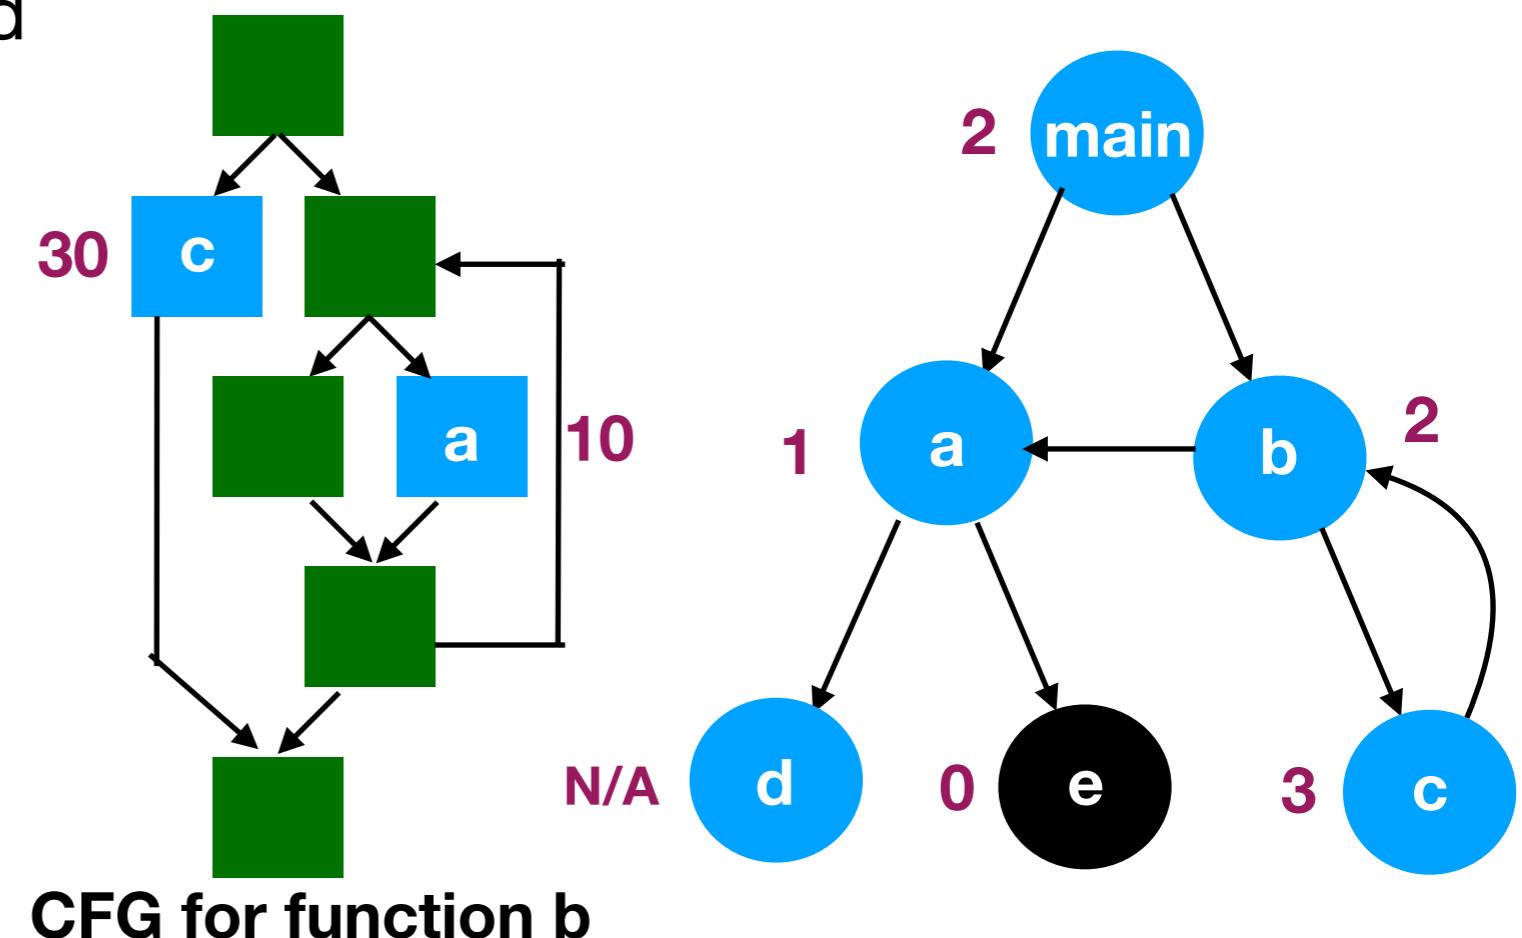
Instrumentation

- **Function-level target distance** using call graph (CG)
- **BB-level target distance** using control-flow graph (CFG)
 1. Identify **target BBs** and assign **distance 0**
 2. Identify BBs that call **functions**



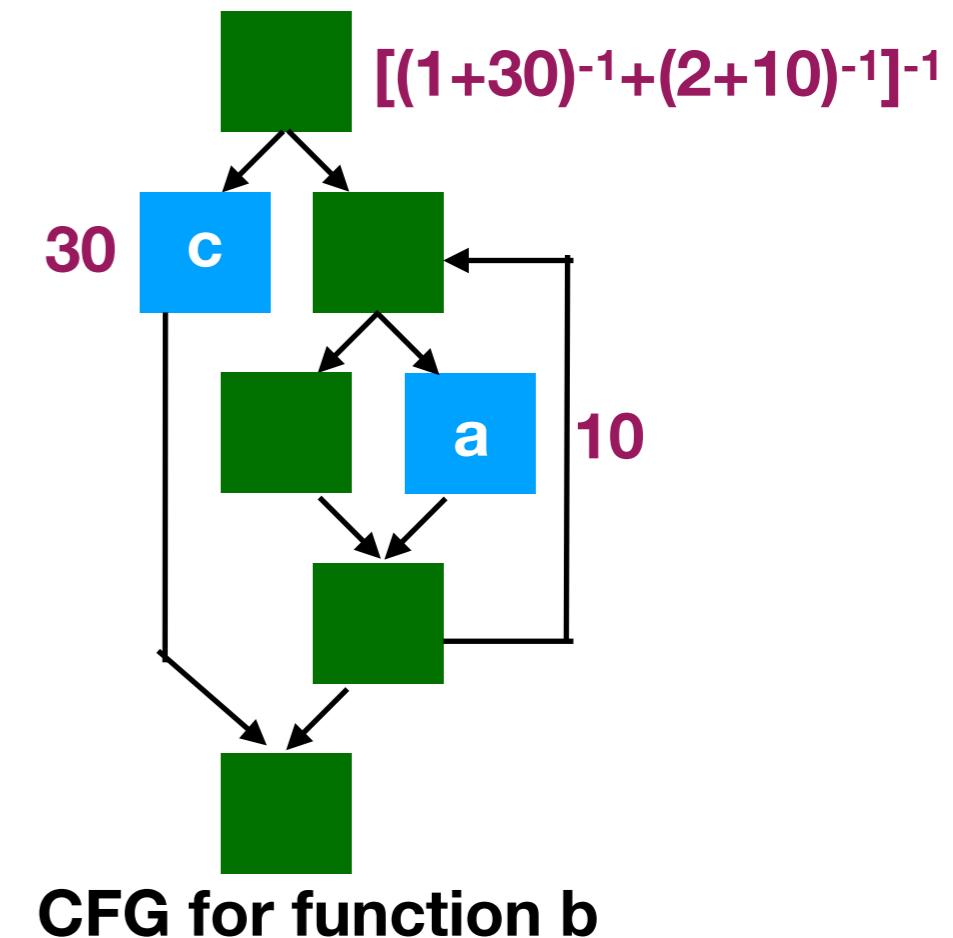
Instrumentation

- **Function-level target distance** using call graph (CG)
- **BB-level target distance** using control-flow graph (CFG)
 1. Identify **target BBs** and assign **distance 0**
 2. Identify BBs that call **functions** and assign **10^*FLTD**



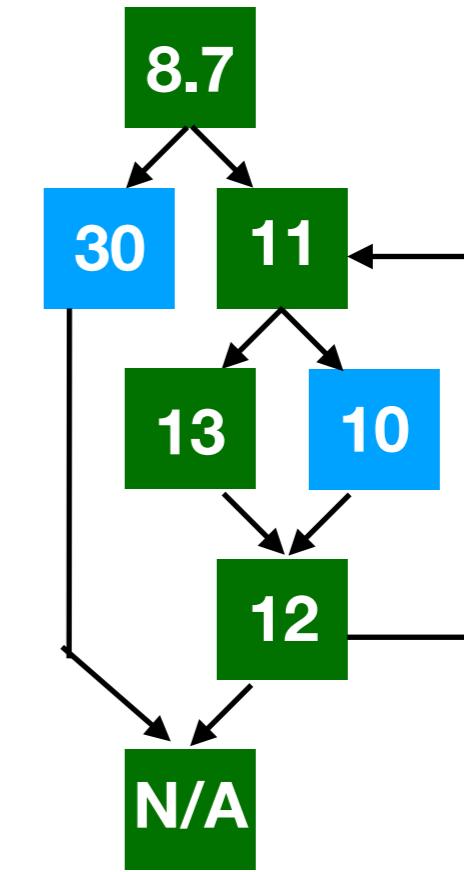
Instrumentation

- **Function-level target distance** using call graph (CG)
- **BB-level target distance** using control-flow graph (CFG)
 1. Identify **target BBs** and assign **distance 0**
 2. Identify BBs that call **functions** and assign **10^*FLTD**
 3. For **each BB**, compute harmonic mean of (length of shortest path to any function-calling BB + 10^*FLTD).



Instrumentation

- **Function-level target distance** using call graph (CG)
- **BB-level target distance** using control-flow graph (CFG)
 1. Identify **target BBs** and assign **distance 0**
 2. Identify BBs that call **functions** and assign **10^*FLTD**
 3. For **each BB**, compute harmonic mean of (length of shortest path to any function-calling BB + 10^*FLTD).

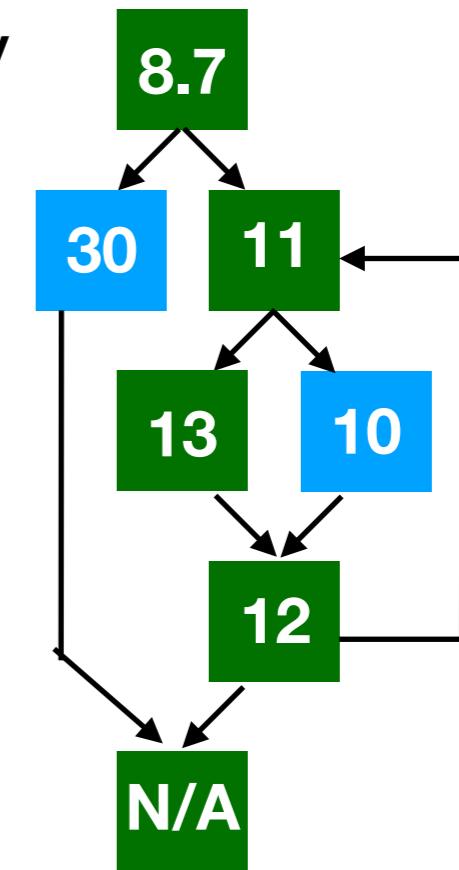


CFG for function b

Runtime



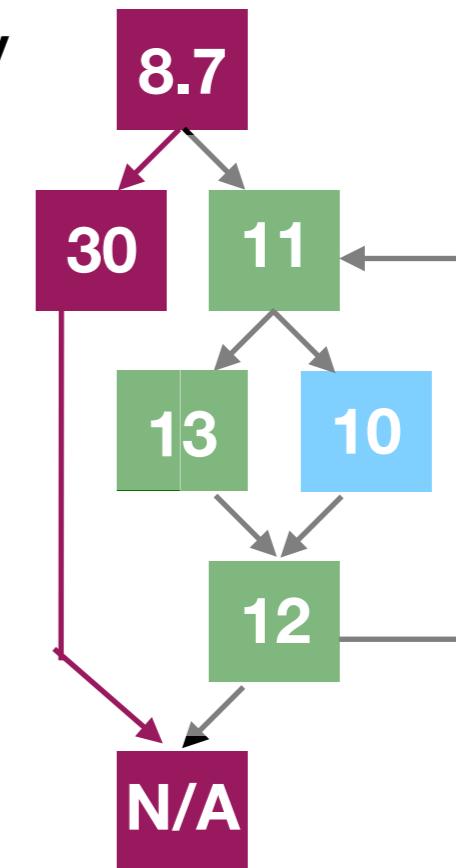
- **Function-level target distance** using call graph (CG)
- **BB-level target distance** using control-flow graph (CFG)
- **Seed distance** from instrumented binary



CFG for function b

Runtime

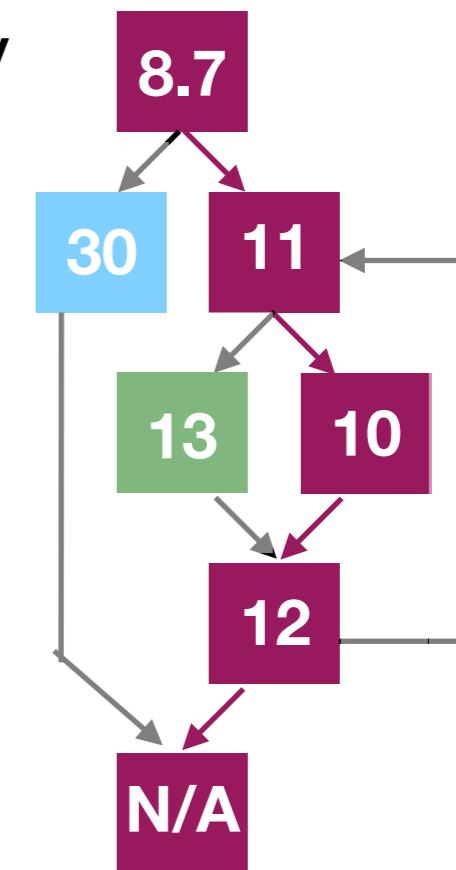
- **Function-level target distance** using call graph (CG)
- **BB-level target distance** using control-flow graph (CFG)
- **Seed distance** from instrumented binary
 - Two 64-bit shared memory entries
 - Aggregated BB-level distance values
 - Number of executed BBs



Seed Distance: 19.4
 $= (8.7+30)/2$

Runtime

- **Function-level target distance** using call graph (CG)
- **BB-level target distance** using control-flow graph (CFG)
- **Seed distance** from instrumented binary
 - Two 64-bit shared memory entries
 - Aggregated BB-level distance values
 - Number of executed BBs

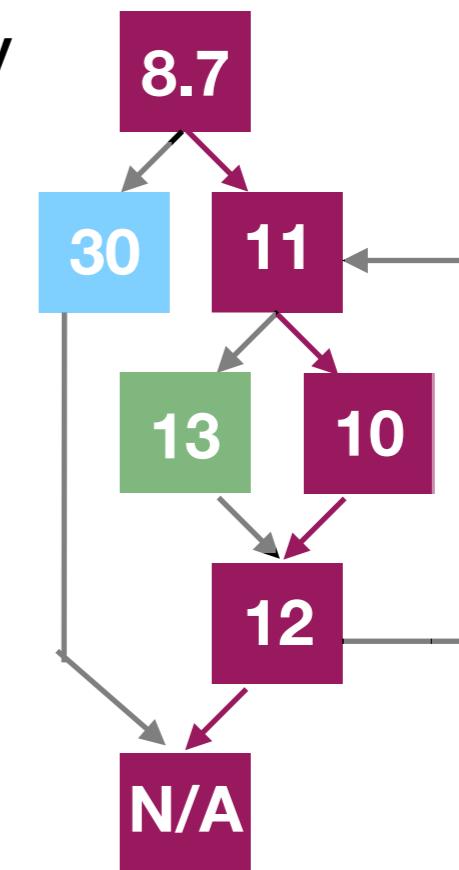


Seed Distance: 10.4
 $= (8.7+11+10+12)/4$

Runtime

- **Function-level target distance** using call graph (CG)
- **BB-level target distance** using control-flow graph (CFG)
- **Seed distance** from instrumented binary
 - Two 64-bit shared memory entries
 - Aggregated BB-level distance values
 - Number of executed BBs

Now that we know how to compute seed distance,
— let's minimise it! —



Seed Distance: 10.4
 $= (8.7+11+10+12)/4$

Directed Fuzzing as Optimisation Problem



- **Background:** Coverage-based Greybox Fuzzing



Seed File

Directed Fuzzing as Optimisation Problem



- **Background:** Coverage-based Greybox Fuzzing

Mutation Operators:

- Bitflips
- Boundary Values
 $(0, 1, -1, \text{INT_MAX}, \text{INT_MIN})$
- Simple arithmetics
(add/subtract 1)
- Block deletion
- Block insertion

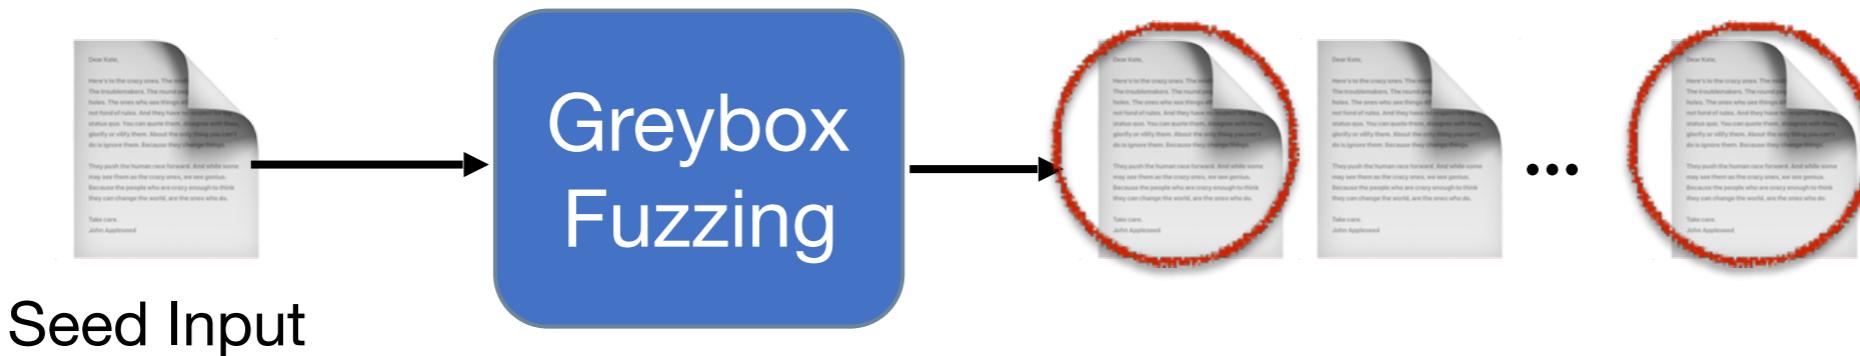


Seed File

Directed Fuzzing as Optimisation Problem

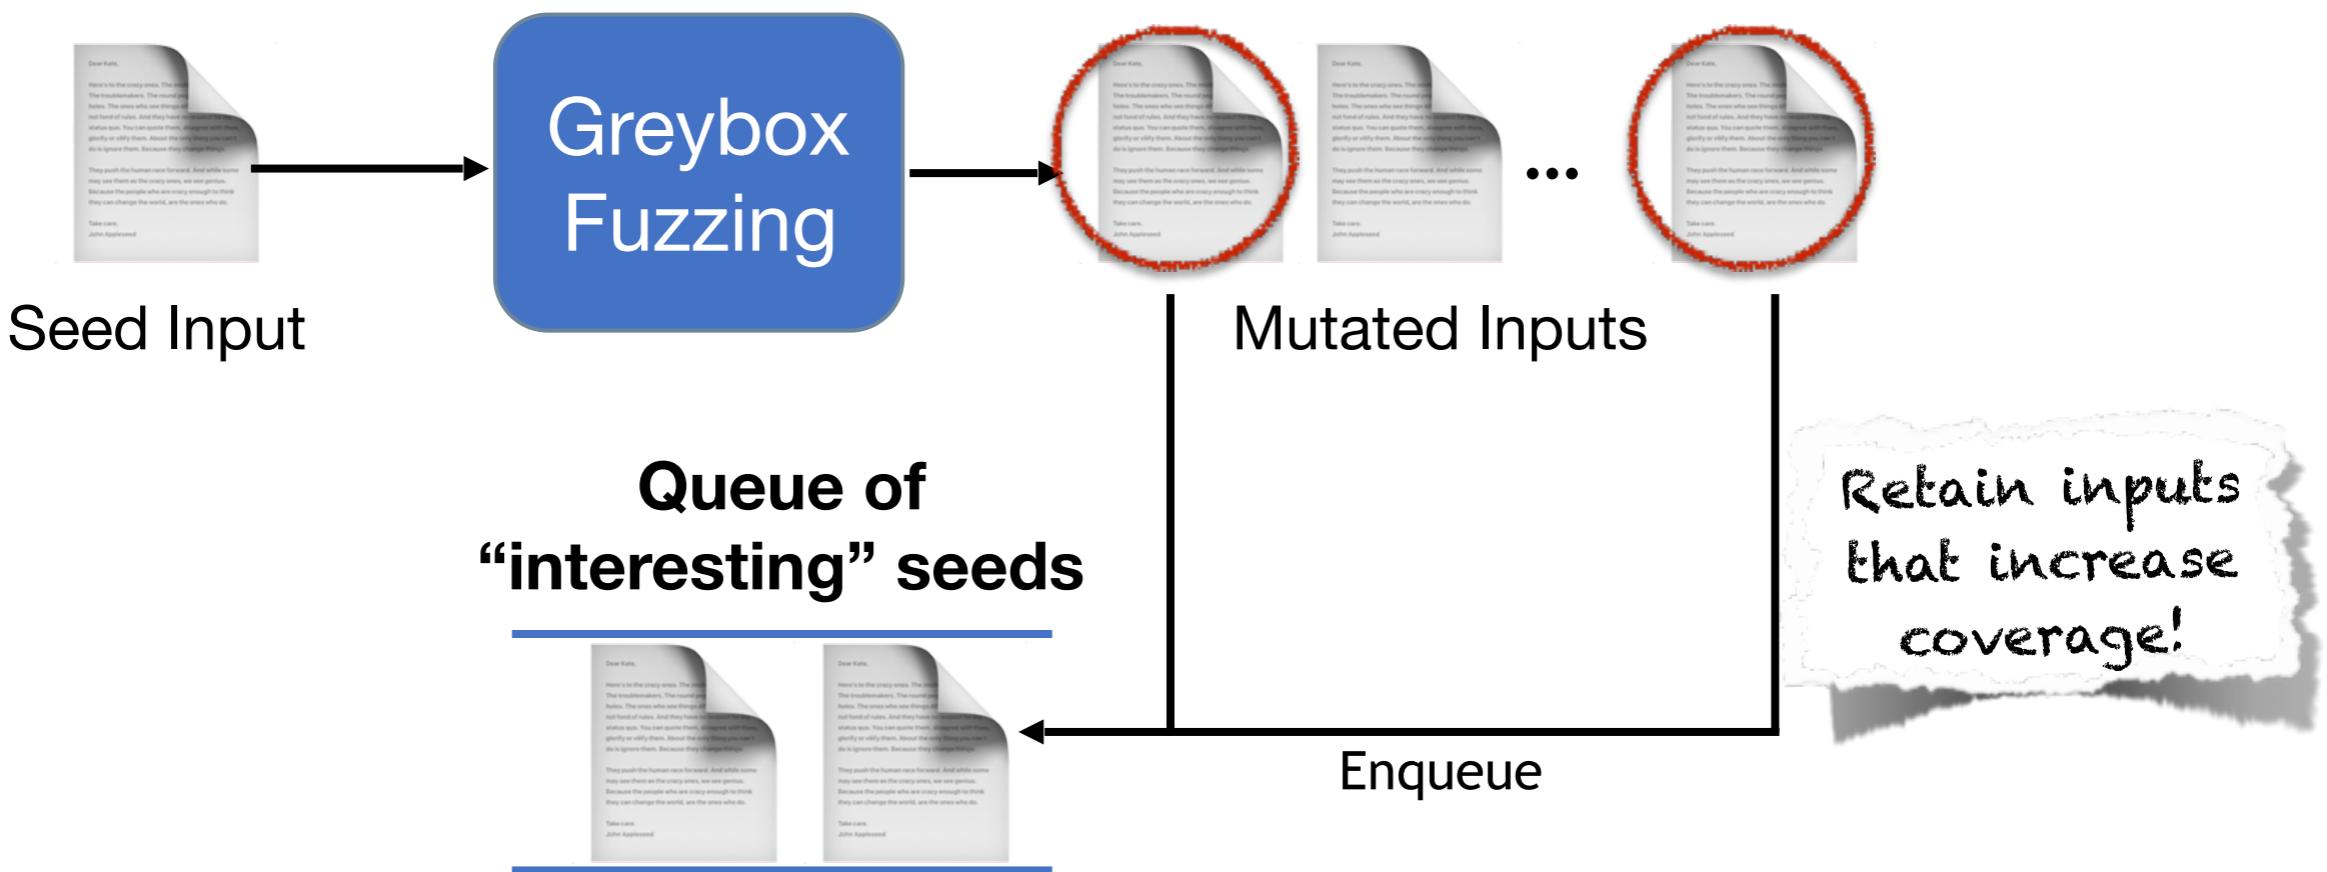


- **Background:** Coverage-based Greybox Fuzzing



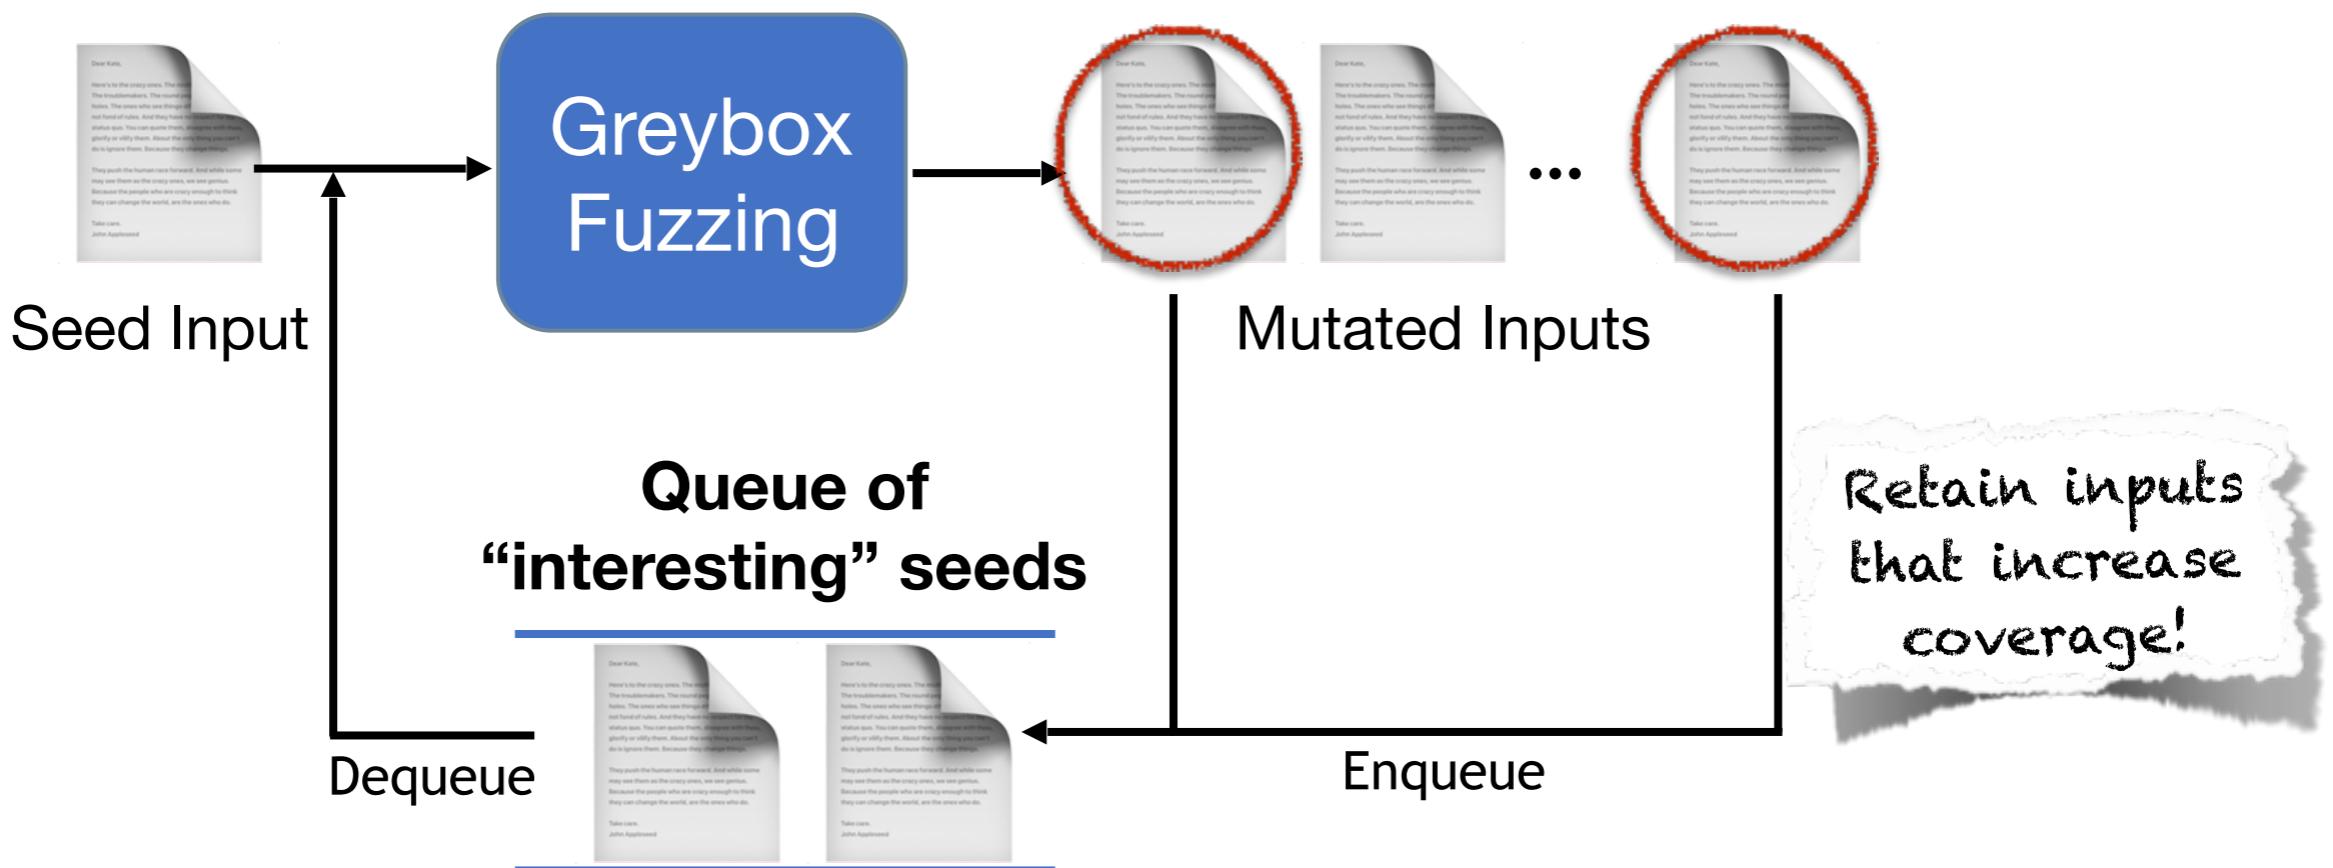
Directed Fuzzing as Optimisation Problem

- **Background:** Coverage-based Greybox Fuzzing



Directed Fuzzing as Optimisation Problem

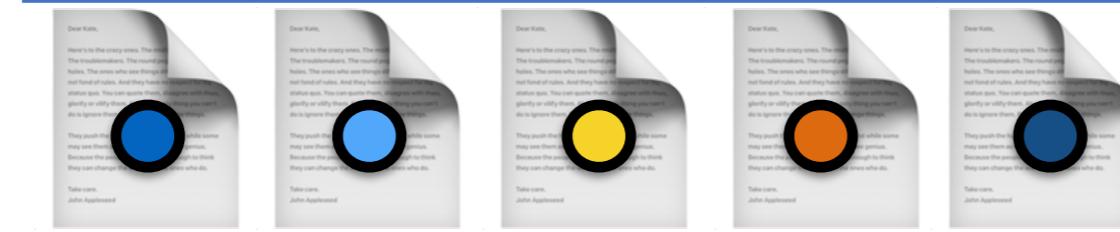
- **Background:** Coverage-based Greybox Fuzzing



Directed Fuzzing as Optimisation Problem

- **Background:** Coverage-based Greybox Fuzzing

Queue of
“interesting” seeds



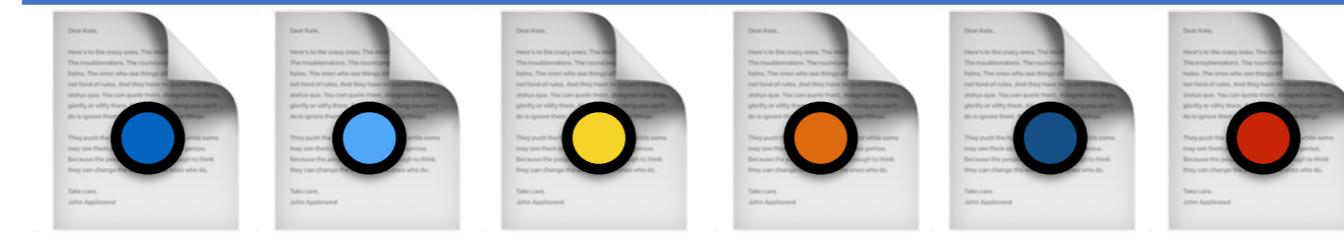
low energy
- generated less
test inputs!



Directed Fuzzing as Optimisation Problem

- **Background:** Coverage-based Greybox Fuzzing

Queue of
“interesting” seeds



high energy
- generated more
test inputs!



Directed Fuzzing as Optimisation Problem



- **Background:** Coverage-based Greybox Fuzzing
 - Seed's energy:
 - Number of inputs generated when chosen for fuzzing
 - Local property: each seed has its own energy
 - Power schedule:
 - Assigns energy to seeds according to a pre-defined formula
- ★ Boosted Greybox Fuzzing (AFLFast CCS'16)
 - Assign **more energy** to seeds exercising **low-frequency paths**.
- ★ Directed Greybox Fuzzing (AFLGo CCS'17)
 - Assign **more energy** to seeds that are **closer to the given targets!**

Directed Fuzzing as Optimisation Problem



- **Background:** Coverage-based Greybox Fuzzing
 - Seed's energy:
 - Number of inputs generated when chosen for fuzzing
 - Local property: each seed has its own energy
 - Power schedule:
 - Assigns energy to seeds according to a pre-defined formula
- ★ **Boosted Greybox Fuzzing (AFLFast CCS'16)**
 - Assign **more energy** to seeds exercising **low-frequency paths**.
- ★ **Directed Greybox Fuzzing (AFLGo CCS'17)**
 - Assign **more energy** to seeds that are **closer to the given targets!**

Directed Fuzzing as Optimisation Problem



- **Directed Greybox Fuzzing**
 - Assign **more energy** to seeds that are **closer** to the given targets!
 - **Problem** (Stochastic Gradient Descent)
 - If we **always** assign **more energy** to **closer** seeds, we **typically** reach only a **local** minimum, but **never** a **global** minimum distance!
 - **Solution** (Simulated Annealing)
 - **Sometimes** assign **more energy** to **further-away** seeds!
 - Approaches **global** minimum distance.

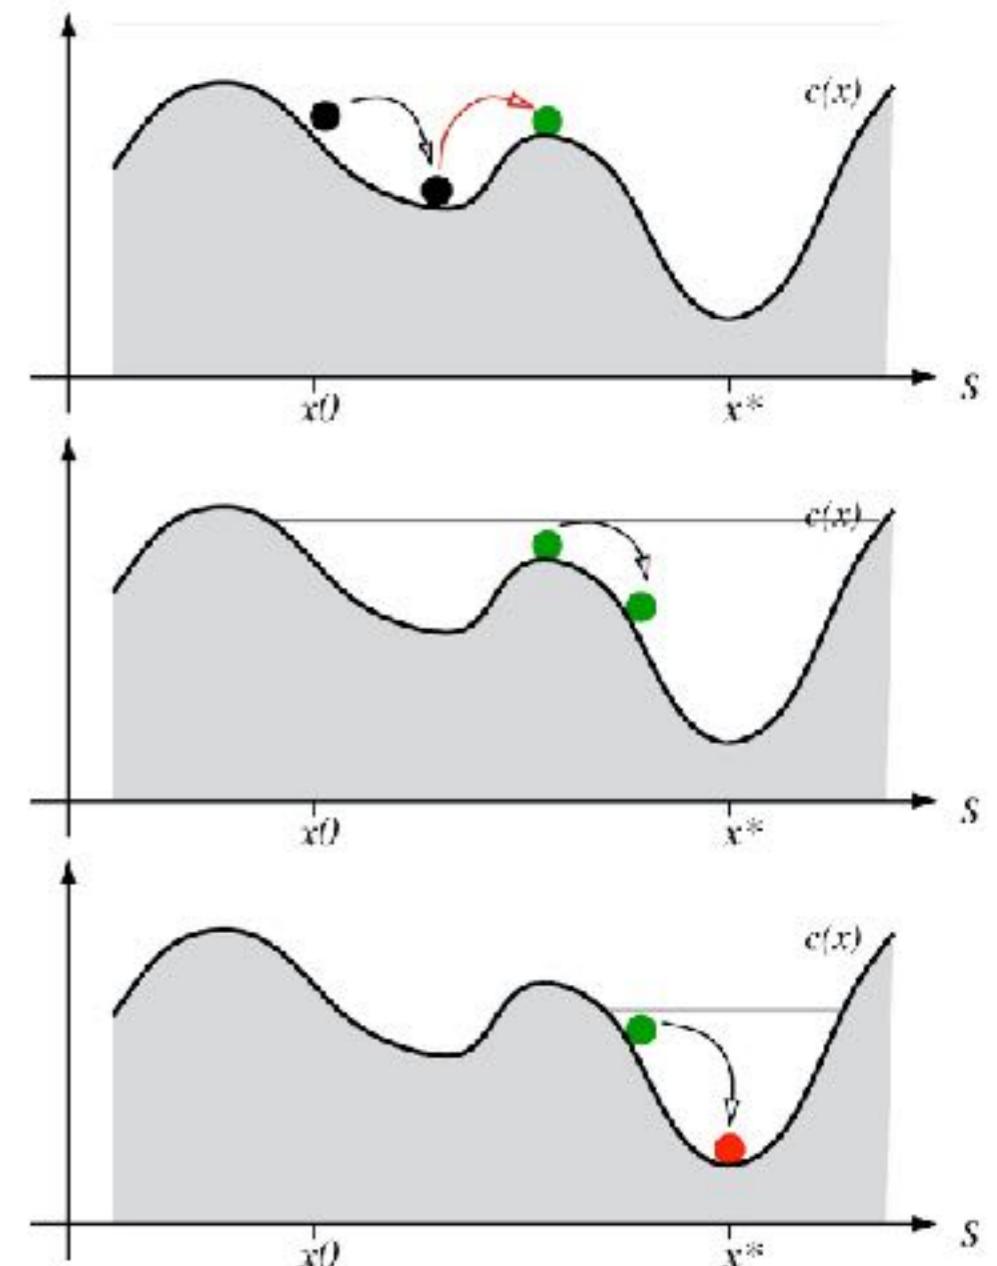
Directed Fuzzing as Optimisation Problem



- **Directed Greybox Fuzzing**
 - Assign **more energy** to seeds that are **closer** to the given targets!
 - **Problem** (Stochastic Gradient Descent)
 - If we **always** assign **more energy** to **closer** seeds, we **typically** reach only a **local** minimum, but **never** a **global** minimum distance!
 - **Solution** (Simulated Annealing)
 - **Sometimes** assign **more energy** to **further-away** seeds!
 - Approaches **global** minimum distance.

Directed Fuzzing as Optimisation Problem

- Simulated Annealing (SA)
 - **Exploration** phase:
 - Energy of **closer** seeds similar to energy of **further-away** seeds
 - **Exploitation** phase:
 - Energy of **closer** seeds is assigned to be **higher** and higher
 - Energy of **further-away** seeds is assigned to be **lower** and lower
 - We are increasing the “importance” of seed distance over time.



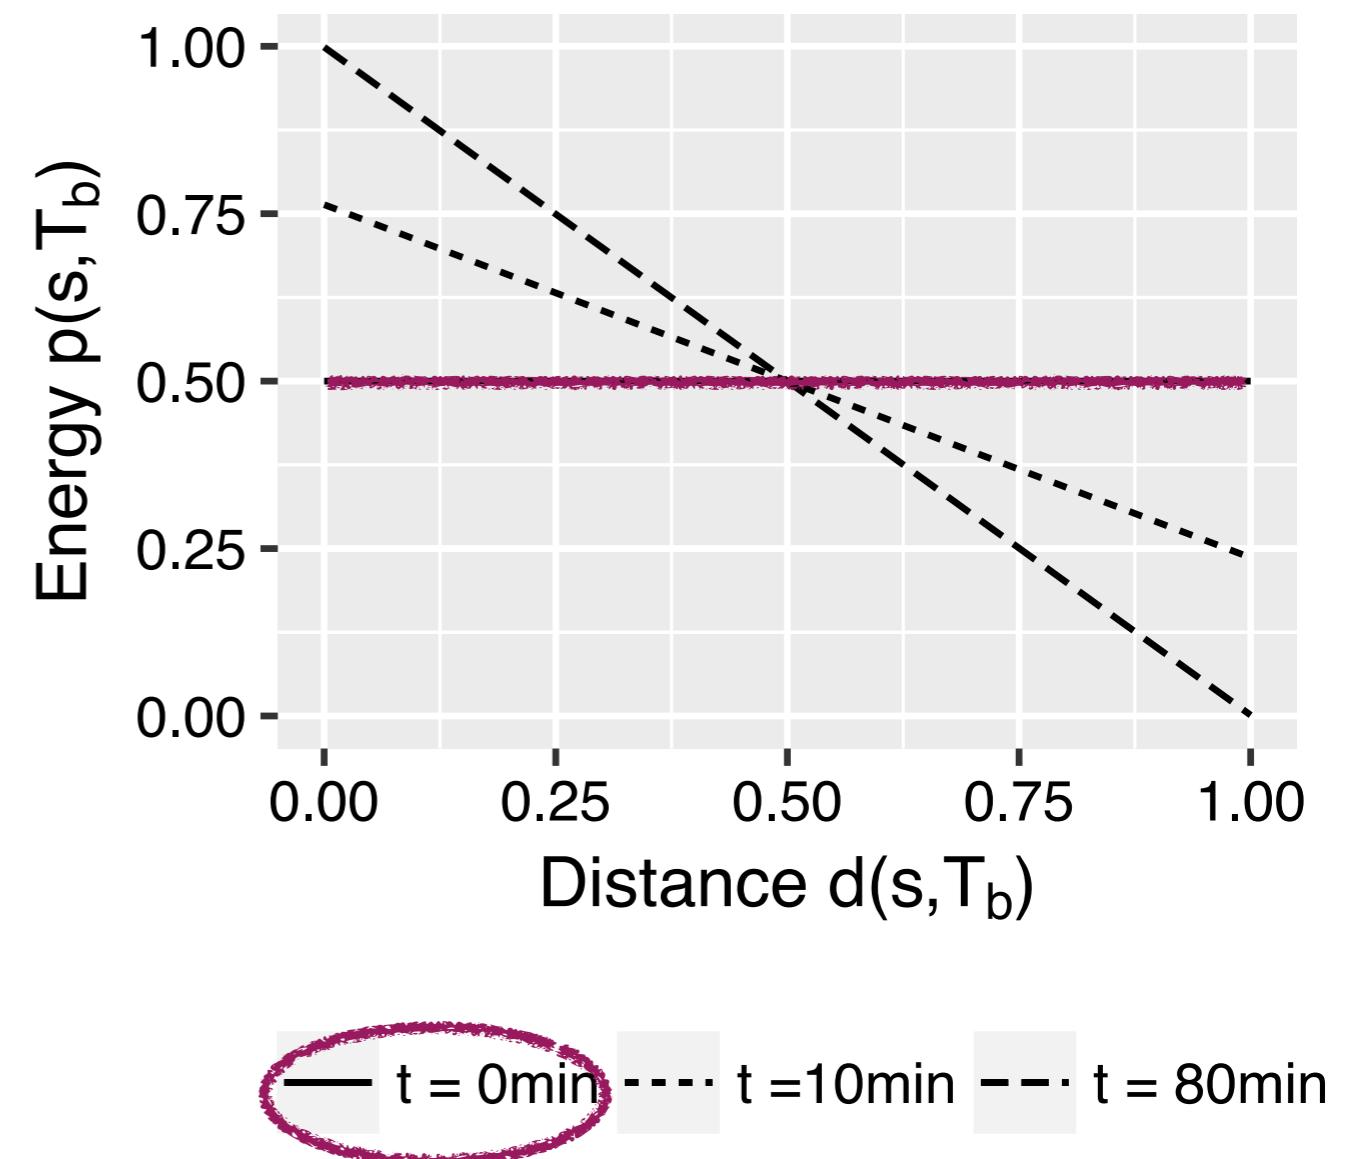
Directed Fuzzing as Optimisation Problem



- Simulated Annealing (SA)
 - Annealing from **metallurgy**: control the cooling of material to reduce defects (e.g., cracks or bubbles) in the material.
 - **Temperature** $T \in [0,1]$ specifies “importance” of distance.
 - At $T=1$, **exploration** (normal AFL)
 - At $T=0$, **exploitation** (gradient descent)
 - **Cooling schedule** controls (global) temperature
 - Classically, exponential cooling.

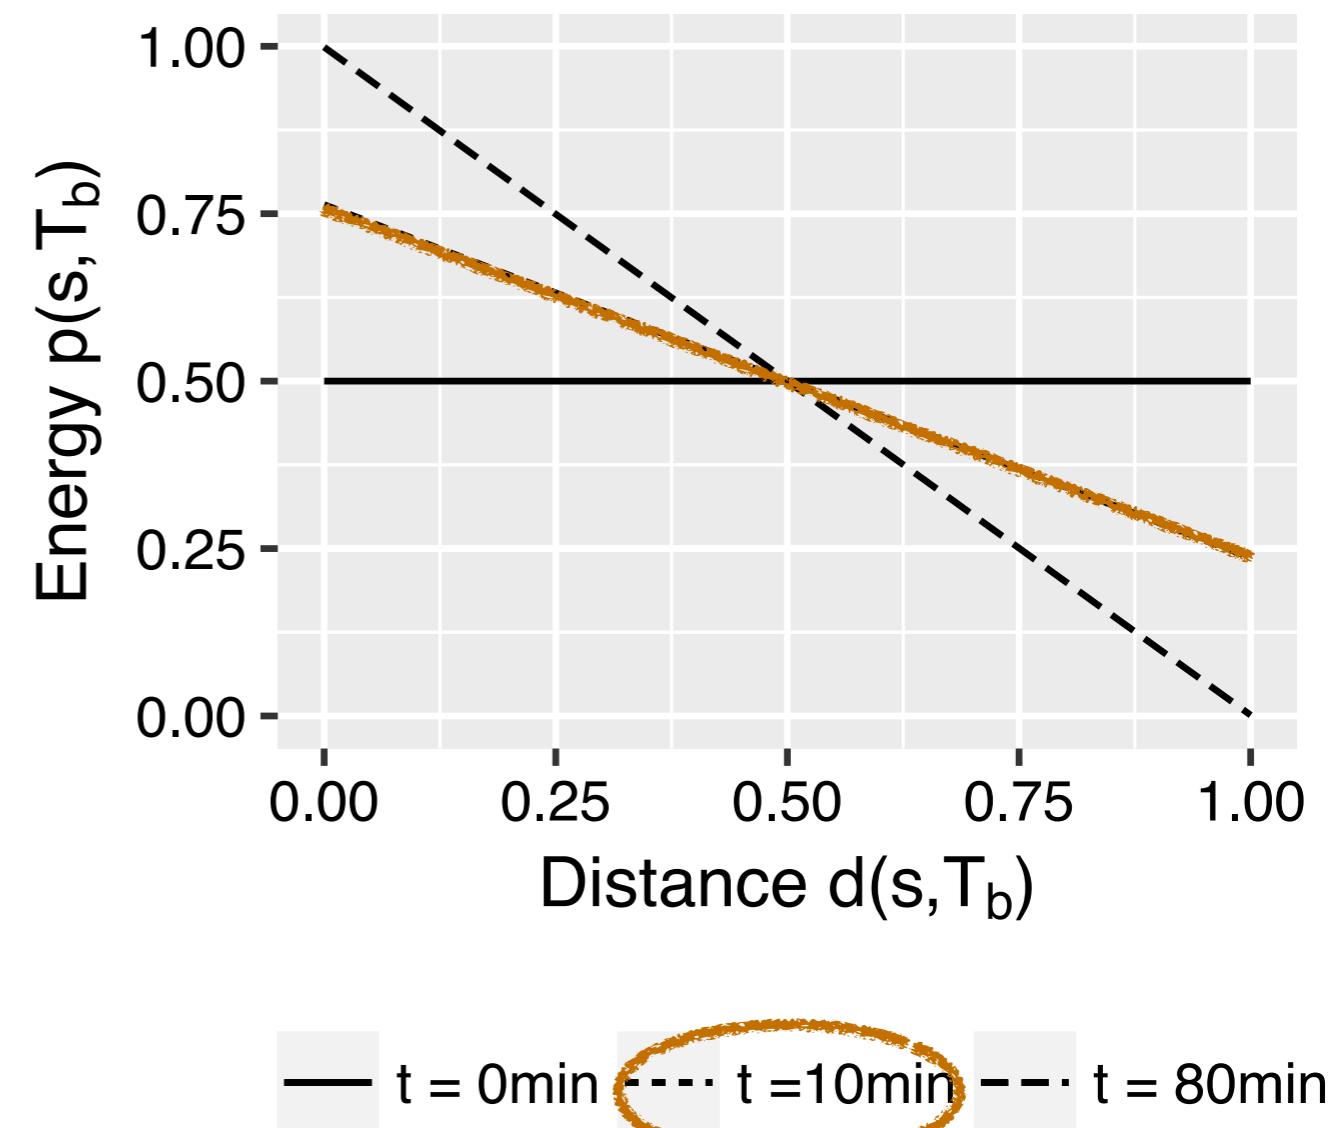
Directed Fuzzing as Optimisation Problem

- Integrating Simulated Annealing as power schedule
 - In the beginning ($t = 0\text{min}$), assign the **same energy to all seeds**.



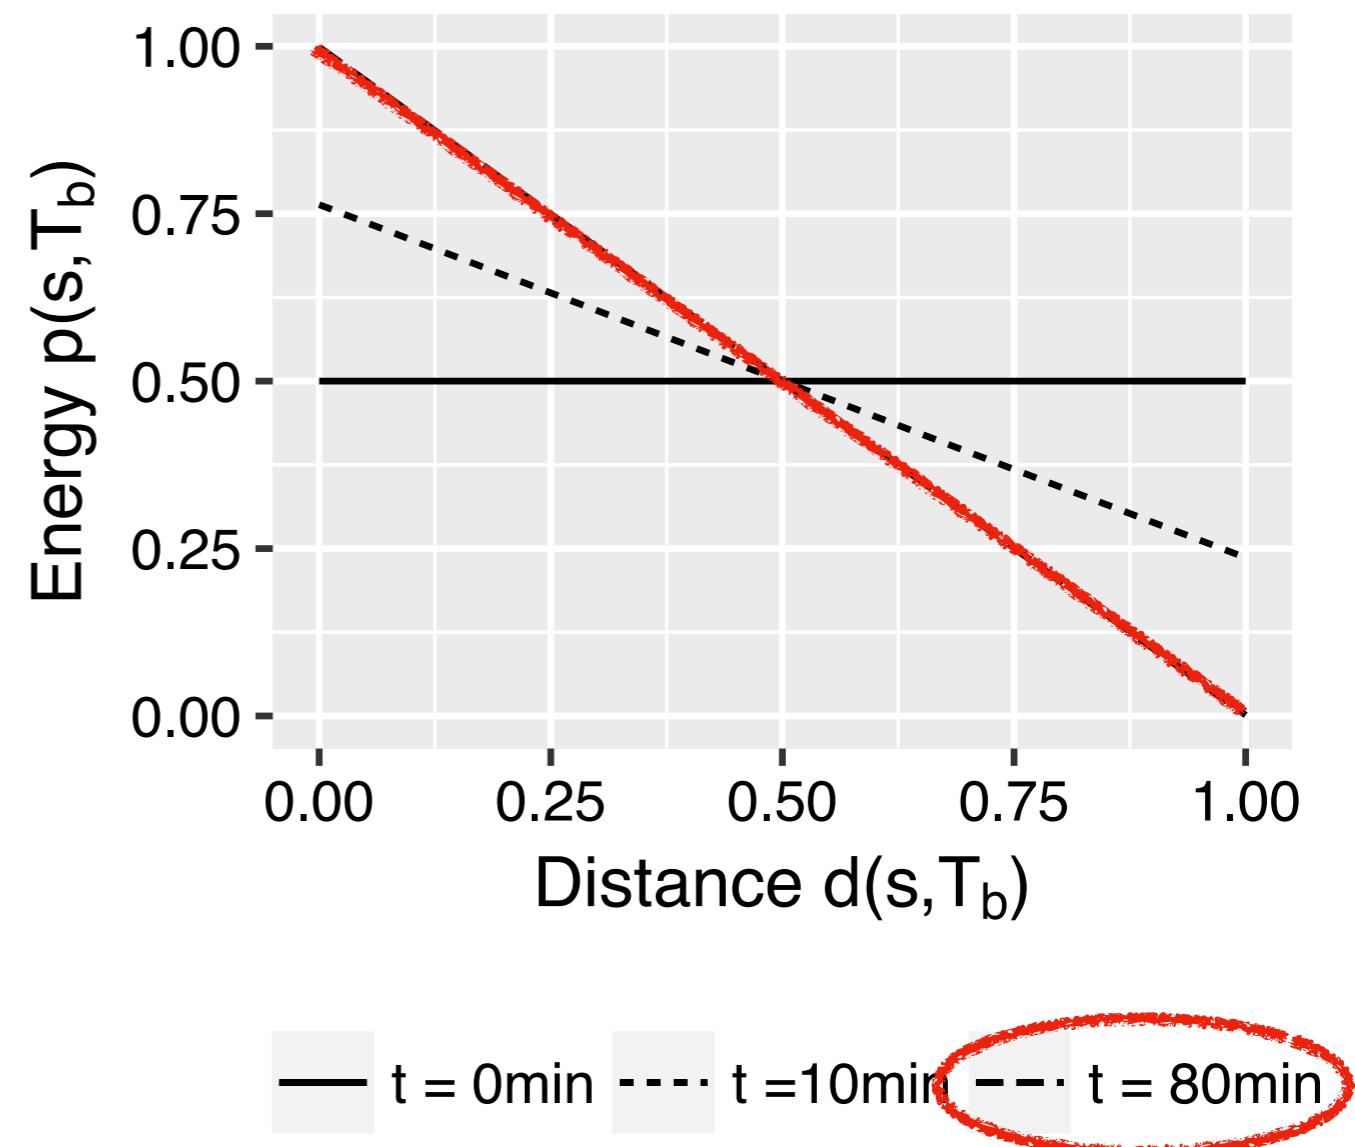
Directed Fuzzing as Optimisation Problem

- Integrating Simulated Annealing as power schedule
 - In the beginning ($t = 0\text{min}$), assign the **same energy to all seeds**.
 - Later ($t=10\text{min}$), assign a **bit more energy** to seeds that are **closer**.



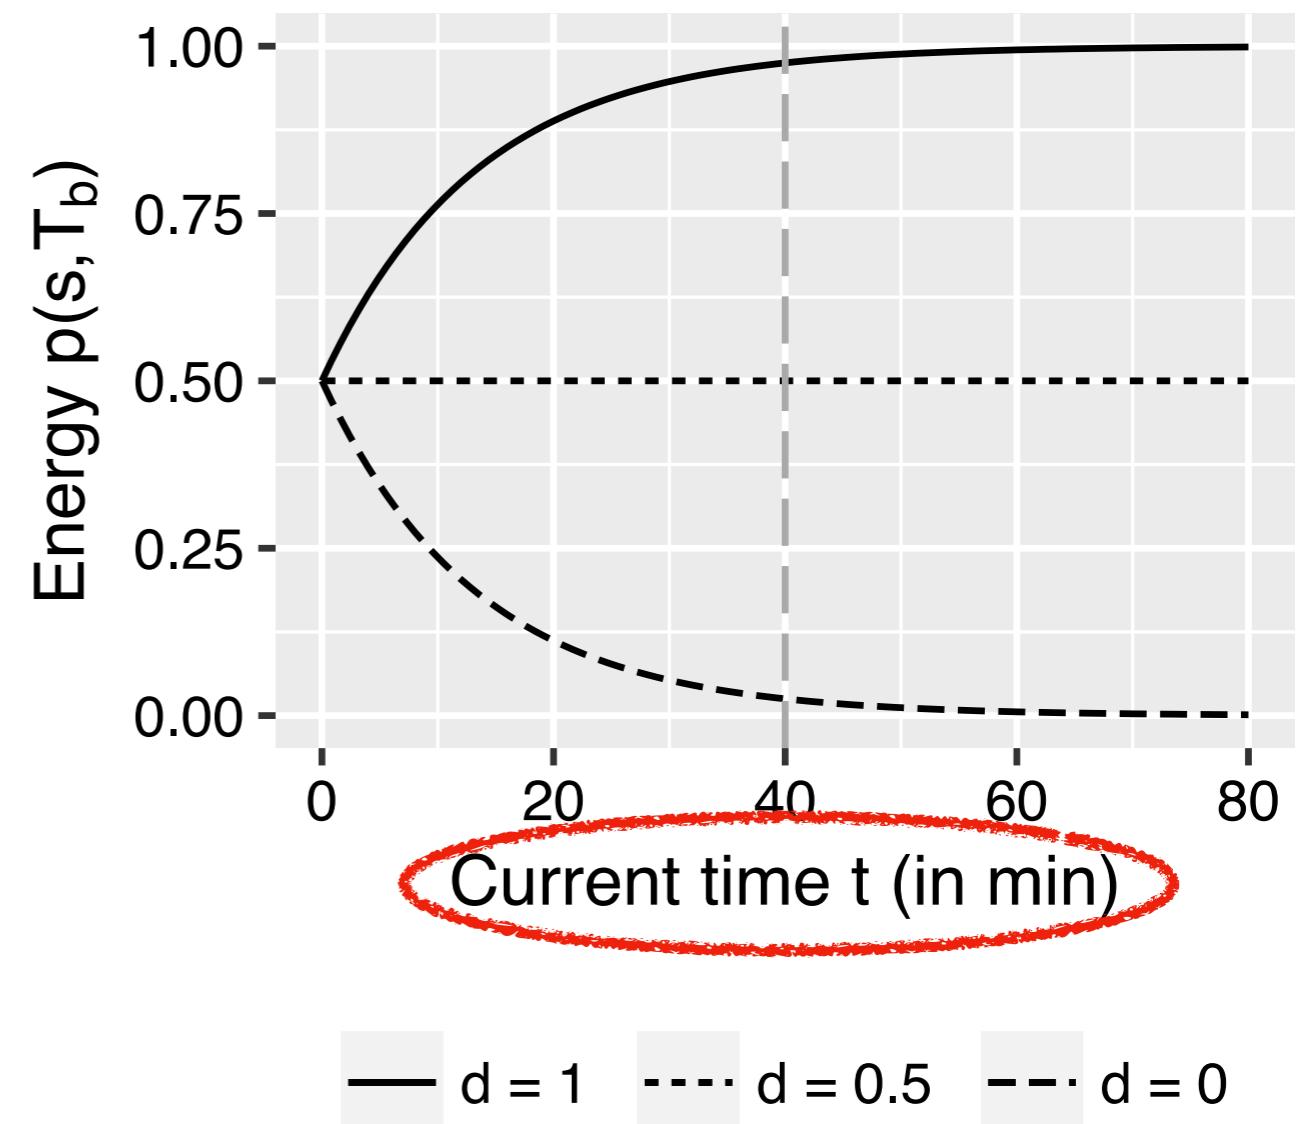
Directed Fuzzing as Optimisation Problem

- Integrating Simulated Annealing as power schedule
 - In the beginning ($t = 0\text{min}$), assign the **same energy to all seeds**.
 - Later ($t=10\text{min}$), assign a **bit more energy** to seeds that are **closer**.
 - At exploitation ($t=80\text{min}$), assign **maximal energy** to seeds that are **closest**.



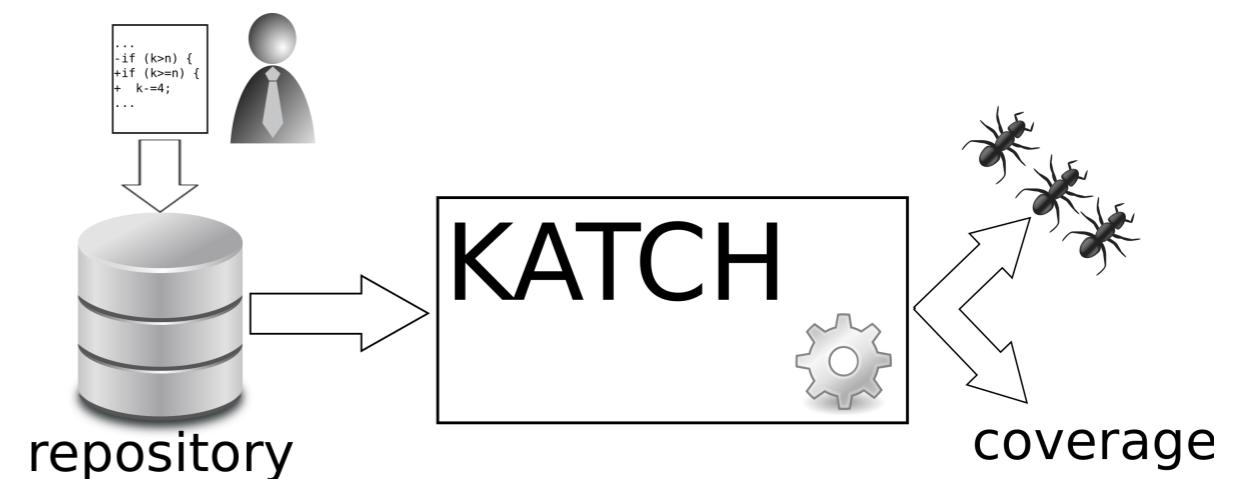
Directed Fuzzing as Optimisation Problem

- Integrating Simulated Annealing as power schedule
 - In the beginning ($t = 0\text{min}$), assign the **same energy to all seeds**.
 - Later ($t=10\text{min}$), assign a **bit more energy** to seeds that are **closer**.
 - At exploitation ($t=80\text{min}$), assign **maximal energy** to seeds that are **closest**.



Results

- **Patch Testing:** Reach changed statements
 - State-of-the-art in patch testing
 - **KATCH** (based on Klee symbolic exec. tool)
 - Experimental Setup
 - Reuse original **KATCH**-benchmark
 - Measure patch coverage (#changed BBs reached)
 - Measure vuln. detection (#errors discovered)



Results

- **Patch Testing:** Reach changed statements
 - State-of-the-art in patch testing
 - **KATCH** (based on Klee symbolic exec. tool)
 - Patch Coverage (#changed BBs reached)

	#Changed Basic Blocks	#Uncovered Changed BBs	KATCH	AFLGo
<i>Binutils</i>	852	702	135	159
<i>Diffutils</i>	166	108	63	64
Sum	1018	810	198	223

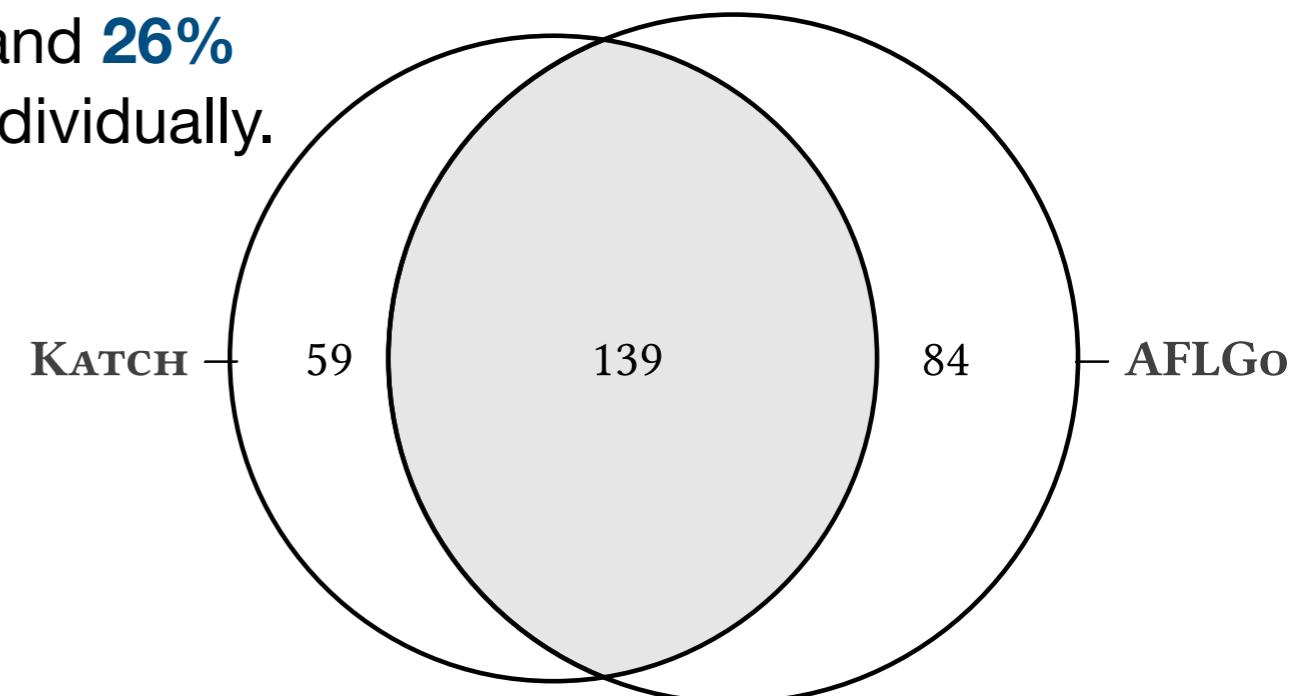
Results

- **Patch Testing:** Reach changed statements
 - State-of-the-art in patch testing
 - **KATCH** (based on Klee symbolic exec. tool)
 - Patch Coverage (#changed BBs reached)
 - While we would expect Klee to take a substantial lead, **AFLGo outperforms KATCH** in terms of patch coverage.

	#Changed Basic Blocks	#Uncovered Changed BBs	KATCH	AFLGO
<i>Binutils</i>	852	702	135	159
<i>Diffutils</i>	166	108	63	64
Sum	1018	810	198	223

Results

- **Patch Testing:** Reach changed statements
 - State-of-the-art in patch testing
 - **KATCH** (based on Klee symbolic exec. tool)
 - Patch Coverage (#changed BBs reached)
 - While we would expect Klee to take a substantial lead, **AFLGo outperforms KATCH** in terms of patch coverage.
 - **BUT: Together** they cover **42%** and **26%** more than **AFLGo** and **KATCH** individually.
They complement each other!



Results

- **Patch Testing:** Reach changed statements
 - State-of-the-art in patch testing
 - **KATCH** (based on Klee symbolic exec. tool)
 - Patch Coverage (#changed BBs reached)
 - Vulnerability Detection (#errors discovered)

	KATCH #Reports	AFLGo		
		#Reports ¹⁴	#New Reports	#CVEs
<i>Binutils</i>	7	4	12	7
<i>Diffutils</i>	0	N/A	1	0
Sum	7	4	13	7

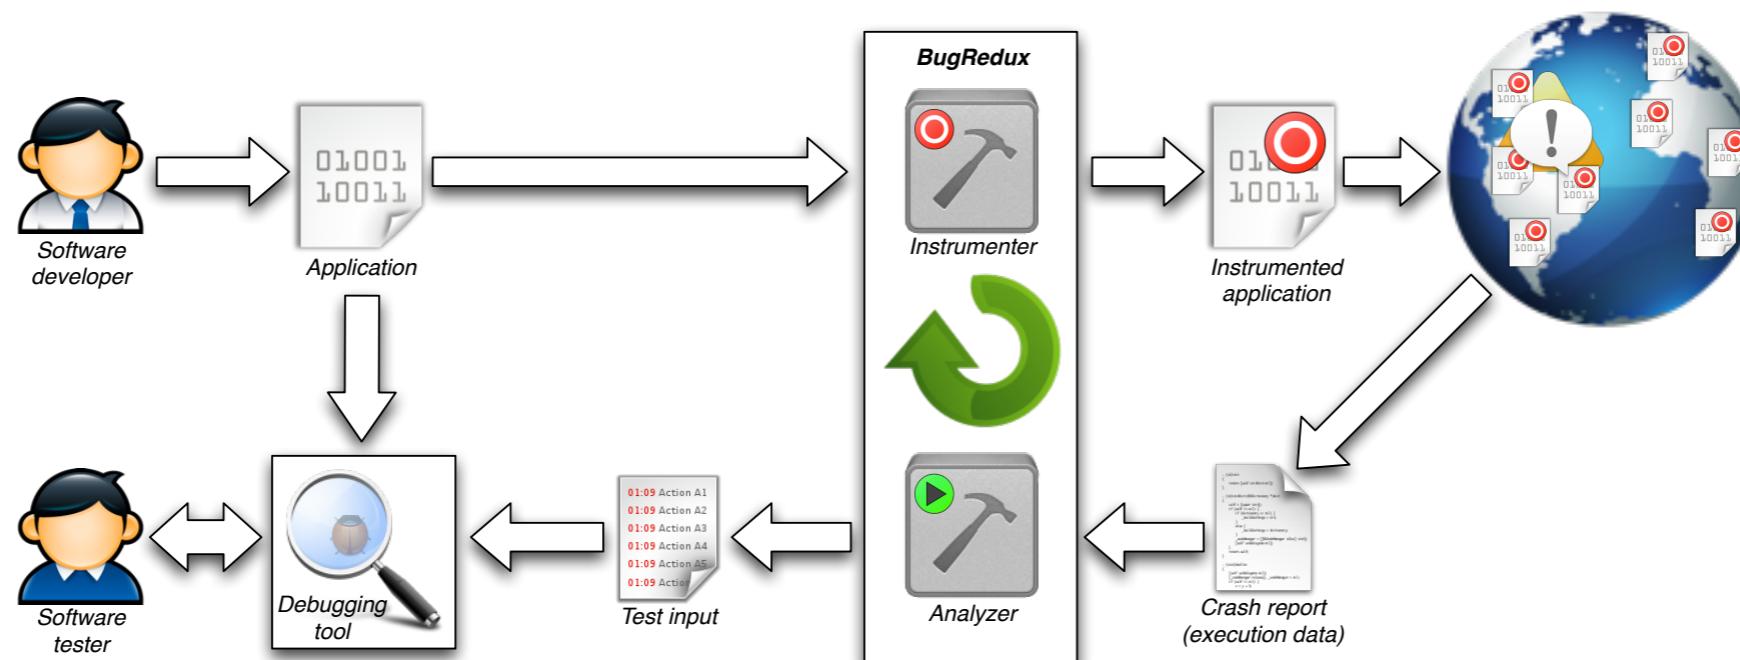
Results

- **Patch Testing:** Reach changed statements
 - State-of-the-art in patch testing
 - **KATCH** (based on Klee symbolic exec. tool)
 - Patch Coverage (#changed BBs reached)
 - Vulnerability Detection (#errors discovered)
 - **AFLGo found 13 previously unreported bugs (7 CVEs)** in addition to 4 of the 7 bugs that were found by **KATCH**.

	KATCH #Reports	AFLGo #Reports ¹⁴	#New Reports	#CVEs
<i>Binutils</i>	7	4	12	7
<i>Diffutils</i>	0	N/A	1	0
Sum	7	4	13	7

Results

- Crash Reproduction: Exercise stack trace
 - State-of-the-art in crash reproduction
 - **BugRedux** (based on Klee symbolic exec. tool)
 - Experimental Setup
 - Reuse original **BugRedux**-benchmark
 - Determine whether or not crash can be reproduced



Results

- **Crash Reproduction:** Exercise stack trace
 - State-of-the-art in crash reproduction
 - **BugRedux** (based on Klee symbolic exec. tool)
 - Experimental Setup
 - Reuse original **BugRedux**-benchmark
 - Determine whether or not crash can be reproduced

Subjects	BugRedux	AFLGo	Comments
sed.fault1	✗	✗	Takes two files as input
sed.fault2	✗	✓	
grep	✗	✓	
gzip.fault1	✗	✓	
gzip.fault2	✗	✓	
ncompress	✓	✓	
polymorph	✓	✓	

AFLGo reproduces
3 times more crashes!

Result Summary

- Our **directed greybox fuzzer** (AFLGo) **outperforms symbolic execution-based directed fuzzers** (KATCH & BugRedux)
 - in terms of **reaching more target locations** and
 - in terms of **detecting more vulnerabilities**,
 - on their own, original benchmark sets.

Motivation



- **Directed Fuzzing:** classical **constraint satisfaction prob.**
 - **Program analysis** to identify **program paths** that reach given program locations.
 - **Symbolic Execution** to derive **path conditions** for any of the identified paths.
 - **Constraint Solving** to find an **input** that
 - **satisfies the path condition** and thus
 - **reaches a program location** that was given.

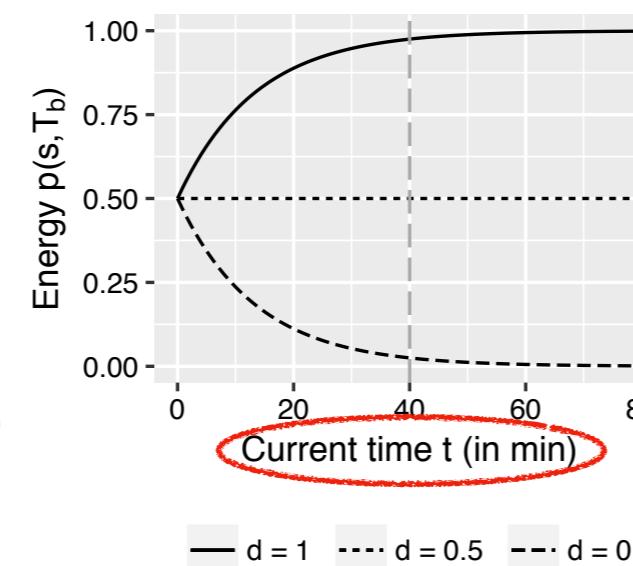
Requires heavy-weight machinery!

Presented by Abhik Roychoudhury

Directed Greybox Fuzzing

Directed Fuzzing as Optimisation Problem

- Integrating Simulated Annealing as power schedule
 - In the beginning ($t = 0\text{min}$), assign the **same energy to all seeds**.
 - Later ($t=10\text{min}$), assign a **bit more energy** to seeds that are **closer**.
 - At exploitation ($t=80\text{min}$), assign **maximal energy** to seeds that are **closest**.



Overview

- **Directed Fuzzing as optimisation problem!**
 1. **Instrumentation Time:**
 1. Extract **call graph** (CG) and **control-flow graphs** (CFGs).
 2. For each **BB**, compute **distance** to target locations.
 3. Instrument program to **aggregate distance values**.
 2. **Runtime, for each input**
 1. collect coverage and distance **information**, and
 2. decide **how long to be fuzzed** based on distance.
 - If input is **closer** to the targets, it is fuzzed for **longer**.
 - If input is **further away** from the targets, it is fuzzed for **shorter**.



Presented by Abhik Roychoudhury

Directed Greybox Fuzzing

Presented by Abhik Roychoudhury

Directed Greybox Fuzzing

Result Summary



- Our **directed greybox fuzzer** (AFLGo) **outperforms*** **symbolic execution-based directed fuzzers** (KATCH & BugRedux)
 - in terms of **reaching more target locations** and
 - in terms of **detecting more vulnerabilities**,
 - on their own, original benchmark sets.

Motivation



- **Directed Fuzzing:** classical **constraint satisfaction prob.**
 - **Program analysis** to identify **program paths** that reach given program locations.
 - **Symbolic Execution** to derive **path conditions** for any of the identified paths.
 - **Constraint Solving** to find an **input** that
 - **satisfies the path condition** and thus
 - **reaches a program location** that was given.

Requires heavy-weight machinery!

Overview

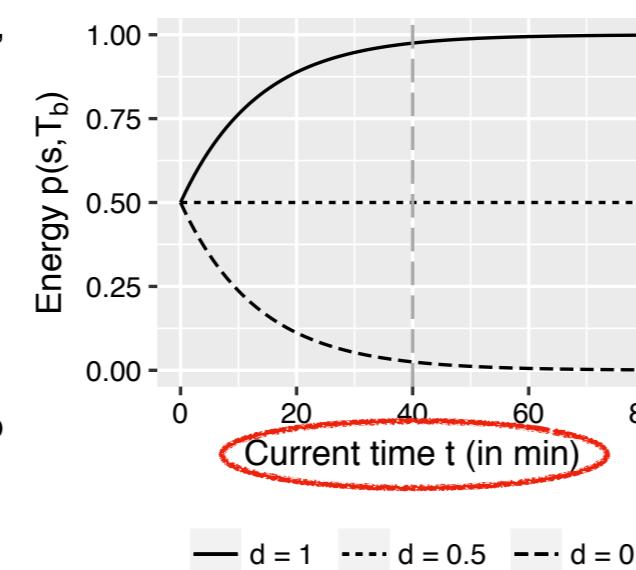
- **Directed Fuzzing as optimisation problem!**
 1. **Instrumentation Time:**
 1. Extract **call graph** (CG) and **control-flow graphs** (CFGs).
 2. For each **BB**, compute **distance** to target locations.
 3. Instrument program to **aggregate distance values**.
 2. **Runtime, for each input**
 1. collect coverage and distance **information**, and
 2. decide **how long to be fuzzed** based on distance.
 - If input is **closer** to the targets, it is fuzzed for **longer**.
 - If input is **further away** from the targets, it is fuzzed for **shorter**.

Questions?

Directed Fuzzing as Optimisation Problem



- Integrating Simulated Annealing as power schedule
 - In the beginning ($t = 0\text{min}$), assign the **same energy to all seeds**.
 - Later ($t=10\text{min}$), assign a **bit more energy** to seeds that are **closer**.
 - At exploitation ($t=80\text{min}$), assign **maximal energy** to seeds that are **closest**.



Result Summary



- Our **directed greybox fuzzer** (AFLGo) **outperforms*** **symbolic execution-based directed fuzzers** (KATCH & BugRedux)
 - in terms of **reaching more target locations** and
 - in terms of **detecting more vulnerabilities**,
 - on their own, original benchmark sets.